



# Hibernate Reference Documentation

Version: 3.0

---

# Table of Contents

Preface .....	viii
<b>1. Quickstart with Tomcat</b> .....	1
1.1. Getting started with Hibernate .....	1
1.2. First persistent class .....	3
1.3. Mapping the cat .....	4
1.4. Playing with cats .....	5
1.5. Finally .....	7
<b>2. Architecture</b> .....	8
2.1. Overview .....	8
2.2. Instance states .....	10
2.3. JMX Integration .....	10
2.4. JCA Support .....	11
<b>3. Configuration</b> .....	12
3.1. Programmatic configuration .....	12
3.2. Obtaining a SessionFactory .....	12
3.3. JDBC connections .....	13
3.4. Optional configuration properties .....	14
3.4.1. SQL Dialects .....	19
3.4.2. Outer Join Fetching .....	20
3.4.3. Binary Streams .....	20
3.4.4. Second-level and query cache .....	20
3.4.5. Transaction strategy configuration .....	20
3.4.6. JNDI-bound SessionFactory .....	21
3.4.7. Query Language Substitution .....	22
3.4.8. Hibernate statistics .....	22
3.5. Logging .....	22
3.6. Implementing a NamingStrategy .....	23
3.7. XML configuration file .....	23
<b>4. Persistent Classes</b> .....	25
4.1. A simple POJO example .....	25
4.1.1. Declare accessors and mutators for persistent fields .....	26
4.1.2. Implement a no-argument constructor .....	26
4.1.3. Provide an identifier property (optional) .....	26
4.1.4. Prefer non-final classes (optional) .....	27
4.2. Implementing inheritance .....	27
4.3. Implementing equals() and hashCode() .....	27
4.4. Dynamic models .....	28
<b>5. Basic O/R Mapping</b> .....	31
5.1. Mapping declaration .....	31
5.1.1. Doctype .....	32
5.1.2. hibernate-mapping .....	32
5.1.3. class .....	33
5.1.4. id .....	35
5.1.4.1. Generator .....	36
5.1.4.2. Hi/lo algorithm .....	37
5.1.4.3. UUID algorithm .....	37
5.1.4.4. Identity columns and sequences .....	37
5.1.4.5. Assigned identifiers .....	38

5.1.4.6. Primary keys assigned by triggers .....	38
5.1.5. composite-id .....	38
5.1.6. discriminator .....	39
5.1.7. version (optional) .....	39
5.1.8. timestamp (optional) .....	40
5.1.9. property .....	40
5.1.10. many-to-one .....	42
5.1.11. one-to-one .....	43
5.1.12. component, dynamic-component .....	44
5.1.13. properties .....	45
5.1.14. subclass .....	46
5.1.15. joined-subclass .....	47
5.1.16. union-subclass .....	48
5.1.17. join .....	48
5.1.18. key .....	49
5.1.19. column and formula elements .....	50
5.1.20. import .....	50
5.1.21. any .....	51
5.2. Hibernate Types .....	51
5.2.1. Entities and values .....	52
5.2.2. Basic value types .....	52
5.2.3. Custom value types .....	53
5.3. SQL quoted identifiers .....	54
5.4. Metadata alternatives .....	54
5.4.1. Using XDoclet markup .....	54
5.4.2. Using JDK 5.0 Annotations .....	56
<b>6. Collection Mapping</b> .....	<b>58</b>
6.1. Persistent collections .....	58
6.2. Collection mappings .....	58
6.2.1. Collection foreign keys .....	59
6.2.2. Collection elements .....	60
6.2.3. Indexed collections .....	60
6.2.4. Collections of values and many-to-many associations .....	61
6.2.5. One-to-many associations .....	62
6.3. Advanced collection mappings .....	63
6.3.1. Sorted collections .....	63
6.3.2. Bidirectional associations .....	64
6.3.3. Ternary associations .....	65
6.3.4. Using an <idbag> .....	65
6.4. Collection examples .....	66
<b>7. Association Mappings</b> .....	<b>69</b>
7.1. Introduction .....	69
7.2. Unidirectional associations .....	69
7.2.1. many to one .....	69
7.2.2. one to one .....	69
7.2.3. one to many .....	70
7.3. Unidirectional associations with join tables .....	71
7.3.1. one to many .....	71
7.3.2. many to one .....	71
7.3.3. one to one .....	72
7.3.4. many to many .....	72
7.4. Bidirectional associations .....	72

7.4.1. one to many / many to one .....	73
7.4.2. one to one .....	73
7.5. Bidirectional associations with join tables .....	74
7.5.1. one to many / many to one .....	74
7.5.2. one to one .....	74
7.5.3. many to many .....	75
<b>8. Component Mapping .....</b>	<b>77</b>
8.1. Dependent objects .....	77
8.2. Collections of dependent objects .....	78
8.3. Components as Map indices .....	79
8.4. Components as composite identifiers .....	79
8.5. Dynamic components .....	81
<b>9. Inheritance Mapping .....</b>	<b>82</b>
9.1. The Three Strategies .....	82
9.1.1. Table per class hierarchy .....	82
9.1.2. Table per subclass .....	82
9.1.3. Table per subclass, using a discriminator .....	83
9.1.4. Mixing table per class hierarchy with table per subclass .....	83
9.1.5. Table per concrete class .....	84
9.1.6. Table per concrete class, using implicit polymorphism .....	84
9.1.7. Mixing implicit polymorphism with other inheritance mappings .....	85
9.2. Limitations .....	86
<b>10. Working with objects .....</b>	<b>88</b>
10.1. Hibernate object states .....	88
10.2. Making objects persistent .....	88
10.3. Loading an object .....	89
10.4. Querying .....	90
10.4.1. Executing queries .....	90
10.4.1.1. Iterating results .....	90
10.4.1.2. Queries that return tuples .....	91
10.4.1.3. Scalar results .....	91
10.4.1.4. Bind parameters .....	91
10.4.1.5. Pagination .....	92
10.4.1.6. Scrollable iteration .....	92
10.4.1.7. Externalizing named queries .....	93
10.4.2. Filtering collections .....	93
10.4.3. Criteria queries .....	93
10.4.4. Queries in native SQL .....	94
10.5. Modifying persistent objects .....	94
10.6. Modifying detached objects .....	94
10.7. Automatic state detection .....	95
10.8. Deleting persistent objects .....	96
10.9. Replicating object between two different datastores .....	96
10.10. Flushing the Session .....	97
10.11. Transitive persistence .....	98
10.12. Using metadata .....	99
<b>11. Transactions And Concurrency .....</b>	<b>100</b>
11.1. Session and transaction scopes .....	100
11.1.1. Unit of work .....	100
11.1.2. Application transactions .....	101
11.1.3. Considering object identity .....	102
11.1.4. Common issues .....	102

11.2. Database transaction demarcation .....	103
11.2.1. Non-managed environment .....	103
11.2.2. Using JTA .....	104
11.2.3. Exception handling .....	105
11.3. Optimistic concurrency control .....	105
11.3.1. Application version checking .....	105
11.3.2. Long session and automatic versioning .....	106
11.3.3. Detached objects and automatic versioning .....	107
11.3.4. Customizing automatic versioning .....	107
11.4. Pessimistic Locking .....	107
<b>12. Interceptors and events</b> .....	109
12.1. Interceptors .....	109
12.2. Event system .....	110
12.3. Hibernate declarative security .....	111
<b>13. Batch processing</b> .....	112
13.1. Batch inserts .....	112
13.2. Batch updates .....	112
<b>14. HQL: The Hibernate Query Language</b> .....	114
14.1. Case Sensitivity .....	114
14.2. The from clause .....	114
14.3. Associations and joins .....	114
14.4. The select clause .....	115
14.5. Aggregate functions .....	116
14.6. Polymorphic queries .....	116
14.7. The where clause .....	117
14.8. Expressions .....	118
14.9. The order by clause .....	120
14.10. The group by clause .....	121
14.11. Subqueries .....	121
14.12. HQL examples .....	122
14.13. Tips & Tricks .....	123
<b>15. Criteria Queries</b> .....	125
15.1. Creating a Criteria instance .....	125
15.2. Narrowing the result set .....	125
15.3. Ordering the results .....	126
15.4. Associations .....	126
15.5. Dynamic association fetching .....	127
15.6. Example queries .....	127
15.7. Projections, aggregation and grouping .....	128
15.8. Detached queries and subqueries .....	129
<b>16. Native SQL</b> .....	130
16.1. Creating a native SQL Query .....	130
16.2. Alias and property references .....	130
16.3. Named SQL queries .....	131
16.3.1. Using return-property to explicitly specify column/alias names .....	131
16.3.2. Using stored procedures for querying .....	132
16.3.2.1. Rules/limitations for using stored procedures .....	132
16.4. Custom SQL for create, update and delete .....	133
16.5. Custom SQL for loading .....	134
<b>17. Filtering data</b> .....	135
17.1. Hibernate filters .....	135
<b>18. XML Mapping</b> .....	137

18.1. Working with XML data .....	137
18.1.1. Specifying XML and class mapping together .....	137
18.1.2. Specifying only an XML mapping .....	137
18.2. XML mapping metadata .....	138
18.3. Manipulating XML data .....	139
<b>19. Improving performance .....</b>	<b>141</b>
19.1. Fetching strategies .....	141
19.1.1. Tuning fetch strategies .....	142
19.1.2. Single-ended association proxies .....	142
19.1.3. Initializing collections and proxies .....	144
19.1.4. Using batch fetching .....	145
19.1.5. Using lazy property fetching .....	145
19.2. The Second Level Cache .....	146
19.2.1. Cache mappings .....	147
19.2.2. Strategy: read only .....	147
19.2.3. Strategy: read/write .....	147
19.2.4. Strategy: nonstrict read/write .....	148
19.2.5. Strategy: transactional .....	148
19.3. Managing the caches .....	148
19.4. The Query Cache .....	149
19.5. Understanding Collection performance .....	150
19.5.1. Taxonomy .....	150
19.5.2. Lists, maps, idbags and sets are the most efficient collections to update .....	151
19.5.3. Bags and lists are the most efficient inverse collections .....	151
19.5.4. One shot delete .....	151
19.6. Monitoring performance .....	152
19.6.1. Monitoring a SessionFactory .....	152
19.6.2. Metrics .....	153
<b>20. Toolset Guide .....</b>	<b>155</b>
20.1. Automatic schema generation .....	155
20.1.1. Customizing the schema .....	155
20.1.2. Running the tool .....	157
20.1.3. Properties .....	157
20.1.4. Using Ant .....	158
20.1.5. Incremental schema updates .....	158
20.1.6. Using Ant for incremental schema updates .....	158
<b>21. Example: Parent/Child .....</b>	<b>159</b>
21.1. A note about collections .....	159
21.2. Bidirectional one-to-many .....	159
21.3. Cascading lifecycle .....	160
21.4. Cascades and unsaved-value .....	161
21.5. Conclusion .....	162
<b>22. Example: Weblog Application .....</b>	<b>163</b>
22.1. Persistent Classes .....	163
22.2. Hibernate Mappings .....	164
22.3. Hibernate Code .....	165
<b>23. Example: Various Mappings .....</b>	<b>169</b>
23.1. Employer/Employee .....	169
23.2. Author/Work .....	170
23.3. Customer/Order/Product .....	172
23.4. Miscellaneous example mappings .....	174
23.4.1. "Typed" one-to-one association .....	174

23.4.2. Composite key example .....	174
23.4.3. Content based discrimination .....	176
23.4.4. Associations on alternate keys .....	177
<b>24. Best Practices .....</b>	<b>179</b>

---

# Preface

Working with object-oriented software and a relational database can be cumbersome and time consuming in today's enterprise environments. Hibernate is an object/relational mapping tool for Java environments. The term object/relational mapping (ORM) refers to the technique of mapping a data representation from an object model to a relational data model with a SQL-based schema.

Hibernate not only takes care of the mapping from Java classes to database tables (and from Java data types to SQL data types), but also provides data query and retrieval facilities and can significantly reduce development time otherwise spent with manual data handling in SQL and JDBC.

Hibernate's goal is to relieve the developer from 95 percent of common data persistence related programming tasks. Hibernate may not be the best solution for data-centric applications that only use stored-procedures to implement the business logic in the database, it is most useful with object-oriented domain models and business logic in the Java-based middle-tier. However, Hibernate can certainly help you to remove or encapsulate vendor-specific SQL code and will help with the common task of result set translation from a tabular representation to a graph of objects.

If you are new to Hibernate and Object/Relational Mapping or even Java, please follow these steps:

1. Read Chapter 1, *Quickstart with Tomcat* for a 30 minute tutorial, using Tomcat.
2. Read Chapter 2, *Architecture* to understand the environments where Hibernate can be used.
3. Have a look at the `eg/` directory in the Hibernate distribution, it contains a simple standalone application. Copy your JDBC driver to the `lib/` directory and edit `etc/hibernate.properties`, specifying correct values for your database. From a command prompt in the distribution directory, type `ant eg` (using Ant), or under Windows, type `build eg`.
4. Use this reference documentation as your primary source of information. Consider reading *Hibernate in Action* (<http://www.manning.com/bauer>) if you need more help with application design or if you prefer a step-by-step tutorial. Also visit <http://caveatemptor.hibernate.org> and download the example application for Hibernate in Action.
5. FAQs are answered on the Hibernate website.
6. Third party demos, examples and tutorials are linked on the Hibernate website.
7. The Community Area on the Hibernate website is a good source for design patterns and various integration solutions (Tomcat, JBoss, Struts, EJB, etc.).

If you have questions, use the user forum linked on the Hibernate website. We also provide a JIRA issue tracking system for bug reports and feature requests. If you are interested in the development of Hibernate, join the developer mailing list. If you are interested in translating this documentation into your language, contact us on the developer mailing list.

Commercial development support, production support and training for Hibernate is available through JBoss Inc. (see <http://www.hibernate.org/SupportTraining/>). Hibernate is a project of the JBoss Professional Open Source product suite.



---

# Chapter 1. Quickstart with Tomcat

## 1.1. Getting started with Hibernate

This tutorial explains a setup of Hibernate 3.0 with the Apache Tomcat servlet container for a web-based application. Hibernate works well in a managed environment with all major J2EE application servers, or even in standalone Java applications. The database system used in this tutorial is PostgreSQL 7.4, support for other database is only a matter of changing the Hibernate SQL dialect configuration and connection properties.

First, we have to copy all required libraries to the Tomcat installation. We use a separate web context (`webapps/quickstart`) for this tutorial, so we've to consider both the global library search path (`TOMCAT/common/lib`) and the classloader at the context level in `webapps/quickstart/WEB-INF/lib` (for JAR files) and `webapps/quickstart/WEB-INF/classes`. We refer to both classloader levels as the global classpath and the context classpath.

Now, copy the libraries to the two classpaths:

1. Copy the JDBC driver for the database to the global classpath. This is required for the DBCP connection pool software which comes bundled with Tomcat. Hibernate uses JDBC connections to execute SQL on the database, so you either have to provide pooled JDBC connections or configure Hibernate to use one of the directly supported pools (C3P0, Proxool). For this tutorial, copy the `pg74jdbc3.jar` library (for PostgreSQL 7.4 and JDK 1.4) to the global classloaders path. If you'd like to use a different database, simply copy its appropriate JDBC driver.
2. Never copy anything else into the global classloader path in Tomcat, or you will get problems with various tools, including Log4j, commons-logging and others. Always use the context classpath for each web application, that is, copy libraries to `WEB-INF/lib` and your own classes and configuration/property files to `WEB-INF/classes`. Both directories are in the context level classpath by default.
3. Hibernate is packaged as a JAR library. The `hibernate3.jar` file should be copied in the context classpath together with other classes of the application. Hibernate requires some 3rd party libraries at runtime, these come bundled with the Hibernate distribution in the `lib/` directory; see Table 1.1, “Hibernate 3rd party libraries”. Copy the required 3rd party libraries to the context classpath.

**Table 1.1. Hibernate 3rd party libraries**

Library	Description
antlr (required)	Hibernate uses ANTLR to produce query parsers, this library is also needed at runtime.
dom4j (required)	Hibernate uses dom4j to parse XML configuration and XML mapping metadata files.
CGLIB (required)	Hibernate uses the code generation library to enhance classes at runtime (in combination with Java reflection).
Commons Collections, Commons Logging (required)	Hibernate uses various utility libraries from the Apache Jakarta Commons project.
EHCache (required)	Hibernate can use various cache providers for the second-level cache. EHCache is the default cache provider if not changed in the configura-

Library	Description
	tion.
Log4j (optional)	Hibernate uses the Commons Logging API, which in turn can use Log4j as the underlying logging mechanism. If the Log4j library is available in the context library directory, Commons Logging will use Log4j and the <code>log4j.properties</code> configuration in the context classpath. An example properties file for Log4j is bundled with the Hibernate distribution. So, copy <code>log4j.jar</code> and the configuration file (from <code>src/</code> ) to your context classpath if you want to see whats going on behind the scenes.
Required or not?	Have a look at the file <code>lib/version.properties</code> in the Hibernate distribution. This is an up-to-date list of 3rd party libraries distributed with Hibernate. You will find all required and optional libraries listed there (note that "buildtime required" here means for Hibernate's build, not your application).

We now set up the database connection pooling and sharing in both Tomcat and Hibernate. This means Tomcat will provide pooled JDBC connections (using its builtin DBCP pooling feature), Hibernate requests theses connections through JNDI. Tomcat binds the connection pool to JNDI, we add a resource declaration to Tomcats main configuration file, `TOMCAT/conf/server.xml`:

```
<Context path="/quickstart" docBase="quickstart">
  <Resource name="jdbc/quickstart" scope="Shareable" type="javax.sql.DataSource"/>
  <ResourceParams name="jdbc/quickstart">
    <parameter>
      <name>factory</name>
      <value>org.apache.commons.dbcp.BasicDataSourceFactory</value>
    </parameter>

    <!-- DBCP database connection settings -->
    <parameter>
      <name>url</name>
      <value>jdbc:postgresql://localhost/quickstart</value>
    </parameter>
    <parameter>
      <name>driverClassName</name><value>org.postgresql.Driver</value>
    </parameter>
    <parameter>
      <name>username</name>
      <value>quickstart</value>
    </parameter>
    <parameter>
      <name>password</name>
      <value>secret</value>
    </parameter>

    <!-- DBCP connection pooling options -->
    <parameter>
      <name>maxWait</name>
      <value>3000</value>
    </parameter>
    <parameter>
      <name>maxIdle</name>
      <value>100</value>
    </parameter>
    <parameter>
      <name>maxActive</name>
      <value>10</value>
    </parameter>
  </ResourceParams>
</Context>
```

The context we configure in this example is named `quickstart`, its base is the `TOMCAT/webapp/quickstart` directory. To access any servlets, call the path `http://localhost:8080/quickstart` in your browser (of course, adding the name of the servlet as mapped in your `web.xml`). You may also go ahead and create a simple servlet now that has an empty `process()` method.

Tomcat provides connections now through JNDI at `java:comp/env/jdbc/quickstart`. If you have trouble getting the connection pool running, refer to the Tomcat documentation. If you get JDBC driver exception messages, try to setup JDBC connection pool without Hibernate first. Tomcat & JDBC tutorials are available on the Web.

Your next step is to configure Hibernate. Hibernate has to know how it should obtain JDBC connections. We use Hibernate's XML-based configuration. The other approach, using a properties file, is equivalent in features, but doesn't offer any advantages. We use the XML configuration because it is usually more convenient. The XML configuration file is placed in the context classpath (`WEB-INF/classes`), as `hibernate.cfg.xml`:

```
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-configuration PUBLIC
    "-//Hibernate/Hibernate Configuration DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-configuration-3.0.dtd">

<hibernate-configuration>

    <session-factory>

        <property name="connection.datasource">java:comp/env/jdbc/quickstart</property>
        <property name="show_sql">false</property>
        <property name="dialect">org.hibernate.dialect.PostgreSQLDialect</property>

        <!-- Mapping files -->
        <mapping resource="Cat.hbm.xml"/>

    </session-factory>

</hibernate-configuration>
```

We turn logging of SQL commands off and tell Hibernate what database SQL dialect is used and where to get the JDBC connections (by declaring the JNDI address of the Tomcat bound pool). The dialect is a required setting, databases differ in their interpretation of the SQL "standard". Hibernate will take care of the differences and comes bundled with dialects for all major commercial and open source databases.

A `SessionFactory` is Hibernate's concept of a single datastore, multiple databases can be used by creating multiple XML configuration files and creating multiple `Configuration` and `SessionFactory` objects in your application.

The last element of the `hibernate.cfg.xml` declares `Cat.hbm.xml` as the name of a Hibernate XML mapping file for the persistent class `Cat`. This file contains the metadata for the mapping of the POJO class `Cat` to a database table (or tables). We'll come back to that file soon. Let's write the POJO class first and then declare the mapping metadata for it.

## 1.2. First persistent class

Hibernate works best with the Plain Old Java Objects (POJOs, sometimes called Plain Ordinary Java Objects) programming model for persistent classes. A POJO is much like a JavaBean, with properties of the class accessible via getter and setter methods, shielding the internal representation from the publicly visible interface (Hibernate can also access fields directly, if needed):

```
package org.hibernate.examples.quickstart;
```

```
public class Cat {  
  
    private String id;  
    private String name;  
    private char sex;  
    private float weight;  
  
    public Cat() {  
    }  
  
    public String getId() {  
        return id;  
    }  
  
    private void setId(String id) {  
        this.id = id;  
    }  
  
    public String getName() {  
        return name;  
    }  
  
    public void setName(String name) {  
        this.name = name;  
    }  
  
    public char getSex() {  
        return sex;  
    }  
  
    public void setSex(char sex) {  
        this.sex = sex;  
    }  
  
    public float getWeight() {  
        return weight;  
    }  
  
    public void setWeight(float weight) {  
        this.weight = weight;  
    }  
  
}
```

Hibernate is not restricted in its usage of property types, all Java JDK types and primitives (like `String`, `char` and `Date`) can be mapped, including classes from the Java collections framework. You can map them as values, collections of values, or associations to other entities. The `id` is a special property that represents the database identifier (primary key) of that class, it is highly recommended for entities like a `Cat`. Hibernate can use identifiers only internally, but we would lose some of the flexibility in our application architecture.

No special interface has to be implemented for persistent classes nor do you have to subclass from a special root persistent class. Hibernate also doesn't require any build time processing, such as byte-code manipulation, it relies solely on Java reflection and runtime class enhancement (through CGLIB). So, without any dependency of the POJO class on Hibernate, we can map it to a database table.

## 1.3. Mapping the cat

The `Cat.hbm.xml` mapping file contains the metadata required for the object/relational mapping. The metadata includes declaration of persistent classes and the mapping of properties (to columns and foreign key relationships to other entities) to database tables.

```
<?xml version="1.0"?>
```

```

<!DOCTYPE hibernate-mapping PUBLIC
    "-//Hibernate/Hibernate Mapping DTD 3.0//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">

<hibernate-mapping>

    <class name="org.hibernate.examples.quickstart.Cat" table="CAT">

        <!-- A 32 hex character is our surrogate key. It's automatically
            generated by Hibernate with the UUID pattern. -->
        <id name="id" type="string" unsaved-value="null" >
            <column name="CAT_ID" sql-type="char(32)" not-null="true"/>
            <generator class="uuid.hex"/>
        </id>

        <!-- A cat has to have a name, but it shouldn't be too long. -->
        <property name="name">
            <column name="NAME" length="16" not-null="true"/>
        </property>

        <property name="sex"/>

        <property name="weight"/>

    </class>

</hibernate-mapping>

```

Every persistent class should have an identifier attribute (actually, only classes representing entities, not dependent value-typed classes, which are mapped as components of an entity). This property is used to distinguish persistent objects: Two cats are equal if `catA.getId().equals(catB.getId())` is true, this concept is called *database identity*. Hibernate comes bundled with various identifier generators for different scenarios (including native generators for database sequences, hi/lo identifier tables, and application assigned identifiers). We use the UUID generator (only recommended for testing, as integer surrogate keys generated by the database should be preferred) and also specify the column `CAT_ID` of the table `CAT` for the Hibernate generated identifier value (as a primary key of the table).

All other properties of `Cat` are mapped to the same table. In the case of the `name` property, we mapped it with an explicit database column declaration. This is especially useful when the database schema is automatically generated (as SQL DDL statements) from the mapping declaration with Hibernate's *SchemaExport* tool. All other properties are mapped using Hibernate's default settings, which is what you need most of the time. The table `CAT` in the database looks like this:

Column	Type	Modifiers
cat_id	character(32)	not null
name	character varying(16)	not null
sex	character(1)	
weight	real	

Indexes: cat\_pkey primary key btree (cat\_id)

You should now create this table in your database manually, and later read Chapter 20, *Toolset Guide* if you want to automate this step with the `hbm2ddl` tool. This tool can create a full SQL DDL, including table definition, custom column type constraints, unique constraints and indexes.

## 1.4. Playing with cats

We're now ready to start Hibernate's *Session*. It is the *persistence manager*, we use it to store and retrieve `Cats` to and from the database. But first, we've to get a *Session* (Hibernate's unit-of-work) from the *SessionFactory*:

```
SessionFactory sessionFactory =
    new Configuration().configure().buildSessionFactory();
```

The call to `configure()` loads the `hibernate.cfg.xml` configuration file and initializes the `Configuration` instance. You can set other properties (and even change the mapping metadata) by accessing the `Configuration` *before* you build the `SessionFactory` (it is immutable). Where do we create the `SessionFactory` and how can we access it in our application?

A `SessionFactory` is usually only build once, e.g. at startup with a *load-on-startup* servlet. This also means you should not keep it in an instance variable in your servlets, but in some other location. Furthermore, we need some kind of *Singleton*, so we can access the `SessionFactory` easily in application code. The approach shown next solves both problems: startup configuration and easy access to a `SessionFactory`.

We implement a `HibernateUtil` helper class:

```
import org.hibernate.*;
import org.hibernate.cfg.*;

public class HibernateUtil {

    private static Log log = LogFactory.getLog(HibernateUtil.class);

    private static final SessionFactory sessionFactory;

    static {
        try {
            // Create the SessionFactory
            sessionFactory = new Configuration().configure().buildSessionFactory();
        } catch (Throwable ex) {
            // Make sure you log the exception, as it might be swallowed
            log.error("Initial SessionFactory creation failed.", ex);
            throw new ExceptionInInitializerError(ex);
        }
    }

    public static final ThreadLocal session = new ThreadLocal();

    public static Session currentSession() {
        Session s = (Session) session.get();
        // Open a new Session, if this Thread has none yet
        if (s == null) {
            s = sessionFactory.openSession();
            session.set(s);
        }
        return s;
    }

    public static void closeSession() {
        Session s = (Session) session.get();
        if (s != null)
            s.close();
        session.set(null);
    }
}
```

This class does not only take care of the `SessionFactory` with its static initializer, but also has a `ThreadLocal` variable which holds the `Session` for the current thread. Make sure you understand the Java concept of a thread-local variable before you try to use this helper. A more complex and powerful `HibernateUtil` class can be found in `CaveatEmptor`, <http://caveatemptor.hibernate.org/>

A `SessionFactory` is threadsafe, many threads can access it concurrently and request `Sessions`. A `Session` is a non-threadsafe object that represents a single unit-of-work with the database. `Sessions` are opened from a `SessionFactory` and are closed when all work is completed. An example in your servlet's `process()` method

might look like this (sans exception handling):

```
Session session = HibernateUtil.currentSession();
Transaction tx = session.beginTransaction();

Cat princess = new Cat();
princess.setName("Princess");
princess.setSex('F');
princess.setWeight(7.4f);

session.save(princess);

tx.commit();
HibernateUtil.closeSession();
```

In a `Session`, every database operation occurs inside a transaction that isolates the database operations (even read-only operations). We use `Hibernate's Transaction API` to abstract from the underlying transaction strategy (in our case, JDBC transactions). This allows our code to be deployed with container-managed transactions (using JTA) without any changes.

Note that you may call `HibernateUtil.currentSession();` as many times as you like, you will always get the current `Session` of this thread. You have to make sure the `Session` is closed after your unit-of-work completes, either in your servlet code or in a servlet filter before the HTTP response is send. The nice side effect of the second option is easy lazy initialization: the `Session` is still open when the view is rendered, so Hibernate can load uninitialized objects while you navigate the current object graph.

Hibernate has various methods that can be used to retrieve objects from the database. The most flexible way is using the `Hibernate Query Language (HQL)`, which is an easy to learn and powerful object-oriented extension to SQL:

```
Transaction tx = session.beginTransaction();

Query query = session.createQuery("select c from Cat as c where c.sex = :sex");
query.setCharacter("sex", 'F');
for (Iterator it = query.iterate(); it.hasNext();) {
    Cat cat = (Cat) it.next();
    out.println("Female Cat: " + cat.getName() );
}

tx.commit();
```

Hibernate also offers an object-oriented *query by criteria* API that can be used to formulate type-safe queries. Hibernate of course uses `PreparedStatement`s and parameter binding for all SQL communication with the database. You may also use Hibernate's direct SQL query feature or get a plain JDBC connection from a `Session` in rare cases.

## 1.5. Finally

We only scratched the surface of Hibernate in this small tutorial. Please note that we don't include any servlet specific code in our examples. You have to create a servlet yourself and insert the Hibernate code as you see fit.

Keep in mind that Hibernate, as a data access layer, is tightly integrated into your application. Usually, all other layers depend on the persistence mechanism. Make sure you understand the implications of this design.

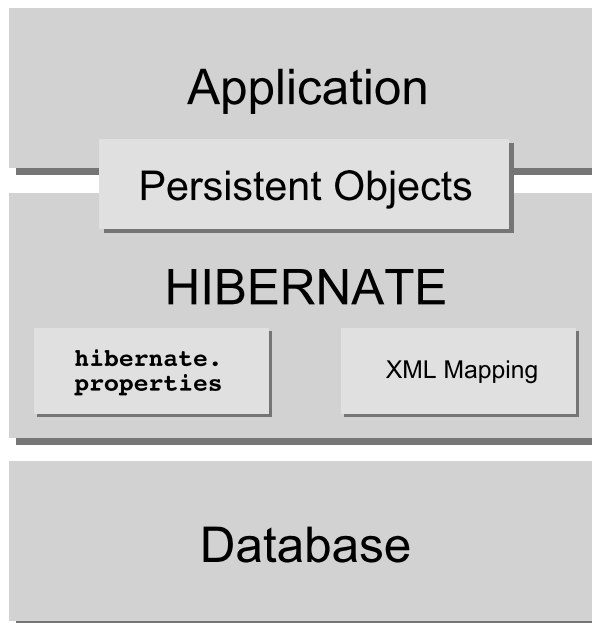
For a more complex application example, see <http://caveatemptor.hibernate.org/> and have a look at other tutorials linked on <http://www.hibernate.org/Documentation>

---

## Chapter 2. Architecture

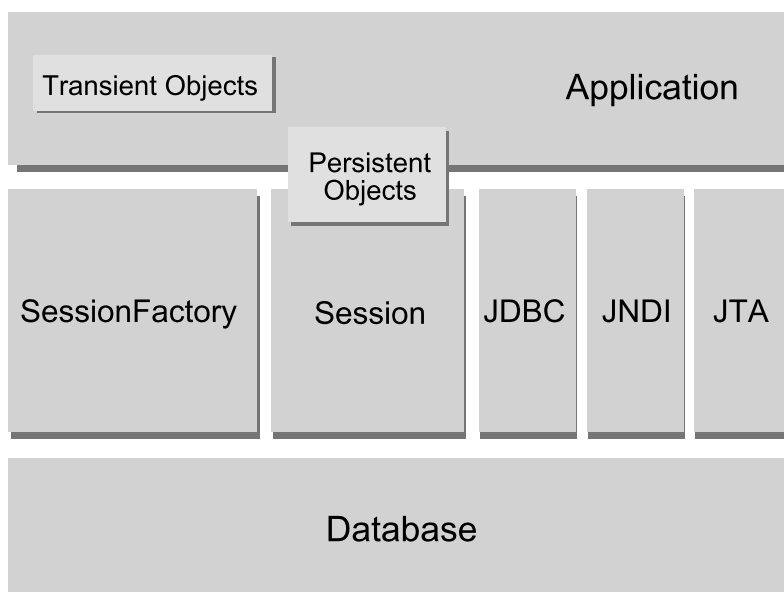
### 2.1. Overview

A (very) high-level view of the Hibernate architecture:



This diagram shows Hibernate using the database and configuration data to provide persistence services (and persistent objects) to the application.

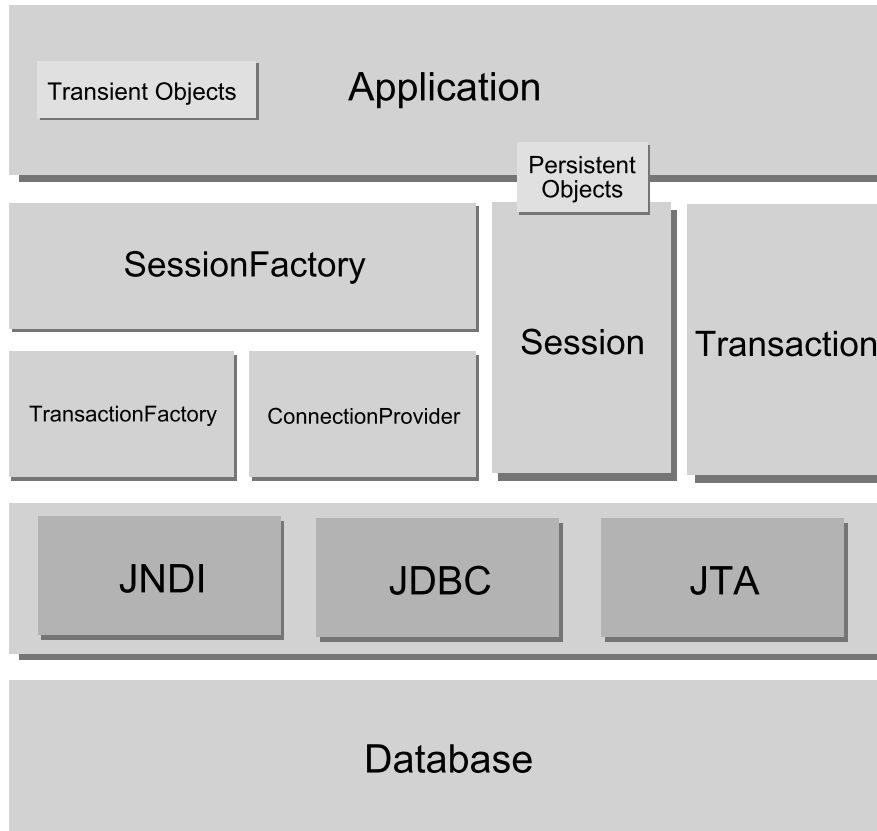
We would like to show a more detailed view of the runtime architecture. Unfortunately, Hibernate is flexible and supports several approaches. We will show the two extremes. The "lite" architecture has the application provide its own JDBC connections and manage its own transactions. This approach uses a minimal subset of Hibernate's APIs:



The "full cream" architecture abstracts the application away from the underlying JDBC/JTA APIs and lets Hi-



bernate take care of the details.



Heres some definitions of the objects in the diagrams:

#### SessionFactory (`org.hibernate.SessionFactory`)

A threadsafe (immutable) cache of compiled mappings for a single database. A factory for `Session` and a client of `ConnectionProvider`. Might hold an optional (second-level) cache of data that is reusable between transactions, at a process- or cluster-level.

#### Session (`org.hibernate.Session`)

A single-threaded, short-lived object representing a conversation between the application and the persistent store. Wraps a JDBC connection. Factory for `Transaction`. Holds a mandatory (first-level) cache of persistent objects, used when navigating the object graph or looking up objects by identifier.

#### Persistent objects and collections

Short-lived, single threaded objects containing persistent state and business function. These might be ordinary JavaBeans/POJOs, the only special thing about them is that they are currently associated with (exactly one) `Session`. As soon as the `Session` is closed, they will be detached and free to use in any application layer (e.g. directly as data transfer objects to and from presentation).

#### Transient and detached objects and collections

Instances of persistent classes that are not currently associated with a `Session`. They may have been instantiated by the application and not (yet) persisted or they may have been instantiated by a closed `Session`.

#### Transaction (`org.hibernate.Transaction`)

(Optional) A single-threaded, short-lived object used by the application to specify atomic units of work. Abstracts application from underlying JDBC, JTA or CORBA transaction. A `Session` might span several `Transactions` in some cases. However, transaction demarcation, either using the underlying API or `Transaction`, is never optional!

`ConnectionProvider` (`org.hibernate.connection.ConnectionProvider`)

(Optional) A factory for (and pool of) JDBC connections. Abstracts application from underlying `DataSource` or `DriverManager`. Not exposed to application, but can be extended/implemented by the developer.

`TransactionFactory` (`org.hibernate.TransactionFactory`)

(Optional) A factory for `Transaction` instances. Not exposed to the application, but can be extended/implemented by the developer.

### *Extension Interfaces*

Hibernate offers many optional extension interfaces you can implement to customize the behavior of your persistence layer. See the API documentation for details.

Given a "lite" architecture, the application bypasses the `Transaction/TransactionFactory` and/or `ConnectionProvider` APIs to talk to JTA or JDBC directly.

## 2.2. Instance states

An instance of a persistent classes may be in one of three different states, which are defined with respect to a *persistence context*. The Hibernate `Session` object is the persistence context:

### *transient*

The instance is not, and has never been associated with any persistence context. It has no persistent identity (primary key value).

### *persistent*

The instance is currently associated with a persistence context. It has a persistent identity (primary key value) and, perhaps, a corresponding row in the database. For a particular persistence context, Hibernate *guarantees* that persistent identity is equivalent to Java identity (in-memory location of the object).

### *detached*

The instance was once associated with a persistence context, but that context was closed, or the instance was serialized to another process. It has a persistent identity and, perhaps, a corresponding row in the database. For detached instances, Hibernate makes no guarantees about the relationship between persistent identity and Java identity.

## 2.3. JMX Integration

JMX is the J2EE standard for management of Java components. Hibernate may be managed via a JMX standard service. We provide an MBean implementation in the distribution, `org.hibernate.jmx.HibernateService`.

For an example how to deploy Hibernate as a JMX service on the JBoss Application Server, please see the JBoss User Guide. On JBoss AS, you also get these benefits if you deploy using JMX:

- *Session Management:* The Hibernate `Session`'s lifecycle can be automatically bound to the scope of a JTA transaction. This means you no longer have to manually open and close the `Session`, this becomes the job of a JBoss EJB interceptor. You also don't have to worry about transaction demarcation in your code anymore (unless you'd like to write a portable persistence layer of course, use the optional Hibernate `Transaction` API for this). You call the `HibernateContext` to access a `Session`.

- *HAR deployment:* Usually you deploy the Hibernate JMX service using a JBoss service deployment descriptor (in an EAR and/or SAR file), it supports all the usual configuration options of a Hibernate `SessionFactory`. However, you still have to name all your mapping files in the deployment descriptor. If you decide to use the optional HAR deployment, JBoss will automatically detect all mapping files in your HAR file.

Consult the JBoss AS user guide for more information about these options.

Another feature available as a JMX service are runtime Hibernate statistics. See Section 3.4.8, “Hibernate statistics”.

## 2.4. JCA Support

Hibernate may also be configured as a JCA connector. Please see the website for more details. Please note that Hibernate JCA support is still considered experimental.

---

## Chapter 3. Configuration

Because Hibernate is designed to operate in many different environments, there are a large number of configuration parameters. Fortunately, most have sensible default values and Hibernate is distributed with an example `hibernate.properties` file in `etc/` that shows the various options. Just put the example file in your classpath and customize it.

### 3.1. Programmatic configuration

An instance of `org.hibernate.cfg.Configuration` represents an entire set of mappings of an application's Java types to an SQL database. The `Configuration` is used to build an (immutable) `SessionFactory`. The mappings are compiled from various XML mapping files.

You may obtain a `Configuration` instance by instantiating it directly and specifying XML mapping documents. If the mapping files are in the classpath, use `addResource()`:

```
Configuration cfg = new Configuration()
    .addResource("Item.hbm.xml")
    .addResource("Bid.hbm.xml");
```

An alternative (sometimes better) way is to specify the mapped class, and let Hibernate find the mapping document for you:

```
Configuration cfg = new Configuration()
    .addClass(org.hibernate.auction.Item.class)
    .addClass(org.hibernate.auction.Bid.class);
```

Then Hibernate will look for mapping files named `/org/hibernate/auction/Item.hbm.xml` and `/org/hibernate/auction/Bid.hbm.xml` in the classpath. This approach eliminates any hardcoded filenames.

A `Configuration` also allows you to specify configuration properties:

```
Configuration cfg = new Configuration()
    .addClass(org.hibernate.auction.Item.class)
    .addClass(org.hibernate.auction.Bid.class)
    .setProperty("hibernate.dialect", "org.hibernate.dialect.MySQLInnoDBDialect")
    .setProperty("hibernate.connection.datasource", "java:comp/env/jdbc/test")
    .setProperty("hibernate.order_updates", "true");
```

This is not the only way to pass configuration properties to Hibernate. The various options include:

1. Pass an instance of `java.util.Properties` to `Configuration.setProperties()`.
2. Place `hibernate.properties` in a root directory of the classpath.
3. Set system properties using `java -Dproperty=value`.
4. Include `<property>` elements in `hibernate.cfg.xml` (discussed later).

`hibernate.properties` is the easiest approach if you want to get started quickly.

The `Configuration` is intended as a startup-time object, to be discarded once a `SessionFactory` is created.

### 3.2. Obtaining a SessionFactory

When all mappings have been parsed by the `Configuration`, the application must obtain a factory for `Session`

instances. This factory is intended to be shared by all application threads:

```
SessionFactory sessions = cfg.buildSessionFactory();
```

Hibernate does allow your application to instantiate more than one `SessionFactory`. This is useful if you are using more than one database.

### 3.3. JDBC connections

Usually, you want to have the `SessionFactory` create and pool JDBC connections for you. If you take this approach, opening a `Session` is as simple as:

```
Session session = sessions.openSession(); // open a new Session
```

As soon as you do something that requires access to the database, a JDBC connection will be obtained from the pool.

For this to work, we need to pass some JDBC connection properties to Hibernate. All Hibernate property names and semantics are defined on the class `org.hibernate.cfg.Environment`. We will now describe the most important settings for JDBC connection configuration.

Hibernate will obtain (and pool) connections using `java.sql.DriverManager` if you set the following properties:

**Table 3.1. Hibernate JDBC Properties**

Property name	Purpose
<code>hibernate.connection.driver_class</code>	<i>jdbc driver class</i>
<code>hibernate.connection.url</code>	<i>jdbc URL</i>
<code>hibernate.connection.username</code>	<i>database user</i>
<code>hibernate.connection.password</code>	<i>database user password</i>
<code>hibernate.connection.pool_size</code>	<i>maximum number of pooled connections</i>

Hibernate's own connection pooling algorithm is however quite rudimentary. It is intended to help you get started and is *not intended for use in a production system* or even for performance testing. You should use a third party pool for best performance and stability. Just replace the `hibernate.connection.pool_size` property with connection pool specific settings. This will turn off Hibernate's internal pool. For example, you might like to use C3P0.

C3P0 is an open source JDBC connection pool distributed along with Hibernate in the `lib` directory. Hibernate will use its `C3P0ConnectionProvider` for connection pooling if you set `hibernate.c3p0.*` properties. If you'd like to use Proxool refer to the packaged `hibernate.properties` and the Hibernate web site for more information.

Here is an example `hibernate.properties` file for C3P0:

```
hibernate.connection.driver_class = org.postgresql.Driver
hibernate.connection.url = jdbc:postgresql://localhost/mydatabase
hibernate.connection.username = myuser
hibernate.connection.password = secret
```

```
hibernate.c3p0.min_size=5
hibernate.c3p0.max_size=20
hibernate.c3p0.timeout=1800
hibernate.c3p0.max_statements=50
hibernate.dialect = org.hibernate.dialect.PostgreSQLDialect
```

For use inside an application server, you should almost always configure Hibernate to obtain connections from an application server `Datasource` registered in JNDI. You'll need to set at least one of the following properties:

**Table 3.2. Hibernate Datasource Properties**

Property name	Purpose
<code>hibernate.connection.datasource</code>	<i>datasource JNDI name</i>
<code>hibernate.jndi.url</code>	<i>URL of the JNDI provider (optional)</i>
<code>hibernate.jndi.class</code>	<i>class of the JNDI InitialContextFactory (optional)</i>
<code>hibernate.connection.username</code>	<i>database user (optional)</i>
<code>hibernate.connection.password</code>	<i>database user password (optional)</i>

Here's an example `hibernate.properties` file for an application server provided JNDI datasource:

```
hibernate.connection.datasource = java:/comp/env/jdbc/test
hibernate.transaction.factory_class = \
    org.hibernate.transaction.JTATransactionFactory
hibernate.transaction.manager_lookup_class = \
    org.hibernate.transaction.JBossTransactionManagerLookup
hibernate.dialect = org.hibernate.dialect.PostgreSQLDialect
```

JDBC connections obtained from a JNDI datasource will automatically participate in the container-managed transactions of the application server.

Arbitrary connection properties may be given by prepending "hibernate.connection" to the property name. For example, you may specify a `charSet` using `hibernate.connection.charSet`.

You may define your own plugin strategy for obtaining JDBC connections by implementing the interface `org.hibernate.connection.ConnectionProvider`. You may select a custom implementation by setting `hibernate.connection.provider_class`.

## 3.4. Optional configuration properties

There are a number of other properties that control the behaviour of Hibernate at runtime. All are optional and have reasonable default values.

*Warning: some of these properties are "system-level" only.* System-level properties can be set only via `java -Dproperty=value` or `hibernate.properties`. They may *not* be set by the other techniques described above.

**Table 3.3. Hibernate Configuration Properties**

Property name	Purpose
<code>hibernate.dialect</code>	The classname of a <code>Hibernate Dialect</code> which allows

Property name	Purpose
	<p>Hibernate to generate SQL optimized for a particular relational database.</p> <p><i>eg. full.classname.of.Dialect</i></p>
<code>hibernate.show_sql</code>	<p>Write all SQL statements to console.</p> <p><i>eg. true   false</i></p>
<code>hibernate.default_schema</code>	<p>Qualify unqualified tablenamees with the given schema/tablespace in generated SQL.</p> <p><i>eg. SCHEMA_NAME</i></p>
<code>hibernate.default_catalog</code>	<p>Qualify unqualified tablenamees with the given catalog in generated SQL.</p> <p><i>eg. CATALOG_NAME</i></p>
<code>hibernate.session_factory_name</code>	<p>The <code>SessionFactory</code> will be automatically bound to this name in JNDI after it has been created.</p> <p><i>eg. jndi/composite/name</i></p>
<code>hibernate.max_fetch_depth</code>	<p>Set a maximum "depth" for the outer join fetch tree for single-ended associations (one-to-one, many-to-one). A 0 disables default outer join fetching.</p> <p><i>eg. recommended values between 0 and 3</i></p>
<code>hibernate.default_batch_fetch_size</code>	<p>Set a default size for Hibernate batch fetching of associations.</p> <p><i>eg. recommended values 4, 8, 16</i></p>
<code>hibernate.default_entity_mode</code>	<p>Set a default mode for entity representation for all sessions opened from this <code>SessionFactory</code></p> <p><i>dynamic-map, dom4j, pojo</i></p>
<code>hibernate.order_updates</code>	<p>Force Hibernate to order SQL updates by the primary key value of the items being updates. This will result in fewer transaction deadlocks in highly concurrent systems.</p> <p><i>eg. true   false</i></p>
<code>hibernate.generate_statistics</code>	<p>If enabled, Hibernate will collect statistics useful for performance tuning.</p> <p><i>eg. true   false</i></p>
<code>hibernate.use_identifer_rollback</code>	<p>If enabled, generated identifier properties will be reset to default values when objects are deleted.</p> <p><i>eg. true   false</i></p>
<code>hibernate.use_sql_comments</code>	<p>If turned on, Hibernate will generate comments inside</p>

Property name	Purpose
	the SQL, for easier debugging, defaults to <code>false</code> .  <i>eg.</i> <code>true</code>   <code>false</code>

**Table 3.4. Hibernate JDBC and Connection Properties**

Property name	Purpose
<code>hibernate.jdbc.fetch_size</code>	A non-zero value determines the JDBC fetch size (calls <code>Statement.setFetchSize()</code> ).
<code>hibernate.jdbc.batch_size</code>	A non-zero value enables use of JDBC2 batch updates by Hibernate.  <i>eg.</i> recommended values between 5 and 30
<code>hibernate.jdbc.batch_versioned_data</code>	Set this property to <code>true</code> if your JDBC driver returns correct row counts from <code>executeBatch()</code> (it is usually safe to turn this option on). Hibernate will then use batched DML for automatically versioned data. Defaults to <code>false</code> .  <i>eg.</i> <code>true</code>   <code>false</code>
<code>hibernate.jdbc.factory_class</code>	Select a custom <code>Batcher</code> . Most applications will not need this configuration property.  <i>eg.</i> <code>classname.of.Batcher</code>
<code>hibernate.jdbc.use_scrollable_resultset</code>	Enables use of JDBC2 scrollable resultsets by Hibernate. This property is only necessary when using user supplied JDBC connections, Hibernate uses connection metadata otherwise.  <i>eg.</i> <code>true</code>   <code>false</code>
<code>hibernate.jdbc.use_streams_for_binary</code>	Use streams when writing/reading binary or serializable types to/from JDBC (system-level property).  <i>eg.</i> <code>true</code>   <code>false</code>
<code>hibernate.jdbc.use_get_generated_keys</code>	Enable use of JDBC3 <code>PreparedStatement.getGeneratedKeys()</code> to retrieve natively generated keys after insert. Requires JDBC3+ driver and JRE1.4+, set to <code>false</code> if your driver has problems with the Hibernate identifier generators. By default, tries to determine the driver capabilities using connection metadata.  <i>eg.</i> <code>true</code>   <code>false</code>
<code>hibernate.connection.provider_class</code>	The classname of a custom <code>ConnectionProvider</code> which provides JDBC connections to Hibernate.  <i>eg.</i> <code>classname.of.ConnectionProvider</code>



Property name	Purpose
<code>hibernate.connection.isolation</code>	Set the JDBC transaction isolation level. Check <code>java.sql.Connection</code> for meaningful values but note that most databases do not support all isolation levels.  <i>eg. 1, 2, 4, 8</i>
<code>hibernate.connection.autocommit</code>	Enables autocommit for JDBC pooled connections (not recommended).  <i>eg. true   false</i>
<code>hibernate.connection.&lt;propertyName&gt;</code>	Pass the JDBC property <code>propertyName</code> to <code>DriverManager.getConnection()</code> .
<code>hibernate.jndi.&lt;propertyName&gt;</code>	Pass the property <code>propertyName</code> to the JNDI <code>InitialContextFactory</code> .

**Table 3.5. Hibernate Cache Properties**

Property name	Purpose
<code>hibernate.cache.provider_class</code>	The classname of a custom <code>CacheProvider</code> .  <i>eg. classname.of.CacheProvider</i>
<code>hibernate.cache.use_minimal_puts</code>	Optimize second-level cache operation to minimize writes, at the cost of more frequent reads. This setting is most useful for clustered caches and, in Hibernate3, is enabled by default for clustered cache implementations.  <i>eg. true   false</i>
<code>hibernate.cache.use_query_cache</code>	Enable the query cache, individual queries still have to be set cacheable.  <i>eg. true   false</i>
<code>hibernate.cache.use_second_level_cache</code>	May be used to completely disable the second level cache, which is enabled by default for classes which specify a <code>&lt;cache&gt;</code> mapping.  <i>eg. true   false</i>
<code>hibernate.cache.query_cache_factory</code>	The classname of a custom <code>QueryCache</code> interface, defaults to the built-in <code>StandardQueryCache</code> .  <i>eg. classname.of.QueryCache</i>
<code>hibernate.cache.region_prefix</code>	A prefix to use for second-level cache region names.  <i>eg. prefix</i>
<code>hibernate.cache.use_structured_entries</code>	Forces Hibernate to store data in the second-level cache in a more human-friendly format.

Property name	Purpose
	<i>eg. true false</i>

**Table 3.6. Hibernate Transaction Properties**

Property name	Purpose
<code>hibernate.transaction.factory_class</code>	The classname of a <code>TransactionFactory</code> to use with Hibernate Transaction API (defaults to <code>JDBCTransactionFactory</code> ).  <i>eg. classname.of.TransactionFactory</i>
<code>jta.UserTransaction</code>	A JNDI name used by <code>JTATransactionFactory</code> to obtain the JTA <code>UserTransaction</code> from the application server.  <i>eg. jndi/composite/name</i>
<code>hibernate.transaction.manager_lookup_class</code>	The classname of a <code>TransactionManagerLookup</code> - required when JVM-level caching is enabled or when using hilo generator in a JTA environment.  <i>eg. classname.of.TransactionManagerLookup</i>
<code>hibernate.transaction.flush_before_completion</code>	If enabled, the session will be automatically flushed during the before completion phase of the transaction. (Very useful when using Hibernate with CMT.)  <i>eg. true   false</i>
<code>hibernate.transaction.auto_close_session</code>	If enabled, the session will be automatically closed during the before completion phase of the transaction. (Very useful when using Hibernate with CMT.)  <i>eg. true   false</i>

**Table 3.7. Miscellaneous Properties**

Property name	Purpose
<code>hibernate.query.factory_class</code>	Chooses the HQL parser implementation.  <i>eg.</i> <code>org.hibernate.hql.ast.ASTQueryTranslatorFactory</code> or <code>org.hibernate.hql.classic.ClassicQueryTranslatorFactory</code>
<code>hibernate.query.substitutions</code>	Mapping from tokens in Hibernate queries to SQL tokens (tokens might be function or literal names, for example).

Property name	Purpose
	<i>eg.</i> <code>hqlLiteral=SQL_LITERAL, hqlFunction=SQLFUNC</code>
<code>hibernate.hbm2ddl.auto</code>	Automatically export schema DDL to the database when the <code>SessionFactory</code> is created. With <code>create-drop</code> , the database schema will be dropped when the <code>SessionFactory</code> is closed explicitly.  <i>eg.</i> <code>update   create   create-drop</code>
<code>hibernate.cglib.use_reflection_optimizer</code>	Enables use of CGLIB instead of runtime reflection (System-level property). Reflection can sometimes be useful when troubleshooting, note that Hibernate always requires CGLIB even if you turn off the optimizer. You can not set this property in <code>hibernate.cfg.xml</code> .  <i>eg.</i> <code>true   false</code>

### 3.4.1. SQL Dialects

You should always set the `hibernate.dialect` property to the correct `org.hibernate.dialect.Dialect` subclass for your database. If you specify a dialect, Hibernate will use sensible defaults for some of the other properties listed above, saving you the effort of specifying them manually.

**Table 3.8. Hibernate SQL Dialects (`hibernate.dialect`)**

RDBMS	Dialect
DB2	<code>org.hibernate.dialect.DB2Dialect</code>
DB2 AS/400	<code>org.hibernate.dialect.DB2400Dialect</code>
DB2 OS390	<code>org.hibernate.dialect.DB2390Dialect</code>
PostgreSQL	<code>org.hibernate.dialect.PostgreSQLDialect</code>
MySQL	<code>org.hibernate.dialect.MySQLDialect</code>
MySQL with InnoDB	<code>org.hibernate.dialect.MySQLInnoDBDialect</code>
MySQL with MyISAM	<code>org.hibernate.dialect.MySQLMyISAMDialect</code>
Oracle (any version)	<code>org.hibernate.dialect.OracleDialect</code>
Oracle 9i/10g	<code>org.hibernate.dialect.Oracle9Dialect</code>
Sybase	<code>org.hibernate.dialect.SybaseDialect</code>
Sybase Anywhere	<code>org.hibernate.dialect.SybaseAnywhereDialect</code>
Microsoft SQL Server	<code>org.hibernate.dialect.SQLServerDialect</code>
SAP DB	<code>org.hibernate.dialect.SAPDBDialect</code>
Informix	<code>org.hibernate.dialect.InformixDialect</code>

RDBMS	Dialect
HypersonicSQL	<code>org.hibernate.dialect.HSQLDialect</code>
Ingres	<code>org.hibernate.dialect.IngresDialect</code>
Progress	<code>org.hibernate.dialect.ProgressDialect</code>
Mckoi SQL	<code>org.hibernate.dialect.MckoiDialect</code>
Interbase	<code>org.hibernate.dialect.InterbaseDialect</code>
Pointbase	<code>org.hibernate.dialect.PointbaseDialect</code>
FrontBase	<code>org.hibernate.dialect.FrontbaseDialect</code>
Firebird	<code>org.hibernate.dialect.FirebirdDialect</code>

### 3.4.2. Outer Join Fetching

If your database supports ANSI, Oracle or Sybase style outer joins, *outer join fetching* will often increase performance by limiting the number of round trips to and from the database (at the cost of possibly more work performed by the database itself). Outer join fetching allows a whole graph of objects connected by many-to-one, one-to-many, many-to-many and one-to-one associations to be retrieved in a single SQL `SELECT`.

Outer join fetching may be disabled *globally* by setting the property `hibernate.max_fetch_depth` to 0. A setting of 1 or higher enables outer join fetching for one-to-one and many-to-one associations which have been mapped with `fetch="join"`.

See Section 19.1, “Fetching strategies” for more information.

### 3.4.3. Binary Streams

Oracle limits the size of `byte` arrays that may be passed to/from its JDBC driver. If you wish to use large instances of `binary` or `serializable` type, you should enable `hibernate.jdbc.use_streams_for_binary`. *This is a system-level setting only.*

### 3.4.4. Second-level and query cache

The properties prefixed by `hibernate.cache` allow you to use a process or cluster scoped second-level cache system with Hibernate. See the Section 19.2, “The Second Level Cache” for more details.

### 3.4.5. Transaction strategy configuration

If you wish to use the Hibernate `Transaction` API instead of directly calling a particular system transaction API, you must specify a factory class for `Transaction` instances by setting the property `hibernate.transaction.factory_class`. The `Transaction` API hides the underlying transaction mechanism and allows Hibernate code to run in managed and non-managed environments.

There are two standard (built-in) choices:

```
org.hibernate.transaction.JDBCTransactionFactory
    delegates to database (JDBC) transactions (default)
```

`org.hibernate.transaction.JTATransactionFactory`

delegates to JTA (if an existing transaction is underway, the `Session` performs its work in that context, otherwise a new transaction is started)

You may also define your own transaction strategies (for a CORBA transaction service, for example).

Some features in Hibernate (i.e. the second level cache) require access to the JTA `TransactionManager` in a management environment. You have to specify how Hibernate should obtain a reference to the `TransactionManager`, since J2EE does not standardize a single mechanism:

**Table 3.9. JTA TransactionManagers**

Transaction Factory	Application Server
<code>org.hibernate.transaction.JBossTransactionManagerLookup</code>	JBoss
<code>org.hibernate.transaction.WeblogicTransactionManagerLookup</code>	Weblogic
<code>org.hibernate.transaction.WebSphereTransactionManagerLookup</code>	WebSphere
<code>org.hibernate.transaction.OrionTransactionManagerLookup</code>	Orion
<code>org.hibernate.transaction.ResinTransactionManagerLookup</code>	Resin
<code>org.hibernate.transaction.JOTMTransactionManagerLookup</code>	JOTM
<code>org.hibernate.transaction.JOnASTransactionManagerLookup</code>	JOnAS
<code>org.hibernate.transaction.JRun4TransactionManagerLookup</code>	JRun4
<code>org.hibernate.transaction.BESTransactionManagerLookup</code>	Borland ES

### 3.4.6. JNDI-bound `SessionFactory`

A JNDI bound Hibernate `SessionFactory` can simplify the lookup of the factory and the creation of new `Sessions`. Note that this is not related to a JNDI bound `Datasource` in a managed environment.

If you wish to have the `SessionFactory` bound to a JNDI namespace, specify a name (eg. `java:hibernate/SessionFactory`) using the property `hibernate.session_factory_name`. If this property is omitted, the `SessionFactory` will not be bound to JNDI. (This is especially useful in environments with a read-only JNDI default implementation, eg. Tomcat.)

When binding the `SessionFactory` to JNDI, Hibernate will use the values of `hibernate.jndi.url`, `hibernate.jndi.class` to instantiate an initial context. If they are not specified, the default `InitialContext` will be used.

Hibernate will automatically place the `SessionFactory` in JNDI after you call `cfg.buildSessionFactory()`. This means you will at least have this call in some startup code (or utility class) in your application, unless you use JMX deployment with the `HibernateService`.

If you use a JNDI `SessionFactory`, an EJB or any other class may obtain the `SessionFactory` using a JNDI lookup. Note that this setup is not necessary if you use the `HibernateUtil` helper class introduced in chapter 1, which acts as a Singleton registry.

### 3.4.7. Query Language Substitution

You may define new Hibernate query tokens using `hibernate.query.substitutions`. For example:

```
hibernate.query.substitutions true=1, false=0
```

would cause the tokens `true` and `false` to be translated to integer literals in the generated SQL.

```
hibernate.query.substitutions toLowercase=LOWER
```

would allow you to rename the SQL `LOWER` function.

### 3.4.8. Hibernate statistics

If you enable `hibernate.generate_statistics`, Hibernate will expose a number of metrics that are useful when tuning a running system via `SessionFactory.getStatistics()`. Hibernate can even be configured to expose these statistics via JMX. Read the Javadoc of the interfaces in `org.hibernate.stats` for more information.

## 3.5. Logging

Hibernate logs various events using Apache commons-logging.

The commons-logging service will direct output to either Apache Log4j (if you include `log4j.jar` in your classpath) or JDK1.4 logging (if running under JDK1.4 or above). You may download Log4j from <http://jakarta.apache.org>. To use Log4j you will need to place a `log4j.properties` file in your classpath, an example properties file is distributed with Hibernate in the `src/` directory.

We strongly recommend that you familiarize yourself with Hibernate's log messages. A lot of work has been put into making the Hibernate log as detailed as possible, without making it unreadable. It is an essential troubleshooting device. The most interesting log categories are the following:

**Table 3.10. Hibernate Log Categories**

Category	Function
<code>org.hibernate.SQL</code>	Log all SQL DML statements as they are executed
<code>org.hibernate.type</code>	Log all JDBC parameters
<code>org.hibernate.tool.hbm2ddl</code>	Log all SQL DDL statements as they are executed
<code>org.hibernate.pretty</code>	Log the state of all entities (max 20 entities) associated with the session at flush time
<code>org.hibernate.cache</code>	Log all second-level cache activity
<code>org.hibernate.transaction</code>	Log transaction related activity
<code>org.hibernate.jdbc</code>	Log all JDBC resource acquisition
<code>org.hibernate.secure</code>	Log all JAAS authorization requests
<code>org.hibernate</code>	Log everything (a lot of information, but very useful for troubleshooting)

When developing applications with Hibernate, you should almost always work with `debug` enabled for the category `org.hibernate.SQL`, or, alternatively, the property `hibernate.show_sql` enabled.

### 3.6. Implementing a `NamingStrategy`

The interface `org.hibernate.cfg.NamingStrategy` allows you to specify a "naming standard" for database objects and schema elements.

You may provide rules for automatically generating database identifiers from Java identifiers or for processing "logical" column and table names given in the mapping file into "physical" table and column names. This feature helps reduce the verbosity of the mapping document, eliminating repetitive noise (TBL\_ prefixes, for example). The default strategy used by Hibernate is quite minimal.

You may specify a different strategy by calling `Configuration.setNamingStrategy()` before adding mappings:

```
SessionFactory sf = new Configuration()
    .setNamingStrategy(ImprovedNamingStrategy.INSTANCE)
    .addFile("Item.hbm.xml")
    .addFile("Bid.hbm.xml")
    .buildSessionFactory();
```

`org.hibernate.cfg.ImprovedNamingStrategy` is a built-in strategy that might be a useful starting point for some applications.

### 3.7. XML configuration file

An alternative approach to configuration is to specify a full configuration in a file named `hibernate.cfg.xml`. This file can be used as a replacement for the `hibernate.properties` file or, if both are present, to override properties.

The XML configuration file is by default expected to be in the root of your `CLASSPATH`. Here is an example:

```
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-configuration PUBLIC
    "-//Hibernate/Hibernate Configuration DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-configuration-3.0.dtd">

<hibernate-configuration>

    <!-- a SessionFactory instance listed as /jndi/name -->
    <session-factory
        name="java:hibernate/SessionFactory">

        <!-- properties -->
        <property name="connection.datasource">java:/comp/env/jdbc/MyDB</property>
        <property name="dialect">org.hibernate.dialect.MySQLDialect</property>
        <property name="show_sql">false</property>
        <property name="transaction.factory_class">
            org.hibernate.transaction.JTATransactionFactory
        </property>
        <property name="jta.UserTransaction">java:comp/UserTransaction</property>

        <!-- mapping files -->
        <mapping resource="org/hibernate/auction/Item.hbm.xml"/>
        <mapping resource="org/hibernate/auction/Bid.hbm.xml"/>

        <!-- cache settings -->
        <class-cache class="org.hibernate.auction.Item" usage="read-write"/>
    </session-factory>
</hibernate-configuration>
```

```
<class-cache class="org.hibernate.auction.Bid" usage="read-only"/>
<collection-cache class="org.hibernate.auction.Item.bids" usage="read-write"/>

</session-factory>

</hibernate-configuration>
```

As you can see, the advantage of this approach is the externalization of the mapping file names to configuration. The `hibernate.cfg.xml` is also more convenient once you have to tune the Hibernate cache. Note that is your choice to use either `hibernate.properties` or `hibernate.cfg.xml`, both are equivalent, except for the above mentioned benefits of using the XML syntax.

With the XML configuration, starting Hibernate is then as simple as

```
SessionFactory sf = new Configuration().configure().buildSessionFactory();
```

You can pick a different XML configuration file using

```
SessionFactory sf = new Configuration()
    .configure("catdb.cfg.xml")
    .buildSessionFactory();
```



---

## Chapter 4. Persistent Classes

Persistent classes are classes in an application that implement the entities of the business problem (e.g. Customer and Order in an E-commerce application). Not all instances of a persistent class are considered to be in the persistent state - an instance may instead be transient or detached.

Hibernate works best if these classes follow some simple rules, also known as the Plain Old Java Object (POJO) programming model. However none of these rules are hard requirements. Indeed, Hibernate3 assumes very little about the nature of your persistent objects. You may express a domain model in other ways: using trees of Map instances, for example.

### 4.1. A simple POJO example

Most Java applications require a persistent class representing felines.

```
package eg;
import java.util.Set;
import java.util.Date;

public class Cat {
    private Long id; // identifier

    private Date birthdate;
    private Color color;
    private char sex;
    private float weight;
    private int litterId;

    private Cat mother;
    private Set kittens = new HashSet();

    private void setId(Long id) {
        this.id=id;
    }
    public Long getId() {
        return id;
    }

    void setBirthdate(Date date) {
        birthdate = date;
    }
    public Date getBirthdate() {
        return birthdate;
    }

    void setWeight(float weight) {
        this.weight = weight;
    }
    public float getWeight() {
        return weight;
    }

    public Color getColor() {
        return color;
    }
    void setColor(Color color) {
        this.color = color;
    }

    void setSex(char sex) {
        this.sex=sex;
    }
    public char getSex() {
        return sex;
    }
}
```

```

    }

    void setLitterId(int id) {
        this.litterId = id;
    }
    public int getLitterId() {
        return litterId;
    }

    void setMother(Cat mother) {
        this.mother = mother;
    }
    public Cat getMother() {
        return mother;
    }
    void setKittens(Set kittens) {
        this.kittens = kittens;
    }
    public Set getKittens() {
        return kittens;
    }

    // addKitten not needed by Hibernate
    public void addKitten(Cat kitten) {
        kitten.setMother(this);
        kitten.setLitterId( kittens.size() );
        kittens.add(kitten);
    }
}

```

There are four main rules to follow here:

#### 4.1.1. Declare accessors and mutators for persistent fields

`Cat` declares accessor methods for all its persistent fields. Many other ORM tools directly persist instance variables. We believe it is far better to decouple this implementation detail from the persistence mechanism. Hibernate persists JavaBeans style properties, and recognizes method names of the form `getFoo`, `isFoo` and `setFoo`. You may however switch to direct field access for particular properties, if needed.

Properties need *not* be declared public - Hibernate can persist a property with a default, `protected` or `private` `get / set` pair.

#### 4.1.2. Implement a no-argument constructor

`Cat` has a no-argument constructor. All persistent classes must have a default constructor (which may be non-public) so Hibernate can instantiate them using `Constructor.newInstance()`. We recommend having a constructor with at least *package* visibility for runtime proxy generation in Hibernate.

#### 4.1.3. Provide an identifier property (optional)

`Cat` has a property called `id`. This property maps to the primary key column of a database table. The property might have been called anything, and its type might have been any primitive type, any primitive "wrapper" type, `java.lang.String` or `java.util.Date`. (If your legacy database table has composite keys, you can even use a user-defined class with properties of these types - see the section on composite identifiers later.)

The identifier property is strictly optional. You can leave them off and let Hibernate keep track of object identifiers internally. We do not recommend this, however.

In fact, some functionality is available only to classes which declare an identifier property:

- Transitive reattachment for detached objects (cascade update or cascade merge) - see Section 10.11, “Transitive persistence”
- `Session.saveOrUpdate()`
- `Session.merge()`

We recommend you declare consistently-named identifier properties on persistent classes. We further recommend that you use a nullable (ie. non-primitive) type.

#### 4.1.4. Prefer non-final classes (optional)

A central feature of Hibernate, *proxies*, depends upon the persistent class being either non-final, or the implementation of an interface that declares all public methods.

You can persist `final` classes that do not implement an interface with Hibernate, but you won't be able to use proxies for lazy association fetching - which will limit your options for performance tuning.

## 4.2. Implementing inheritance

A subclass must also observe the first and second rules. It inherits its identifier property from the superclass, `Cat`.

```
package eg;

public class DomesticCat extends Cat {
    private String name;

    public String getName() {
        return name;
    }
    protected void setName(String name) {
        this.name=name;
    }
}
```

## 4.3. Implementing `equals()` and `hashCode()`

You have to override the `equals()` and `hashCode()` methods if you

- intend to put instances of persistent classes in a `Set` (the recommended way to represent many-valued associations) *and*
- intend to use reattachment of detached instances

Hibernate guarantees equivalence of persistent identity (database row) and Java identity only inside a particular session scope. So as soon as we mix instances retrieved in different sessions, we must implement `equals()` and `hashCode()` if we wish to have meaningful semantics for `Sets`.

The most obvious way is to implement `equals()/hashCode()` by comparing the identifier value of both objects. If the value is the same, both must be the same database row, they are therefore equal (if both are added to a `Set`, we will only have one element in the `Set`). Unfortunately, we can't use that approach with generated identifiers! Hibernate will only assign identifier values to objects that are persistent, a newly created instance will not have any identifier value! Furthermore, if an instance is unsaved and currently in a `Set`, saving it will assign

an identifier value to the object. If `equals()` and `hashCode()` are based on the identifier value, the hash code would change, breaking the contract of the `Set`. See the Hibernate website for a full discussion of this problem. Note that this is not a Hibernate issue, but normal Java semantics of object identity and equality.

We recommend implementing `equals()` and `hashCode()` using *Business key equality*. Business key equality means that the `equals()` method compares only the properties that form the business key, a key that would identify our instance in the real world (a *natural* candidate key):

```
public class Cat {

    ...
    public boolean equals(Object other) {
        if (this == other) return true;
        if ( !(other instanceof Cat) ) return false;

        final Cat cat = (Cat) other;

        if ( !cat.getLitterId().equals( getLitterId() ) ) return false;
        if ( !cat.getMother().equals( getMother() ) ) return false;

        return true;
    }

    public int hashCode() {
        int result;
        result = getMother().hashCode();
        result = 29 * result + getLitterId();
        return result;
    }

}
```

Note that a business key does not have to be as solid as a database primary key candidate (see Section 11.1.3, “Considering object identity”). Immutable or unique properties are usually good candidates for a business key.

## 4.4. Dynamic models

*Note that the following features are currently considered experimental and may change in the near future.*

Persistent entities don't necessarily have to be represented as POJO classes or as JavaBean objects at runtime. Hibernate also supports dynamic models (using `Maps` of `Maps` at runtime) and the representation of entities as DOM4J trees. With this approach, you don't write persistent classes, only mapping files.

By default, Hibernate works in normal POJO mode. You may set a default entity representation mode for a particular `SessionFactory` using the `default_entity_mode` configuration option (see Table 3.3, “Hibernate Configuration Properties”).

The following examples demonstrates the representation using `Maps`. First, in the mapping file, an `entity-name` has to be declared instead of (or in addition to) a class name:

```
<hibernate-mapping>

    <class entity-name="Customer">

        <id name="id"
            type="long"
            column="ID">
            <generator class="sequence" />
        </id>

        <property name="name"
            column="NAME"
```

```

        type="string" />

        <property name="address"
            column="ADDRESS"
            type="string" />

        <many-to-one name="organization"
            column="ORGANIZATION_ID"
            class="Organization" />

        <bag name="orders"
            inverse="true"
            lazy="false"
            cascade="all">
            <key column="CUSTOMER_ID" />
            <one-to-many class="Order" />
        </bag>

    </class>
</hibernate-mapping>

```

Note that even though associations are declared using target class names, the target type of an associations may also be a dynamic entity instead of a POJO.

After setting the default entity mode to `dynamic-map` for the `SessionFactory`, we can at runtime work with Maps Of Maps:

```

Session s = openSession();
Transaction tx = s.beginTransaction();
Session s = openSession();

// Create a customer
Map david = new HashMap();
david.put("name", "David");

// Create an organization
Map foobar = new HashMap();
foobar.put("name", "Foobar Inc.");

// Link both
david.put("organization", foobar);

// Save both
s.save("Customer", david);
s.save("Organization", foobar);

tx.commit();
s.close();

```

The advantages of a dynamic mapping are quick turnaround time for prototyping without the need for entity class implementation. However, you lose compile-time type checking and will very likely deal with many exceptions at runtime. Thanks to the Hibernate mapping, the database schema can easily be normalized and sound, allowing to add a proper domain model implementation on top later on.

Entity representation modes can also be set on a per `Session` basis:

```

Session dynamicSession = pojoSession.getSession(EntityMode.MAP);

// Create a customer
Map david = new HashMap();
david.put("name", "David");
dynamicSession.save("Customer", david);
...
dynamicSession.flush();
dynamicSession.close();

```

```
...  
// Continue on pojoSession
```

Please note that the call to `getSession()` using an `EntityMode` is on the `Session` API, not the `SessionFactory`. That way, the new `Session` shares the underlying JDBC connection, transaction, and other context information. This means you don't have to call `flush()` and `close()` on the secondary `Session`, and also leave the transaction and connection handling to the primary unit of work.

More information about the XML representation capabilities can be found in Chapter 18, *XML Mapping*.

TODO: Document user-extension framework in the property and proxy packages

---

# Chapter 5. Basic O/R Mapping

## 5.1. Mapping declaration

Object/relational mappings are usually defined in an XML document. The mapping document is designed to be readable and hand-editable. The mapping language is Java-centric, meaning that mappings are constructed around persistent class declarations, not table declarations.

Note that, even though many Hibernate users choose to write the XML by hand, a number of tools exist to generate the mapping document, including XDoclet, Middlegen and AndroMDA.

Lets kick off with an example mapping:

```
<?xml version="1.0"?>
<!DOCTYPE hibernate-mapping PUBLIC
    "-//Hibernate/Hibernate Mapping DTD 3.0//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">

<hibernate-mapping package="eg">

    <class name="Cat"
        table="cats"
        discriminator-value="C">

        <id name="id">
            <generator class="native"/>
        </id>

        <discriminator column="subclass"
            type="character"/>

        <property name="weight"/>

        <property name="birthdate"
            type="date"
            not-null="true"
            update="false"/>

        <property name="color"
            type="eg.types.ColorUserType"
            not-null="true"
            update="false"/>

        <property name="sex"
            not-null="true"
            update="false"/>

        <property name="litterId"
            column="litterId"
            update="false"/>

        <many-to-one name="mother"
            column="mother_id"
            update="false"/>

        <set name="kittens"
            inverse="true"
            order-by="litter_id">
            <key column="mother_id"/>
            <one-to-many class="Cat"/>
        </set>

        <subclass name="DomesticCat"
            discriminator-value="D">
```

```

        <property name="name"
                type="string"/>

    </subclass>

</class>

<class name="Dog">
    <!-- mapping for Dog could go here -->
</class>

</hibernate-mapping>

```

We will now discuss the content of the mapping document. We will only describe the document elements and attributes that are used by Hibernate at runtime. The mapping document also contains some extra optional attributes and elements that affect the database schemas exported by the schema export tool. (For example the `not-null` attribute.)

### 5.1.1. Doctype

All XML mappings should declare the doctype shown. The actual DTD may be found at the URL above, in the directory `hibernate-x.x.x/src/org/hibernate` or in `hibernate3.jar`. Hibernate will always look for the DTD in its classpath first. If you experience lookups of the DTD using an Internet connection, check your DTD declaration against the contents of your classpath.

### 5.1.2. hibernate-mapping

This element has several optional attributes. The `schema` and `catalog` attributes specify that tables referred to in this mapping belong to the named schema and/or catalog. If specified, `tablename`s will be qualified by the given schema and catalog names. If missing, `tablename`s will be unqualified. The `default-cascade` attribute specifies what cascade style should be assumed for properties and collections which do not specify a `cascade` attribute. The `auto-import` attribute lets us use unqualified class names in the query language, by default.

```

<hibernate-mapping
    schema="schemaName"                (1)
    catalog="catalogName"              (2)
    default-cascade="cascade_style"    (3)
    default-access="field|property|ClassName" (4)
    default-lazy="true|false"          (5)
    auto-import="true|false"           (6)
    package="package.name"            (7)
/>

```

- (1) `schema` (optional): The name of a database schema.
- (2) `catalog` (optional): The name of a database catalog.
- (3) `default-cascade` (optional - defaults to `none`): A default cascade style.
- (4) `default-access` (optional - defaults to `property`): The strategy Hibernate should use for accessing all properties. Can be a custom implementation of `PropertyAccessor`.
- (5) `default-lazy` (optional - defaults to `true`): The default value for unspecified `lazy` attributes of class and collection mappings.
- (6) `auto-import` (optional - defaults to `true`): Specifies whether we can use unqualified class names (of classes in this mapping) in the query language.
- (7) `package` (optional): Specifies a package prefix to assume for unqualified class names in the mapping document.



If you have two persistent classes with the same (unqualified) name, you should set `auto-import="false"`. Hibernate will throw an exception if you attempt to assign two classes to the same "imported" name.

Note that the `hibernate-mapping` element allows you to nest several persistent `<class>` mappings, as shown above. It is however good practice (and expected by some tools) to map only a single persistent class (or a single class hierarchy) in one mapping file and name it after the persistent superclass, e.g. `Cat.hbm.xml`, `Dog.hbm.xml`, or if using inheritance, `Animal.hbm.xml`.

### 5.1.3. class

You may declare a persistent class using the `class` element:

```
<class
    name="ClassName"                                (1)
    table="tableName"                               (2)
    discriminator-value="discriminator_value"        (3)
    mutable="true|false"                             (4)
    schema="owner"                                   (5)
    catalog="catalog"                               (6)
    proxy="ProxyInterface"                           (7)
    dynamic-update="true|false"                       (8)
    dynamic-insert="true|false"                       (9)
    select-before-update="true|false"                 (10)
    polymorphism="implicit|explicit"                 (11)
    where="arbitrary sql where condition"            (12)
    persister="PersisterClass"                       (13)
    batch-size="N"                                   (14)
    optimistic-lock="none|version|dirty|all"         (15)
    lazy="true|false"                                (16)
    entity-name="EntityName"                         (17)
    catalog="catalog"                               (18)
    check="arbitrary sql check condition"            (19)
    rowid="rowid"                                    (20)
    subselect="SQL expression"                       (21)
    abstract="true|false"                            (22)
    entity-name="EntityName"
/>
```

- (1) `name` (optional): The fully qualified Java class name of the persistent class (or interface). If this attribute is missing, it is assumed that the mapping is for a non-POJO entity.
- (2) `table` (optional - defaults to the unqualified class name): The name of its database table.
- (3) `discriminator-value` (optional - defaults to the class name): A value that distinguishes individual subclasses, used for polymorphic behaviour. Acceptable values include `null` and `not null`.
- (4) `mutable` (optional, defaults to `true`): Specifies that instances of the class are (not) mutable.
- (5) `schema` (optional): Override the schema name specified by the root `<hibernate-mapping>` element.
- (6) `catalog` (optional): Override the catalog name specified by the root `<hibernate-mapping>` element.
- (7) `proxy` (optional): Specifies an interface to use for lazy initializing proxies. You may specify the name of the class itself.
- (8) `dynamic-update` (optional, defaults to `false`): Specifies that `UPDATE SQL` should be generated at runtime and contain only those columns whose values have changed.
- (9) `dynamic-insert` (optional, defaults to `false`): Specifies that `INSERT SQL` should be generated at runtime and contain only the columns whose values are not null.
- (10) `select-before-update` (optional, defaults to `false`): Specifies that Hibernate should *never* perform an `SQL UPDATE` unless it is certain that an object is actually modified. In certain cases (actually, only when a transient object has been associated with a new session using `update()`), this means that Hibernate will perform an extra `SQL SELECT` to determine if an `UPDATE` is actually required.
- (11) `polymorphism` (optional, defaults to `implicit`): Determines whether implicit or explicit query polymorphism is used.

- (12) `where` (optional): specify an arbitrary SQL `WHERE` condition to be used when retrieving objects of this class
- (13) `persister` (optional): Specifies a custom `ClassPersister`.
- (14) `batch-size` (optional, defaults to 1) specify a "batch size" for fetching instances of this class by identifier.
- (15) `optimistic-lock` (optional, defaults to `version`): Determines the optimistic locking strategy.
- (16) `lazy` (optional): Lazy fetching may be completely disabled by setting `lazy="false"`.
- (17) `entity-name` (optional): Hibernate3 allows a class to be mapped multiple times (to different tables, potentially), and allows entity mappings that are represented by Maps or XML at the Java level. In these cases, you should provide an explicit arbitrary name for the entity. See Section 4.4, "Dynamic models" and Chapter 18, *XML Mapping* for more information.
- (18) `catalog` (optional): The name of a database catalog used for this class and its table.
- (19) `check` (optional): A SQL expression used to generate a multi-row *check* constraint for automatic schema generation.
- (20) `rowid` (optional): Hibernate can use so called ROWIDs on databases which support. E.g. on Oracle, Hibernate can use the `rowid` extra column for fast updates if you set this option to `rowid`. A ROWID is an implementation detail and represents the physical location of a stored tuple.
- (21) `subselect` (optional): Maps an immutable and read-only entity to a database subselect. Useful if you want to have a view instead of a base table, but don't. See below for more information.
- (22) `abstract` (optional): Used to mark abstract superclasses in `<union-subclass>` hierarchies.

It is perfectly acceptable for the named persistent class to be an interface. You would then declare implementing classes of that interface using the `<subclass>` element. You may persist any *static* inner class. You should specify the class name using the standard form ie. eg.`Foo$Bar`.

Immutable classes, `mutable="false"`, may not be updated or deleted by the application. This allows Hibernate to make some minor performance optimizations.

The optional `proxy` attribute enables lazy initialization of persistent instances of the class. Hibernate will initially return CGLIB proxies which implement the named interface. The actual persistent object will be loaded when a method of the proxy is invoked. See "Proxies for Lazy Initialization" below.

*Implicit* polymorphism means that instances of the class will be returned by a query that names any superclass or implemented interface or the class and that instances of any subclass of the class will be returned by a query that names the class itself. *Explicit* polymorphism means that class instances will be returned only by queries that explicitly name that class and that queries that name the class will return only instances of subclasses mapped inside this `<class>` declaration as a `<subclass>` or `<joined-subclass>`. For most purposes the default, `polymorphism="implicit"`, is appropriate. Explicit polymorphism is useful when two different classes are mapped to the same table (this allows a "lightweight" class that contains a subset of the table columns).

The `persister` attribute lets you customize the persistence strategy used for the class. You may, for example, specify your own subclass of `org.hibernate.persister.EntityPersister` or you might even provide a completely new implementation of the interface `org.hibernate.persister.ClassPersister` that implements persistence via, for example, stored procedure calls, serialization to flat files or LDAP. See `org.hibernate.test.CustomPersister` for a simple example (of "persistence" to a `Hashtable`).

Note that the `dynamic-update` and `dynamic-insert` settings are not inherited by subclasses and so may also be specified on the `<subclass>` or `<joined-subclass>` elements. These settings may increase performance in some cases, but might actually decrease performance in others. Use judiciously.

Use of `select-before-update` will usually decrease performance. It is very useful to prevent a database update trigger being called unnecessarily if you reattach a graph of detached instances to a `Session`.

If you enable `dynamic-update`, you will have a choice of optimistic locking strategies:

- `version` check the version/timestamp columns

- `all` check all columns
- `dirty` check the changed columns, allowing some concurrent updates
- `none` do not use optimistic locking

We *very* strongly recommend that you use version/timestamp columns for optimistic locking with Hibernate. This is the optimal strategy with respect to performance and is the only strategy that correctly handles modifications made to detached instances (ie. when `Session.merge()` is used).

There is no difference between a view and a base table for a Hibernate mapping, as expected this is transparent at the database level (note that some DBMS don't support views properly, especially with updates). Sometimes you want to use a view, but can't create one in the database (ie. with a legacy schema). In this case, you can map an immutable and read-only entity to a given SQL subselect expression:

```
<class name="Summary">
  <subselect>
    select item.name, max(bid.amount), count(*)
    from item
    join bid on bid.item_id = item.id
    group by item.name
  </subselect>
  <synchronize table="item"/>
  <synchronize table="bid"/>
  <id name="name"/>
  ...
</class>
```

Declare the tables to synchronize this entity with, ensuring that auto-flush happens correctly, and that queries against the derived entity do not return stale data. The `<subselect>` is available as both as an attribute and a nested mapping element.

#### 5.1.4. id

Mapped classes *must* declare the primary key column of the database table. Most classes will also have a Java-Beans-style property holding the unique identifier of an instance. The `<id>` element defines the mapping from that property to the primary key column.

```
<id
  name="propertyName" (1)
  type="typename" (2)
  column="column_name" (3)
  unsaved-value="null|any|none|undefined|id_value" (4)
  access="field|property|ClassName"> (5)

  <generator class="generatorClass"/>
</id>
```

- (1) `name` (optional): The name of the identifier property.
- (2) `type` (optional): A name that indicates the Hibernate type.
- (3) `column` (optional - defaults to the property name): The name of the primary key column.
- (4) `unsaved-value` (optional - defaults to a "sensible" value): An identifier property value that indicates that an instance is newly instantiated (unsaved), distinguishing it from detached instances that were saved or loaded in a previous session.
- (5) `access` (optional - defaults to `property`): The strategy Hibernate should use for accessing the property value.

If the `name` attribute is missing, it is assumed that the class has no identifier property.

The `unsaved-value` attribute is important! If the identifier property of your class does not default to the normal Java default value (null or zero), then you should specify the actual default.

There is an alternative `<composite-id>` declaration to allow access to legacy data with composite keys. We strongly discourage its use for anything else.

## Generator

The optional `<generator>` child element names a Java class used to generate unique identifiers for instances of the persistent class. If any parameters are required to configure or initialize the generator instance, they are passed using the `<param>` element.

```
<id name="id" type="long" column="cat_id">
  <generator class="org.hibernate.id.TableHiLoGenerator">
    <param name="table">uid_table</param>
    <param name="column">next_hi_value_column</param>
  </generator>
</id>
```

All generators implement the interface `org.hibernate.id.IdentifierGenerator`. This is a very simple interface; some applications may choose to provide their own specialized implementations. However, Hibernate provides a range of built-in implementations. There are shortcut names for the built-in generators:

### increment

generates identifiers of type `long`, `short` or `int` that are unique only when no other process is inserting data into the same table. *Do not use in a cluster.*

### identity

supports identity columns in DB2, MySQL, MS SQL Server, Sybase and HypersonicSQL. The returned identifier is of type `long`, `short` or `int`.

### sequence

uses a sequence in DB2, PostgreSQL, Oracle, SAP DB, McKoi or a generator in Interbase. The returned identifier is of type `long`, `short` or `int`

### hilo

uses a hi/lo algorithm to efficiently generate identifiers of type `long`, `short` or `int`, given a table and column (by default `hibernate_unique_key` and `next_hi` respectively) as a source of hi values. The hi/lo algorithm generates identifiers that are unique only for a particular database.

### seqhilo

uses a hi/lo algorithm to efficiently generate identifiers of type `long`, `short` or `int`, given a named database sequence.

### uuid

uses a 128-bit UUID algorithm to generate identifiers of type `string`, unique within a network (the IP address is used). The UUID is encoded as a string of hexadecimal digits of length 32.

### guid

uses a database-generated GUID string on MS SQL Server and MySQL.

### native

picks `identity`, `sequence` or `hilo` depending upon the capabilities of the underlying database.

### assigned

lets the application to assign an identifier to the object before `save()` is called. This is the default strategy if no `<generator>` element is specified.

`select`

retrieves a primary key assigned by a database trigger by selecting the row by some unique key and retrieving the primary key value.

`foreign`

uses the identifier of another associated object. Usually used in conjunction with a `<one-to-one>` primary key association.

## Hi/lo algorithm

The `hilo` and `seqhilo` generators provide two alternate implementations of the hi/lo algorithm, a favorite approach to identifier generation. The first implementation requires a "special" database table to hold the next available "hi" value. The second uses an Oracle-style sequence (where supported).

```
<id name="id" type="long" column="cat_id">
  <generator class="hilo">
    <param name="table">hi_value</param>
    <param name="column">next_value</param>
    <param name="max_lo">100</param>
  </generator>
</id>
```

```
<id name="id" type="long" column="cat_id">
  <generator class="seqhilo">
    <param name="sequence">hi_value</param>
    <param name="max_lo">100</param>
  </generator>
</id>
```

Unfortunately, you can't use `hilo` when supplying your own `Connection` to Hibernate. When Hibernate is using an application server datasource to obtain connections enlisted with JTA, you must properly configure the `hibernate.transaction.manager_lookup_class`.

## UUID algorithm

The UUID contains: IP address, startup time of the JVM (accurate to a quarter second), system time and a counter value (unique within the JVM). It's not possible to obtain a MAC address or memory address from Java code, so this is the best we can do without using JNI.

## Identity columns and sequences

For databases which support identity columns (DB2, MySQL, Sybase, MS SQL), you may use `identity` key generation. For databases that support sequences (DB2, Oracle, PostgreSQL, Interbase, McKoi, SAP DB) you may use `sequence` style key generation. Both these strategies require two SQL queries to insert a new object.

```
<id name="id" type="long" column="person_id">
  <generator class="sequence">
    <param name="sequence">person_id_sequence</param>
  </generator>
</id>
```

```
<id name="id" type="long" column="person_id" unsaved-value="0">
  <generator class="identity"/>
</id>
```

For cross-platform development, the `native` strategy will choose from the `identity`, `sequence` and `hilo` strategies, dependant upon the capabilities of the underlying database.

## Assigned identifiers

If you want the application to assign identifiers (as opposed to having Hibernate generate them), you may use the `assigned` generator. This special generator will use the identifier value already assigned to the object's identifier property. This generator is used when the primary key is a natural key instead of a surrogate key. This is the default behavior if you do not specify a `<generator>` element.

Choosing the `assigned` generator makes Hibernate use `unsaved-value="undefined"`, forcing Hibernate to go to the database to determine if an instance is transient or detached, unless there is a version or timestamp property, or you define `Interceptor.isUnsaved()`.

## Primary keys assigned by triggers

For legacy schemas only (Hibernate does not generate DDL with triggers).

```
<id name="id" type="long" column="person_id">
  <generator class="select">
    <param name="key">socialSecurityNumber</param>
  </generator>
</id>
```

In the above example, there is a unique valued property named `socialSecurityNumber` defined by the class, as a natural key, and a surrogate key named `person_id` whose value is generated by a trigger.

### 5.1.5. composite-id

```
<composite-id
  name="propertyName"
  class="ClassName"
  unsaved-value="undefined|any|none"
  access="field|property|ClassName">

  <key-property name="propertyName" type="typename" column="column_name" />
  <key-many-to-one name="propertyName" class="ClassName" column="column_name" />
  .....
</composite-id>
```

For a table with a composite key, you may map multiple properties of the class as identifier properties. The `<composite-id>` element accepts `<key-property>` property mappings and `<key-many-to-one>` mappings as child elements.

```
<composite-id>
  <key-property name="medicareNumber" />
  <key-property name="dependent" />
</composite-id>
```

Your persistent class *must* override `equals()` and `hashCode()` to implement composite identifier equality. It must also implement `Serializable`.

Unfortunately, this approach to composite identifiers means that a persistent object is its own identifier. There is no convenient "handle" other than the object itself. You must instantiate an instance of the persistent class itself and populate its identifier properties before you can `load()` the persistent state associated with a composite key. We will describe a much more convenient approach where the composite identifier is implemented as a

separate class in Section 8.4, “Components as composite identifiers”. The attributes described below apply only to this alternative approach:

- `name` (optional): A property of component type that holds the composite identifier (see next section).
- `class` (optional - defaults to the property type determined by reflection): The component class used as a composite identifier (see next section).
- `unsaved-value` (optional - defaults to `undefined`): Indicates that transient instances should be considered newly instantiated, if set to `any`, or detached, if set to `none`. It is best to leave the default value in all cases.

### 5.1.6. discriminator

The `<discriminator>` element is required for polymorphic persistence using the table-per-class-hierarchy mapping strategy and declares a discriminator column of the table. The discriminator column contains marker values that tell the persistence layer what subclass to instantiate for a particular row. A restricted set of types may be used: `string`, `character`, `integer`, `byte`, `short`, `boolean`, `yes_no`, `true_false`.

```
<discriminator
  column="discriminator_column"           (1)
  type="discriminator_type"               (2)
  force="true|false"                      (3)
  insert="true|false"                     (4)
  formula="arbitrary sql expression"      (5)
/>
```

- (1) `column` (optional - defaults to `class`) the name of the discriminator column.
- (2) `type` (optional - defaults to `string`) a name that indicates the Hibernate type
- (3) `force` (optional - defaults to `false`) "force" Hibernate to specify allowed discriminator values even when retrieving all instances of the root class.
- (4) `insert` (optional - defaults to `true`) set this to `false` if your discriminator column is also part of a mapped composite identifier. (Tells Hibernate to not include the column in SQL `INSERTS`.)
- (5) `formula` (optional) an arbitrary SQL expression that is executed when a type has to be evaluated. Allows content-based discrimination.

Actual values of the discriminator column are specified by the `discriminator-value` attribute of the `<class>` and `<subclass>` elements.

The `force` attribute is (only) useful if the table contains rows with "extra" discriminator values that are not mapped to a persistent class. This will not usually be the case.

Using the `formula` attribute you can declare an arbitrary SQL expression that will be used to evaluate the type of a row:

```
<discriminator
  formula="case when CLASS_TYPE in ('a', 'b', 'c') then 0 else 1 end"
  type="integer"/>
```

### 5.1.7. version (optional)

The `<version>` element is optional and indicates that the table contains versioned data. This is particularly useful if you plan to use *long transactions* (see below).

```
<version
  column="version_column"                 (1)
  name="propertyName"                    (2)
  type="typename"                         (3)
  access="field|property|ClassName"       (4)
  unsaved-value="null|negative|undefined" (5)
```

```
</>
```

- (1) `column` (optional - defaults to the property name): The name of the column holding the version number.
- (2) `name`: The name of a property of the persistent class.
- (3) `type` (optional - defaults to `integer`): The type of the version number.
- (4) `access` (optional - defaults to `property`): The strategy Hibernate should use for accessing the property value.
- (5) `unsaved-value` (optional - defaults to `undefined`): A version property value that indicates that an instance is newly instantiated (unsaved), distinguishing it from detached instances that were saved or loaded in a previous session. (`undefined` specifies that the identifier property value should be used.)

Version numbers may be of Hibernate type `long`, `integer`, `short`, `timestamp` or `calendar`.

A version or timestamp property should never be null for a detached instance, so Hibernate will detect any instance with a null version or timestamp as transient, no matter what other `unsaved-value` strategies are specified. *Declaring a nullable version or timestamp property is an easy way to avoid any problems with transitive reattachment in Hibernate, especially useful for people using assigned identifiers or composite keys!*

### 5.1.8. timestamp (optional)

The optional `<timestamp>` element indicates that the table contains timestamped data. This is intended as an alternative to versioning. Timestamps are by nature a less safe implementation of optimistic locking. However, sometimes the application might use the timestamps in other ways.

```
<timestamp
  column="timestamp_column"      (1)
  name="propertyName"          (2)
  access="field|property|ClassName" (3)
  unsaved-value="null|undefined" (4)
/>
```

- (1) `column` (optional - defaults to the property name): The name of a column holding the timestamp.
- (2) `name`: The name of a JavaBeans style property of Java type `Date` or `Timestamp` of the persistent class.
- (3) `access` (optional - defaults to `property`): The strategy Hibernate should use for accessing the property value.
- (4) `unsaved-value` (optional - defaults to `null`): A version property value that indicates that an instance is newly instantiated (unsaved), distinguishing it from detached instances that were saved or loaded in a previous session. (`undefined` specifies that the identifier property value should be used.)

Note that `<timestamp>` is equivalent to `<version type="timestamp">`.

### 5.1.9. property

The `<property>` element declares a persistent, JavaBean style property of the class.

```
<property
  name="propertyName"          (1)
  column="column_name"         (2)
  type="typename"              (3)
  update="true|false"          (4)
  insert="true|false"          (4)
  formula="arbitrary SQL expression" (5)
  access="field|property|ClassName" (6)
  lazy="true|false"            (7)
  unique="true|false"          (8)
  not-null="true|false"        (9)
/>
```



```
optimistic-lock="true|false" (10)
/>
```

- (1) **name**: the name of the property, with an initial lowercase letter.
- (2) **column** (optional - defaults to the property name): the name of the mapped database table column. This may also be specified by nested `<column>` element(s).
- (3) **type** (optional): a name that indicates the Hibernate type.
- (4) **update**, **insert** (optional - defaults to `true`): specifies that the mapped columns should be included in SQL `UPDATE` and/or `INSERT` statements. Setting both to `false` allows a pure "derived" property whose value is initialized from some other property that maps to the same column(s) or by a trigger or other application.
- (5) **formula** (optional): an SQL expression that defines the value for a *computed* property. Computed properties do not have a column mapping of their own.
- (6) **access** (optional - defaults to `property`): The strategy Hibernate should use for accessing the property value.
- (7) **lazy** (optional - defaults to `false`): Specifies that this property should be fetched lazily when the instance variable is first accessed (requires build-time bytecode instrumentation).
- (8) **unique** (optional): Enable the DDL generation of a unique constraint for the columns. Also, allow this to be the target of a `property-ref`.
- (9) **not-null** (optional): Enable the DDL generation of a nullability constraint for the columns.
- (10) **optimistic-lock** (optional - defaults to `true`): Specifies that updates to this property do or do not require acquisition of the optimistic lock. In other words, determines if a version increment should occur when this property is dirty.

*typename* could be:

1. The name of a Hibernate basic type (eg. `integer`, `string`, `character`, `date`, `timestamp`, `float`, `binary`, `serializable`, `object`, `blob`).
2. The name of a Java class with a default basic type (eg. `int`, `float`, `char`, `java.lang.String`, `java.util.Date`, `java.lang.Integer`, `java.sql.Clob`).
3. The name of a serializable Java class.
4. The class name of a custom type (eg. `com.illflow.type.MyCustomType`).

If you do not specify a type, Hibernate will use reflection upon the named property to take a guess at the correct Hibernate type. Hibernate will try to interpret the name of the return class of the property getter using rules 2, 3, 4 in that order. However, this is not always enough. In certain cases you will still need the `type` attribute. (For example, to distinguish between `Hibernate.DATE` and `Hibernate.TIMESTAMP`, or to specify a custom type.)

The `access` attribute lets you control how Hibernate will access the property at runtime. By default, Hibernate will call the property `get/set` pair. If you specify `access="field"`, Hibernate will bypass the `get/set` pair and access the field directly, using reflection. You may specify your own strategy for property access by naming a class that implements the interface `org.hibernate.property.PropertyAccessor`.

An especially powerful feature are derived properties. These properties are by definition read-only, the property value is computed at load time. You declare the computation as a SQL expression, this translates to a `SELECT` clause subquery in the SQL query that loads an instance:

```
<property name="totalPrice"
  formula="( SELECT SUM (li.quantity*p.price) FROM LineItem li, Product p
            WHERE li.productId = p.productId
            AND li.customerId = customerId
            AND li.orderNumber = orderNumber )"/>
```

Note that you can reference the entities own table by not declaring an alias on a particular column (`customerId`

in the given example). Also note that you can use the nested `<formula>` mapping element if you don't like to use the attribute.

### 5.1.10. many-to-one

An ordinary association to another persistent class is declared using a `many-to-one` element. The relational model is a many-to-one association: a foreign key in one table is referencing the primary key column(s) of the target table.

```
<many-to-one
  name="propertyName"                (1)
  column="column_name"               (2)
  class="ClassName"                  (3)
  cascade="cascade_style"            (4)
  fetch="join|select"                (5)
  update="true|false"                (6)
  insert="true|false"                (6)
  property-ref="propertyNameFromAssociatedClass" (7)
  access="field|property|ClassName" (8)
  unique="true|false"                (9)
  not-null="true|false"              (10)
  optimistic-lock="true|false"       (11)
  lazy="true|false"                  (12)
  entity-name="EntityName"
/>
```

- (1) `name`: The name of the property.
- (2) `column` (optional): The name of the foreign key column. This may also be specified by nested `<column>` element(s).
- (3) `class` (optional - defaults to the property type determined by reflection): The name of the associated class.
- (4) `cascade` (optional): Specifies which operations should be cascaded from the parent object to the associated object.
- (5) `fetch` (optional - defaults to `select`): Chooses between outer-join fetching or sequential select fetching.
- (6) `update`, `insert` (optional - defaults to `true`) specifies that the mapped columns should be included in SQL `UPDATE` and/or `INSERT` statements. Setting both to `false` allows a pure "derived" association whose value is initialized from some other property that maps to the same column(s) or by a trigger or other application.
- (7) `property-ref`: (optional) The name of a property of the associated class that is joined to this foreign key. If not specified, the primary key of the associated class is used.
- (8) `access` (optional - defaults to `property`): The strategy Hibernate should use for accessing the property value.
- (9) `unique` (optional): Enable the DDL generation of a unique constraint for the foreign-key column. Also, allow this to be the target of a `property-ref`. This makes the association multiplicity effectively one to one.
- (10) `not-null` (optional): Enable the DDL generation of a nullability constraint for the foreign key columns.
- (11) `optimistic-lock` (optional - defaults to `true`): Specifies that updates to this property do or do not require acquisition of the optimistic lock. In other words, determines if a version increment should occur when this property is dirty.
- (12) `lazy` (optional - defaults to `false`): Specifies that this property should be fetched lazily when the instance variable is first accessed (requires build-time bytecode instrumentation). Note that this does not influence Hibernate's proxy behavior - like the `lazy` attribute on class or collection mappings, but uses interception for deferred loading.

Setting a value of the `cascade` attribute to any meaningful value other than `none` will propagate certain operations to the associated object. The meaningful values are the names of Hibernate's basic operations, `create`, `merge`, `delete`, `save-update`, `evict`, `replicate`, `lock`, `refresh`, as well as the special values `delete-`

orphan and all and comma-separated combinations of operation names, for example, cascade="create,merge,evict" or cascade="all,delete-orphan". See Section 10.11, "Transitive persistence" for a full explanation.

A typical many-to-one declaration looks as simple as this:

```
<many-to-one name="product" class="Product" column="PRODUCT_ID"/>
```

The `property-ref` attribute should only be used for mapping legacy data where a foreign key refers to a unique key of the associated table other than the primary key. This is an ugly relational model. For example, suppose the `Product` class had a unique serial number, that is not the primary key. (The `unique` attribute controls Hibernate's DDL generation with the `SchemaExport` tool.)

```
<property name="serialNumber" unique="true" type="string" column="SERIAL_NUMBER"/>
```

Then the mapping for `OrderItem` might use:

```
<many-to-one name="product" property-ref="serialNumber" column="PRODUCT_SERIAL_NUMBER"/>
```

This is certainly not encouraged, however.

If the referenced unique key comprises multiple properties of the associated entity, you should map the referenced properties inside a named `<properties>` element.

### 5.1.11. one-to-one

A one-to-one association to another persistent class is declared using a `one-to-one` element.

```
<one-to-one
    name="propertyName" (1)
    class="ClassName" (2)
    cascade="cascade_style" (3)
    constrained="true|false" (4)
    fetch="join|select" (5)
    property-ref="propertyNameFromAssociatedClass" (6)
    access="field|property|ClassName" (7)
    formula="any SQL expression" (8)
    entity-name="EntityName"
/>
```

- (1) `name`: The name of the property.
- (2) `class` (optional - defaults to the property type determined by reflection): The name of the associated class.
- (3) `cascade` (optional) specifies which operations should be cascaded from the parent object to the associated object.
- (4) `constrained` (optional) specifies that a foreign key constraint on the primary key of the mapped table references the table of the associated class. This option affects the order in which `save()` and `delete()` are cascaded, and determines whether the association may be proxied (it is also used by the schema export tool).
- (5) `fetch` (optional - defaults to `select`): Chooses between outer-join fetching or sequential select fetching.
- (6) `property-ref`: (optional) The name of a property of the associated class that is joined to the primary key of this class. If not specified, the primary key of the associated class is used.
- (7) `access` (optional - defaults to `property`): The strategy Hibernate should use for accessing the property value.
- (8) `formula` (optional): Almost all one to one associations map to the primary key of the owning entity. In the rare case that this is not the case, you may specify a some other column, columns or expression to join on

using an SQL formula. (See `org.hibernate.test.onetooneformula` for an example.)

There are two varieties of one-to-one association:

- primary key associations
- unique foreign key associations

Primary key associations don't need an extra table column; if two rows are related by the association then the two table rows share the same primary key value. So if you want two objects to be related by a primary key association, you must make sure that they are assigned the same identifier value!

For a primary key association, add the following mappings to `Employee` and `Person`, respectively.

```
<one-to-one name="person" class="Person"/>
```

```
<one-to-one name="employee" class="Employee" constrained="true"/>
```

Now we must ensure that the primary keys of related rows in the `PERSON` and `EMPLOYEE` tables are equal. We use a special Hibernate identifier generation strategy called `foreign`:

```
<class name="person" table="PERSON">
  <id name="id" column="PERSON_ID">
    <generator class="foreign">
      <param name="property">employee</param>
    </generator>
  </id>
  ...
  <one-to-one name="employee"
    class="Employee"
    constrained="true"/>
</class>
```

A newly saved instance of `Person` is then assigned the same primary key value as the `Employee` instance referred with the `employee` property of that `Person`.

Alternatively, a foreign key with a unique constraint, from `Employee` to `Person`, may be expressed as:

```
<many-to-one name="person" class="Person" column="PERSON_ID" unique="true"/>
```

And this association may be made bidirectional by adding the following to the `Person` mapping:

```
<one-to-one name="employee" class="Employee" property-ref="person"/>
```

### 5.1.12. component, dynamic-component

The `<component>` element maps properties of a child object to columns of the table of a parent class. Components may, in turn, declare their own properties, components or collections. See "Components" below.

```
<component
  name="propertyName"                (1)
  class="className"                  (2)
  insert="true|false"                 (3)
  update="true|false"                 (4)
  access="field|property|ClassName" (5)
  lazy="true|false"                   (6)
  optimistic-lock="true|false"        (7)
  unique="true|false"                  (8)
>
```

```

    <property ...../>
    <many-to-one .... />
    .....
</component>

```

- (1) **name**: The name of the property.
- (2) **class** (optional - defaults to the property type determined by reflection): The name of the component (child) class.
- (3) **insert**: Do the mapped columns appear in SQL `INSERTS`?
- (4) **update**: Do the mapped columns appear in SQL `UPDATES`?
- (5) **access** (optional - defaults to `property`): The strategy Hibernate should use for accessing the property value.
- (6) **lazy** (optional - defaults to `false`): Specifies that this component should be fetched lazily when the instance variable is first accessed (requires build-time bytecode instrumentation).
- (7) **optimistic-lock** (optional - defaults to `true`): Specifies that updates to this component do or do not require acquisition of the optimistic lock. In other words, determines if a version increment should occur when this property is dirty.
- (8) **unique** (optional - defaults to `false`): Specifies that a unique constraint exists upon all mapped columns of the component.

The child `<property>` tags map properties of the child class to table columns.

The `<component>` element allows a `<parent>` subelement that maps a property of the component class as a reference back to the containing entity.

The `<dynamic-component>` element allows a `Map` to be mapped as a component, where the property names refer to keys of the map, see Section 8.5, “Dynamic components”.

### 5.1.13. properties

The `<properties>` element allows the definition of a named, logical grouping of properties of a class. The most important use of the construct is that it allows a combination of properties to be the target of a `property-ref`. It is also a convenient way to define a multi-column unique constraint.

```

<properties
  name="logicalName"           (1)
  insert="true|false"         (2)
  update="true|false"         (3)
  optimistic-lock="true|false" (4)
  unique="true|false"         (5)
>

  <property ...../>
  <many-to-one .... />
  .....
</properties>

```

- (1) **name**: The logical name of the grouping - *not* an actual property name.
- (2) **insert**: Do the mapped columns appear in SQL `INSERTS`?
- (3) **update**: Do the mapped columns appear in SQL `UPDATES`?
- (4) **optimistic-lock** (optional - defaults to `true`): Specifies that updates to these properties do or do not require acquisition of the optimistic lock. In other words, determines if a version increment should occur when these properties are dirty.
- (5) **unique** (optional - defaults to `false`): Specifies that a unique constraint exists upon all mapped columns of the component.

For example, if we have the following `<properties>` mapping:

```
<class name="Person">
  <id name="personNumber" />
  ...
  <properties name="name"
    unique="true" update="false">
    <property name="firstName" />
    <property name="initial" />
    <property name="lastName" />
  </properties>
</class>
```

Then we might have some legacy data association which refers to this unique key of the `Person` table, instead of to the primary key:

```
<many-to-one name="person"
  class="Person" property-ref="name">
  <column name="firstName" />
  <column name="initial" />
  <column name="lastName" />
</many-to-one>
```

We don't recommend the use of this kind of thing outside the context of mapping legacy data.

### 5.1.14. subclass

Finally, polymorphic persistence requires the declaration of each subclass of the root persistent class. For the table-per-class-hierarchy mapping strategy, the `<subclass>` declaration is used.

```
<subclass
  name="ClassName" (1)
  discriminator-value="discriminator_value" (2)
  proxy="ProxyInterface" (3)
  lazy="true|false" (4)
  dynamic-update="true|false"
  dynamic-insert="true|false"
  entity-name="EntityName">

  <property .... />
  .....
</subclass>
```

- (1) `name`: The fully qualified class name of the subclass.
- (2) `discriminator-value` (optional - defaults to the class name): A value that distinguishes individual subclasses.
- (3) `proxy` (optional): Specifies a class or interface to use for lazy initializing proxies.
- (4) `lazy` (optional, defaults to `true`): Setting `lazy="false"` disables the use of lazy fetching.

Each subclass should declare its own persistent properties and subclasses. `<version>` and `<id>` properties are assumed to be inherited from the root class. Each subclass in a hierarchy must define a unique `discriminator-value`. If none is specified, the fully qualified Java class name is used.

It is possible to define `subclass`, `union-subclass`, and `joined-subclass` mappings in separate mapping documents, directly beneath `hibernate-mapping`. This allows you to extend a class hierarchy just by adding a new mapping file. You must specify an `extends` attribute in the subclass mapping, naming a previously mapped superclass. Note: Previously this feature made the ordering of the mapping documents important. Since Hibernate3, the ordering of mapping files does not matter when using the `extends` keyword. The ordering inside a single mapping file still needs to be defined as superclasses before subclasses.

```
<hibernate-mapping>
  <subclass name="DomesticCat" extends="Cat" discriminator-value="D">
    <property name="name" type="string" />
  </subclass>
</hibernate-mapping>
```

For information about inheritance mappings, see Chapter 9, *Inheritance Mapping*.

### 5.1.15. joined-subclass

Alternatively, each subclass may be mapped to its own table (table-per-subclass mapping strategy). Inherited state is retrieved by joining with the table of the superclass. We use the `<joined-subclass>` element.

```
<joined-subclass
  name="ClassName" (1)
  table="tablename" (2)
  proxy="ProxyInterface" (3)
  lazy="true|false" (4)
  dynamic-update="true|false"
  dynamic-insert="true|false"
  schema="schema"
  catalog="catalog"
  extends="SuperclassName"
  persister="ClassName"
  subselect="SQL expression"
  entity-name="EntityName">

  <key .... >

  <property .... />
  .....
</joined-subclass>
```

- (1) `name`: The fully qualified class name of the subclass.
- (2) `table`: The name of the subclass table.
- (3) `proxy` (optional): Specifies a class or interface to use for lazy initializing proxies.
- (4) `lazy` (optional, defaults to `true`): Setting `lazy="false"` disables the use of lazy fetching.

No discriminator column is required for this mapping strategy. Each subclass must, however, declare a table column holding the object identifier using the `<key>` element. The mapping at the start of the chapter would be re-written as:

```
<?xml version="1.0"?>
<!DOCTYPE hibernate-mapping PUBLIC
  "-//Hibernate/Hibernate Mapping DTD//EN"
  "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">

<hibernate-mapping package="eg">

  <class name="Cat" table="CATS">
    <id name="id" column="uid" type="long">
      <generator class="hilo"/>
    </id>
    <property name="birthdate" type="date"/>
    <property name="color" not-null="true"/>
    <property name="sex" not-null="true"/>
    <property name="weight"/>
    <many-to-one name="mate"/>
    <set name="kittens">
      <key column="MOTHER"/>
      <one-to-many class="Cat"/>
    </set>
    <joined-subclass name="DomesticCat" table="DOMESTIC_CATS">
```

```

        <key column="CAT"/>
        <property name="name" type="string"/>
    </joined-subclass>
</class>

<class name="eg.Dog">
    <!-- mapping for Dog could go here -->
</class>

</hibernate-mapping>

```

For information about inheritance mappings, see Chapter 9, *Inheritance Mapping*.

### 5.1.16. union-subclass

A third option is to map only the concrete classes of an inheritance hierarchy to tables, (the table-per-concrete-class strategy) where each table defines all persistent state of the class, including inherited state. In Hibernate, it is not absolutely necessary to explicitly map such inheritance hierarchies. You can simply map each class with a separate `<class>` declaration. However, if you wish use polymorphic associations (e.g. an association to the superclass of your hierarchy), you need to use the `<union-subclass>` mapping.

```

<union-subclass
    name="ClassName"                (1)
    table="tablename"              (2)
    proxy="ProxyInterface"         (3)
    lazy="true|false"              (4)
    dynamic-update="true|false"
    dynamic-insert="true|false"
    schema="schema"
    catalog="catalog"
    extends="SuperclassName"
    abstract="true|false"
    persister="ClassName"
    subselect="SQL expression"
    entity-name="EntityName">

    <property .... />
    .....
</union-subclass>

```

- (1) name: The fully qualified class name of the subclass.
- (2) table: The name of the subclass table.
- (3) proxy (optional): Specifies a class or interface to use for lazy initializing proxies.
- (4) lazy (optional, defaults to true): Setting lazy="false" disables the use of lazy fetching.

No discriminator column or key column is required for this mapping strategy.

For information about inheritance mappings, see Chapter 9, *Inheritance Mapping*.

### 5.1.17. join

Using the `<join>` element, it is possible to map properties of one class to several tables.

```

<join
    table="tablename"              (1)
    schema="owner"                 (2)
    catalog="catalog"              (3)
    fetch="join|select"            (4)
    inverse="true|false"           (5)
    optional="true|false">        (6)

```



```

    <key ... />

    <property ... />
    ...
</join>

```

- (1) `table`: The name of the joined table.
- (2) `schema` (optional): Override the schema name specified by the root `<hibernate-mapping>` element.
- (3) `catalog` (optional): Override the catalog name specified by the root `<hibernate-mapping>` element.
- (4) `fetch` (optional - defaults to `join`): If set to `join`, the default, Hibernate will use an inner join to retrieve a `<join>` defined by a class or its superclasses and an outer join for a `<join>` defined by a subclass. If set to `select` then Hibernate will use a sequential select for a `<join>` defined on a subclass, which will be issued only if a row turns out to represent an instance of the subclass. Inner joins will still be used to retrieve a `<join>` defined by the class and its superclasses.
- (5) `inverse` (optional - defaults to `false`): If enabled, Hibernate will not try to insert or update the properties defined by this join.
- (6) `optional` (optional - defaults to `false`): If enabled, Hibernate will insert a row only if the properties defined by this join are non-null and will always use an outer join to retrieve the properties.

For example, the address information for a person can be mapped to a separate table (while preserving value type semantics for all properties):

```

<class name="Person"
  table="PERSON">

  <id name="id" column="PERSON_ID">...</id>

  <join table="ADDRESS">
    <key column="ADDRESS_ID"/>
    <property name="address"/>
    <property name="zip"/>
    <property name="country"/>
  </join>
  ...

```

This feature is often only useful for legacy data models, we recommend fewer tables than classes and a fine-grained domain model. However, it is useful for switching between inheritance mapping strategies in a single hierarchy, as explained later.

### 5.1.18. key

We've seen the `<key>` element crop up a few times now. It appears anywhere the parent mapping element defines a join to a new table, and defines the foreign key in the joined table, that references the primary key of the original table.

```

<key
  column="columnname"                (1)
  on-delete="noaction|cascade"       (2)
  property-ref="propertyName"       (3)
  not-null="true|false"              (4)
  update="true|false"                (5)
  unique="true|false"                (6)
/>

```

- (1) `column` (optional): The name of the foreign key column. This may also be specified by nested `<column>` element(s).
- (2) `on-delete` (optional, defaults to `noaction`): Specifies whether the foreign key constraint has database-level cascade delete enabled.

- (3) `property-ref` (optional): Specifies that the foreign key refers to columns that are not the primary key of the original table. (Provided for legacy data.)
- (4) `not-null` (optional): Specifies that the foreign key columns are not nullable (this is implied whenever the foreign key is also part of the primary key).
- (5) `update` (optional): Specifies that the foreign key should never be updated (this is implied whenever the foreign key is also part of the primary key).
- (6) `unique` (optional): Specifies that the foreign key should have a unique constraint (this is implied whenever the foreign key is also the primary key).

We recommend that for systems where delete performance is important, all keys should be defined `on-delete="cascade"`, and Hibernate will use a database-level `ON CASCADE DELETE` constraint, instead of many individual `DELETE` statements. Be aware that this feature bypasses Hibernate's usual optimistic locking strategy for versioned data.

The `not-null` and `update` attributes are useful when mapping a unidirectional one to many association. If you map a unidirectional one to many to a non-nullable foreign key, you *must* declare the key column using `<key not-null="true">`.

### 5.1.19. column and formula elements

Any mapping element which accepts a `column` attribute will alternatively accept a `<column>` subelement. Likewise, `<formula>` is an alternative to the `formula` attribute.

```
<column
  name="column_name"
  length="N"
  precision="N"
  scale="N"
  not-null="true|false"
  unique="true|false"
  unique-key="multicolumn_unique_key_name"
  index="index_name"
  sql-type="sql_type_name"
  check="SQL expression"/>
```

```
<formula>SQL expression</formula>
```

`column` and `formula` attributes may even be combined within the same property or association mapping to express, for example, exotic join conditions.

```
<many-to-one name="homeAddress" class="Address"
  insert="false" update="false">
  <column name="person_id" not-null="true" length="10"/>
  <formula>'MAILING'</formula>
</many-to-one>
```

### 5.1.20. import

Suppose your application has two persistent classes with the same name, and you don't want to specify the fully qualified (package) name in Hibernate queries. Classes may be "imported" explicitly, rather than relying upon `auto-import="true"`. You may even import classes and interfaces that are not explicitly mapped.

```
<import class="java.lang.Object" rename="Universe"/>
```

```
<import
  class="ClassName"           (1)
  rename="ShortName"         (2)
```

```
</>
```

- (1) `class`: The fully qualified class name of any Java class.
- (2) `rename` (optional - defaults to the unqualified class name): A name that may be used in the query language.

### 5.1.21. any

There is one further type of property mapping. The `<any>` mapping element defines a polymorphic association to classes from multiple tables. This type of mapping always requires more than one column. The first column holds the type of the associated entity. The remaining columns hold the identifier. It is impossible to specify a foreign key constraint for this kind of association, so this is most certainly not meant as the usual way of mapping (polymorphic) associations. You should use this only in very special cases (eg. audit logs, user session data, etc).

The `meta-type` attribute lets the application specify a custom type that maps database column values to persistent classes which have identifier properties of the type specified by `id-type`. You must specify the mapping from values of the meta-type to class names.

```
<any name="being" id-type="long" meta-type="string">
  <meta-value value="TBL_ANIMAL" class="Animal"/>
  <meta-value value="TBL_HUMAN" class="Human"/>
  <meta-value value="TBL_ALIEN" class="Alien"/>
  <column name="table_name"/>
  <column name="id"/>
</any>
```

```
<any
  name="propertyName"                (1)
  id-type="idtypename"                (2)
  meta-type="metatypename"           (3)
  cascade="cascade_style"             (4)
  access="field|property|ClassName"   (5)
  optimistic-lock="true|false"        (6)
>
  <meta-value ... />
  <meta-value ... />
  .....
  <column .... />
  <column .... />
  .....
</any>
```

- (1) `name`: the property name.
- (2) `id-type`: the identifier type.
- (3) `meta-type` (optional - defaults to `string`): Any type that is allowed for a discriminator mapping.
- (4) `cascade` (optional- defaults to `none`): the cascade style.
- (5) `access` (optional - defaults to `property`): The strategy Hibernate should use for accessing the property value.
- (6) `optimistic-lock` (optional - defaults to `true`): Specifies that updates to this property do or do not require acquisition of the optimistic lock. In other words, define if a version increment should occur if this property is dirty.

## 5.2. Hibernate Types

### 5.2.1. Entities and values

To understand the behaviour of various Java language-level objects with respect to the persistence service, we need to classify them into two groups:

An *entity* exists independently of any other objects holding references to the entity. Contrast this with the usual Java model where an unreferenced object is garbage collected. Entities must be explicitly saved and deleted (except that saves and deletions may be *cascaded* from a parent entity to its children). This is different from the ODMG model of object persistence by reachability - and corresponds more closely to how application objects are usually used in large systems. Entities support circular and shared references. They may also be versioned.

An entity's persistent state consists of references to other entities and instances of *value* types. Values are primitives, collections (not what's inside a collection), components and certain immutable objects. Unlike entities, values (in particular collections and components) *are* persisted and deleted by reachability. Since value objects (and primitives) are persisted and deleted along with their containing entity they may not be independently versioned. Values have no independent identity, so they cannot be shared by two entities or collections.

Up until now, we've been using the term "persistent class" to refer to entities. We will continue to do that. Strictly speaking, however, not all user-defined classes with persistent state are entities. A *component* is a user defined class with value semantics. A Java property of type `java.lang.String` also has value semantics. Given this definition, we can say that all types (classes) provided by the JDK have value type semantics in Java, while user-defined types may be mapped with entity or value type semantics. This decision is up to the application developer. A good hint for an entity class in a domain model are shared references to a single instance of that class, while composition or aggregation usually translates to a value type.

We'll revisit both concepts throughout the documentation.

The challenge is to map the Java type system (and the developers' definition of entities and value types) to the SQL/database type system. The bridge between both systems is provided by Hibernate: for entities we use `<class>`, `<subclass>` and so on. For value types we use `<property>`, `<component>`, etc, usually with a `type` attribute. The value of this attribute is the name of a Hibernate *mapping type*. Hibernate provides many mappings (for standard JDK value types) out of the box. You can write your own mapping types and implement your custom conversion strategies as well, as you'll see later.

All built-in Hibernate types except collections support null semantics.

### 5.2.2. Basic value types

The built-in *basic mapping types* may be roughly categorized into

`integer, long, short, float, double, character, byte, boolean, yes_no, true_false`

Type mappings from Java primitives or wrapper classes to appropriate (vendor-specific) SQL column types. `boolean`, `yes_no` and `true_false` are all alternative encodings for a Java `boolean` or `java.lang.Boolean`.

`string`

A type mapping from `java.lang.String` to `VARCHAR` (or Oracle `VARCHAR2`).

`date, time, timestamp`

Type mappings from `java.util.Date` and its subclasses to SQL types `DATE`, `TIME` and `TIMESTAMP` (or equivalent).

`calendar, calendar_date`

Type mappings from `java.util.Calendar` to SQL types `TIMESTAMP` and `DATE` (or equivalent).

`big_decimal`

A type mapping from `java.math.BigDecimal` to `NUMERIC` (or Oracle `NUMBER`).

`locale, timezone, currency`

Type mappings from `java.util.Locale`, `java.util.TimeZone` and `java.util.Currency` to `VARCHAR` (or Oracle `VARCHAR2`). Instances of `Locale` and `Currency` are mapped to their ISO codes. Instances of `TimeZone` are mapped to their ID.

`class`

A type mapping from `java.lang.Class` to `VARCHAR` (or Oracle `VARCHAR2`). A `Class` is mapped to its fully qualified name.

`binary`

Maps byte arrays to an appropriate SQL binary type.

`text`

Maps long Java strings to a SQL `CLOB` or `TEXT` type.

`serializable`

Maps serializable Java types to an appropriate SQL binary type. You may also indicate the Hibernate type `serializable` with the name of a serializable Java class or interface that does not default to a basic type.

`clob, blob`

Type mappings for the JDBC classes `java.sql.Clob` and `java.sql.Blob`. These types may be inconvenient for some applications, since the blob or clob object may not be reused outside of a transaction. (Furthermore, driver support is patchy and inconsistent.)

Unique identifiers of entities and collections may be of any basic type except `binary`, `blob` and `clob`. (Composite identifiers are also allowed, see below.)

The basic value types have corresponding `Type` constants defined on `org.hibernate.Hibernate`. For example, `Hibernate.STRING` represents the `string` type.

### 5.2.3. Custom value types

It is relatively easy for developers to create their own value types. For example, you might want to persist properties of type `java.lang.BigInteger` to `VARCHAR` columns. Hibernate does not provide a built-in type for this. But custom types are not limited to mapping a property (or collection element) to a single table column. So, for example, you might have a Java property `getName()/setName()` of type `java.lang.String` that is persisted to the columns `FIRST_NAME`, `INITIAL`, `SURNAME`.

To implement a custom type, implement either `org.hibernate.UserType` or `org.hibernate.CompositeUserType` and declare properties using the fully qualified classname of the type. Check out `org.hibernate.test.DoubleStringType` to see the kind of things that are possible.

```
<property name="twoStrings" type="org.hibernate.test.DoubleStringType">
  <column name="first_string"/>
  <column name="second_string"/>
</property>
```

Notice the use of `<column>` tags to map a property to multiple columns.

The `CompositeUserType`, `EnhancedUserType`, `UserCollectionType`, and `UserVersionType` interfaces provide

support for more specialized uses.

You may even supply parameters to a `UserType` in the mapping file. To do this, your `UserType` must implement the `org.hibernate.usertype.ParameterizedType` interface. To supply parameters to your custom type, you can use the `<type>` element in your mapping files.

```
<property name="priority">
  <type name="com.mycompany.usertypes.DefaultValueIntegerType">
    <param name="default">0</param>
  </type>
</property>
```

The `UserType` can now retrieve the value for the parameter named `default` from the `Properties` object passed to it.

If you use a certain `UserType` very often, it may be useful to define a shorter name for it. You can do this using the `<typedef>` element. `Typedefs` assign a name to a custom type, and may also contain a list of default parameter values if the type is parameterized.

```
<typedef class="com.mycompany.usertypes.DefaultValueIntegerType" name="default_zero">
  <param name="default">0</param>
</typedef>
```

```
<property name="priority" type="default_zero"/>
```

It is also possible to override the parameters supplied in a `typedef` on a case-by-case basis by using type parameters on the property mapping.

Even though Hibernate's rich range of built-in types and support for components means you will very rarely *need* to use a custom type, it is nevertheless considered good form to use custom types for (non-entity) classes that occur frequently in your application. For example, a `MonetaryAmount` class is a good candidate for a `CompositeUserType`, even though it could easily be mapped as a component. One motivation for this is abstraction. With a custom type, your mapping documents would be future-proofed against possible changes in your way of representing monetary values.

## 5.3. SQL quoted identifiers

You may force Hibernate to quote an identifier in the generated SQL by enclosing the table or column name in backticks in the mapping document. Hibernate will use the correct quotation style for the `SQL Dialect` (usually double quotes, but brackets for SQL Server and backticks for MySQL).

```
<class name="LineItem" table="`Line Item`">
  <id name="id" column="`Item Id`"/><generator class="assigned"/></id>
  <property name="itemNumber" column="`Item #`"/>
  ...
</class>
```

## 5.4. Metadata alternatives

XML isn't for everyone, and so there are some alternative ways to define O/R mapping metadata in Hibernate.

### 5.4.1. Using XDoclet markup

Many Hibernate users prefer to embed mapping information directly in sourcecode using XDoclet `@hibernate.tags`. We will not cover this approach in this document, since strictly it is considered part of XDoclet. However, we include the following example of the `Cat` class with XDoclet mappings.

```
package eg;
import java.util.Set;
import java.util.Date;

/**
 * @hibernate.class
 * table="CATS"
 */
public class Cat {
    private Long id; // identifier
    private Date birthdate;
    private Cat mother;
    private Set kittens;
    private Color color;
    private char sex;
    private float weight;

    /**
     * @hibernate.id
     * generator-class="native"
     * column="CAT_ID"
     */
    public Long getId() {
        return id;
    }
    private void setId(Long id) {
        this.id=id;
    }

    /**
     * @hibernate.many-to-one
     * column="PARENT_ID"
     */
    public Cat getMother() {
        return mother;
    }
    void setMother(Cat mother) {
        this.mother = mother;
    }

    /**
     * @hibernate.property
     * column="BIRTH_DATE"
     */
    public Date getBirthdate() {
        return birthdate;
    }
    void setBirthdate(Date date) {
        birthdate = date;
    }

    /**
     * @hibernate.property
     * column="WEIGHT"
     */
    public float getWeight() {
        return weight;
    }
    void setWeight(float weight) {
        this.weight = weight;
    }

    /**
     * @hibernate.property
     * column="COLOR"
     * not-null="true"
     */
}
```

```

    public Color getColor() {
        return color;
    }
    void setColor(Color color) {
        this.color = color;
    }
    /**
     * @hibernate.set
     * inverse="true"
     * order-by="BIRTH_DATE"
     * @hibernate.collection-key
     * column="PARENT_ID"
     * @hibernate.collection-one-to-many
     */
    public Set getKittens() {
        return kittens;
    }
    void setKittens(Set kittens) {
        this.kittens = kittens;
    }
    // addKitten not needed by Hibernate
    public void addKitten(Cat kitten) {
        kittens.add(kitten);
    }

    /**
     * @hibernate.property
     * column="SEX"
     * not-null="true"
     * update="false"
     */
    public char getSex() {
        return sex;
    }
    void setSex(char sex) {
        this.sex=sex;
    }
}

```

See the Hibernate web site for more examples of XDoclet and Hibernate.

### 5.4.2. Using JDK 5.0 Annotations

JDK 5.0 introduced XDoclet-style annotations at the language level, type-safe and checked at compile time. This mechanism is more powerful than XDoclet annotations and better supported by tools and IDEs. IntelliJ IDEA, for example, supports auto-completion and syntax highlighting of JDK 5.0 annotations. The new revision of the EJB specification (JSR-220) uses JDK 5.0 annotations as the primary metadata mechanism for entity beans. Hibernate3 implements the `EntityManager` of JSR-220 (the persistence API), support for mapping metadata is available via the *Hibernate Annotations* package, as a separate download. Both EJB3 (JSR-220) and Hibernate3 metadata is supported.

This is an example of a POJO class annotated as an EJB entity bean:

```

@Entity(access = AccessType.FIELD)
public class Customer implements Serializable {

    @Id;
    Long id;

    String firstName;
    String lastName;
    Date birthday;

    @Transient
    Integer age;
}

```



```
@Dependent
private Address homeAddress;

@OneToMany(cascade=CascadeType.ALL,
           targetEntity="Order")
@JoinColumn(name="CUSTOMER_ID")
Set orders;

// Getter/setter and business methods
}
```

Note that support for JDK 5.0 Annotations (and JSR-220) is still work in progress and not completed.

---

# Chapter 6. Collection Mapping

## 6.1. Persistent collections

Hibernate requires that persistent collection-valued fields be declared as an interface type, for example:

```
public class Product {
    private String serialNumber;
    private Set parts = new HashSet();

    public Set getParts() { return parts; }
    void setParts(Set parts) { this.parts = parts; }
    public String getSerialNumber() { return serialNumber; }
    void setSerialNumber(String sn) { serialNumber = sn; }
}
```

The actual interface might be `java.util.Set`, `java.util.Collection`, `java.util.List`, `java.util.Map`, `java.util.SortedSet`, `java.util.SortedMap` or ... anything you like! (Where "anything you like" means you will have to write an implementation of `org.hibernate.usertype.UserCollectionType`.)

Notice how we initialized the instance variable with an instance of `HashSet`. This is the best way to initialize collection valued properties of newly instantiated (non-persistent) instances. When you make the instance persistent - by calling `persist()`, for example - Hibernate will actually replace the `HashSet` with an instance of Hibernate's own implementation of `Set`. Watch out for errors like this:

```
Cat cat = new DomesticCat();
Cat kitten = new DomesticCat();
....
Set kittens = new HashSet();
kittens.add(kitten);
cat.setKittens(kittens);
session.persist(cat);
kittens = cat.getKittens(); // Okay, kittens collection is a Set
(HashSet) cat.getKittens(); // Error!
```

The persistent collections injected by Hibernate behave like `HashMap`, `HashSet`, `TreeMap`, `TreeSet` or `ArrayList`, depending upon the interface type.

Collections instances have the usual behavior of value types. They are automatically persisted when referenced by a persistent object and automatically deleted when unreferenced. If a collection is passed from one persistent object to another, its elements might be moved from one table to another. Two entities may not share a reference to the same collection instance. Due to the underlying relational model, collection-valued properties do not support null value semantics; Hibernate does not distinguish between a null collection reference and an empty collection.

You shouldn't have to worry much about any of this. Use persistent collections the same way you use ordinary Java collections. Just make sure you understand the semantics of bidirectional associations (discussed later).

## 6.2. Collection mappings

The Hibernate mapping element used for mapping a collection depends upon the type of the interface. For example, a `<set>` element is used for mapping properties of type `Set`.

```
<class name="Product">
    <id name="serialNumber" column="productSerialNumber"/>
```

```

    <set name="parts">
      <key column="productSerialNumber" not-null="true"/>
      <one-to-many class="Part"/>
    </set>
  </class>

```

Apart from `<set>`, there is also `<list>`, `<map>`, `<bag>`, `<array>` and `<primitive-array>` mapping elements. The `<map>` element is representative:

```

<map
  name="propertyName"                                (1)
  table="table_name"                                  (2)
  schema="schema_name"                               (3)
  lazy="true|false"                                   (4)
  inverse="true|false"                                (5)
  cascade="all|none|save-update|delete|all-delete-orphan" (6)
  sort="unsorted|natural|comparatorClass"             (7)
  order-by="column_name asc|desc"                     (8)
  where="arbitrary sql where condition"               (9)
  fetch="join|select|subselect"                      (10)
  batch-size="N"                                       (11)
  access="field|property|ClassName"                   (12)
  optimistic-lock="true|false"                       (13)
>
  <key .... />
  <map-key .... />
  <element .... />
</map>

```

- (1) name the collection property name
- (2) table (optional - defaults to property name) the name of the collection table (not used for one-to-many associations)
- (3) schema (optional) the name of a table schema to override the schema declared on the root element
- (4) lazy (optional - defaults to true) enable lazy initialization (not available for arrays)
- (5) inverse (optional - defaults to false) mark this collection as the "inverse" end of a bidirectional association
- (6) cascade (optional - defaults to none) enable operations to cascade to child entities
- (7) sort (optional) specify a sorted collection with `natural` sort order, or a given comparator class
- (8) order-by (optional, JDK1.4 only) specify a table column (or columns) that define the iteration order of the Map, Set or bag, together with an optional `asc` or `desc`
- (9) where (optional) specify an arbitrary SQL WHERE condition to be used when retrieving or removing the collection (useful if the collection should contain only a subset of the available data)
- (10) fetch (optional, defaults to `select`) Choose between outer-join fetching, fetching by sequential select, and fetching by sequential subselect. Only one collection may be fetched by outer join per SQL SELECT.
- (11) batch-size (optional, defaults to 1) specify a "batch size" for lazily fetching instances of this collection.
- (12) access (optional - defaults to `property`): The strategy Hibernate should use for accessing the property value.
- (12) optimistic-lock (optional - defaults to `true`): Species that changes to the state of the collection results in increment of the owning entity's version. (For one to many associations, it is often reasonable to disable this setting.)

### 6.2.1. Collection foreign keys

Collection instances are distinguished in the database by the foreign key of the entity that owns the collection. This foreign key is referred to as the *collection key column* (or columns) of the collection table. The collection key column is mapped by the `<key>` element.

There may be a nullability constraint on the foreign key column. For most collections, this is implied. For uni-directional one to many associations, the foreign key column is nullable by default, so you might need to specify `not-null="true"`.

```
<key column="productSerialNumber" not-null="true"/>
```

The foreign key constraint may use `ON DELETE CASCADE`.

```
<key column="productSerialNumber" on-delete="cascade"/>
```

See the previous chapter for a full definition of the `<key>` element.

## 6.2.2. Collection elements

Collections may contain almost any other Hibernate type, including all basic types, custom types, components, and of course, references to other entities. This is an important distinction: an object in a collection might be handled with "value" semantics (its lifecycle fully depends on the collection owner) or it might be a reference to another entity, with its own lifecycle. In the latter case, only the "link" between the two objects is considered to be state held by the collection.

The contained type is referred to as the *collection element type*. Collection elements are mapped by `<element>` or `<composite-element>`, or in the case of entity references, with `<one-to-many>` or `<many-to-many>`. The first two map elements with value semantics, the next two are used to map entity associations.

## 6.2.3. Indexed collections

All collection mappings, except those with set and bag semantics, need an *index column* in the collection table - a column that maps to an array index, or `List` index, or `Map` key. The index of a `Map` may be of any basic type, mapped with `<map-key>`, it may be an entity reference mapped with `<map-key-many-to-many>`, or it may be a composite type, mapped with `<composite-map-key>`. The index of an array or list is always of type `integer` and is mapped using the `<list-index>` element. The mapped column contains sequential integers (numbered from zero, by default).

```
<list-index
  column="column_name"                (1)
  base="0|1|..." />
```

- (1) `column_name` (required): The name of the column holding the collection index values.
- (1) `base` (optional, defaults to 0): The value of the index column that corresponds to the first element of the list or array.

```
<map-key
  column="column_name"                (1)
  formula="any SQL expression"        (2)
  type="type_name"                    (3)
  length="N" />
```

- (1) `column` (optional): The name of the column holding the collection index values.
- (2) `formula` (optional): A SQL formula used to evaluate the key of the map.
- (3) `type` (optional, defaults to `integer`): The type of the collection index.

```
<map-key-many-to-many
  column="column_name"                (1)
  formula="any SQL expression"        (2)(3)
  class="ClassName" />
```

```
</>
```

- (1) `column` (optional): The name of the foreign key column for the collection index values.
- (2) `formula` (optional): A SQL formula used to evaluate the foreign key of the map key.
- (3) `class` (required): The entity class used as the collection index.

If your table doesn't have an index column, and you still wish to use `List` as the property type, you should map the property as a Hibernate `<bag>`. A bag does not retain its order when it is retrieved from the database, but it may be optionally sorted or ordered.

There are quite a range of mappings that can be generated for collections, covering many common relational models. We suggest you experiment with the schema generation tool to get a feeling for how various mapping declarations translate to database tables.

## 6.2.4. Collections of values and many-to-many associations

Any collection of values or many-to-many association requires a dedicated *collection table* with a foreign key column or columns, *collection element column* or columns and possibly an index column or columns.

For a collection of values, we use the `<element>` tag.

```
<element
    column="column_name"           (1)
    formula="any SQL expression"  (2)
    type="typename"               (3)
    length="N"
    precision="N"
    scale="N"
    not-null="true|false"
    unique="true|false"
/>
```

- (1) `column` (optional): The name of the column holding the collection element values.
- (2) `formula` (optional): An SQL formula used to evaluate the element.
- (3) `type` (required): The type of the collection element.

A *many-to-many association* is specified using the `<many-to-many>` element.

```
<many-to-many
    column="column_name"           (1)
    formula="any SQL expression"  (2)
    class="ClassName"             (3)
    fetch="select|join"           (4)
    unique="true|false"           (5)
    entity-name="EntityName"
/>
```

- (1) `column` (optional): The name of the element foreign key column.
- (2) `formula` (optional): An SQL formula used to evaluate the element foreign key value.
- (3) `class` (required): The name of the associated class.
- (4) `fetch` (optional - defaults to `join`): enables outer-join or sequential select fetching for this association. This is a special case; for full eager fetching (in a single `SELECT`) of an entity and its many-to-many relationships to other entities, you would enable `join` fetching not only of the collection itself, but also with this attribute on the `<many-to-many>` nested element.
- (5) `unique` (optional): Enable the DDL generation of a unique constraint for the foreign-key column. This makes the association multiplicity effectively one to many.

Some examples, first, a set of strings:

```
<set name="names" table="person_names">
  <key column="person_id"/>
  <element column="person_name" type="string"/>
</set>
```

A bag containing integers (with an iteration order determined by the `order-by` attribute):

```
<bag name="sizes"
      table="item_sizes"
      order-by="size asc">
  <key column="item_id"/>
  <element column="size" type="integer"/>
</bag>
```

An array of entities - in this case, a many to many association:

```
<array name="addresses"
        table="PersonAddress"
        cascade="create">
  <key column="personId"/>
  <list-index column="sortOrder"/>
  <many-to-many column="addressId" class="Address"/>
</array>
```

A map from string indices to dates:

```
<map name="holidays"
      table="holidays"
      schema="dbo"
      order-by="hol_name asc">
  <key column="id"/>
  <map-key column="hol_name" type="string"/>
  <element column="hol_date" type="date"/>
</map>
```

A list of components (discussed in the next chapter):

```
<list name="carComponents"
      table="CarComponents">
  <key column="carId"/>
  <list-index column="sortOrder"/>
  <composite-element class="CarComponent">
    <property name="price"/>
    <property name="type"/>
    <property name="serialNumber" column="serialNum"/>
  </composite-element>
</list>
```

## 6.2.5. One-to-many associations

A *one to many association* links the tables of two classes via a foreign key, with no intervening collection table. This mapping loses certain semantics of normal Java collections:

- An instance of the contained entity class may not belong to more than one instance of the collection
- An instance of the contained entity class may not appear at more than one value of the collection index

An association from `Product` to `Part` requires existence of a foreign key column and possibly an index column to the `Part` table. A `<one-to-many>` tag indicates that this is a one to many association.

```
<one-to-many                                (1)
    class="ClassName"
    entity-name="EntityName" />
```

(1) `class` (required): The name of the associated class.

Notice that the `<one-to-many>` element does not need to declare any columns. Nor is it necessary to specify the table name anywhere.

*Very important note:* If the foreign key column of a `<one-to-many>` association is declared `NOT NULL`, you must declare the `<key>` mapping `not-null="true"` or use a *bidirectional association* with the collection mapping marked `inverse="true"`. See the discussion of bidirectional associations later in this chapter.

This example shows a map of `Part` entities by name (where `partName` is a persistent property of `Part`). Notice the use of a formula-based index.

```
<map name="parts"
    cascade="all">
    <key column="productId" not-null="true" />
    <map-key formula="partName" />
    <one-to-many class="Part" />
</map>
```

## 6.3. Advanced collection mappings

### 6.3.1. Sorted collections

Hibernate supports collections implementing `java.util.SortedMap` and `java.util.SortedSet`. You must specify a comparator in the mapping file:

```
<set name="aliases"
    table="person_aliases"
    sort="natural">
    <key column="person" />
    <element column="name" type="string" />
</set>

<map name="holidays" sort="my.custom.HolidayComparator">
    <key column="year_id" />
    <map-key column="hol_name" type="string" />
    <element column="hol_date" type="date" />
</map>
```

Allowed values of the `sort` attribute are `unsorted`, `natural` and the name of a class implementing `java.util.Comparator`.

Sorted collections actually behave like `java.util.TreeSet` or `java.util.TreeMap`.

If you want the database itself to order the collection elements use the `order-by` attribute of `set`, `bag` or `map` mappings. This solution is only available under JDK 1.4 or higher (it is implemented using `LinkedHashSet` or `LinkedHashMap`). This performs the ordering in the SQL query, not in memory.

```
<set name="aliases" table="person_aliases" order-by="lower(name) asc">
    <key column="person" />
    <element column="name" type="string" />
</set>

<map name="holidays" order-by="hol_date, hol_name">
```

```

<key column="year_id" />
<map-key column="hol_name" type="string" />
<element column="hol_date" type="date" />
</map>

```

Note that the value of the `order-by` attribute is an SQL ordering, not a HQL ordering!

Associations may even be sorted by some arbitrary criteria at runtime using a `collection filter()`.

```
sortedUsers = s.createFilter( group.getUsers(), "order by this.name" ).list();
```

### 6.3.2. Bidirectional associations

A *bidirectional association* allows navigation from both "ends" of the association. Two kinds of bidirectional association are supported:

one-to-many

set or bag valued at one end, single-valued at the other

many-to-many

set or bag valued at both ends

You may specify a bidirectional many-to-many association simply by mapping two many-to-many associations to the same database table and declaring one end as *inverse* (which one is your choice, but it can not be an indexed collection).

Here's an example of a bidirectional many-to-many association; each category can have many items and each item can be in many categories:

```

<class name="Category">
  <id name="id" column="CATEGORY_ID" />
  ...
  <bag name="items" table="CATEGORY_ITEM">
    <key column="CATEGORY_ID" />
    <many-to-many class="Item" column="ITEM_ID" />
  </bag>
</class>

<class name="Item">
  <id name="id" column="CATEGORY_ID" />
  ...

  <!-- inverse end -->
  <bag name="categories" table="CATEGORY_ITEM" inverse="true">
    <key column="ITEM_ID" />
    <many-to-many class="Category" column="CATEGORY_ID" />
  </bag>
</class>

```

Changes made only to the inverse end of the association are *not* persisted. This means that Hibernate has two representations in memory for every bidirectional association, one link from A to B and another link from B to A. This is easier to understand if you think about the Java object model and how we create a many-to-many relationship in Java:

```

category.getItems().add(item);           // The category now "knows" about the relationship
item.getCategories().add(category);       // The item now "knows" about the relationship

session.persist(item);                    // The relationship won't be saved!

```



```
session.persist(category); // The relationship will be saved
```

The non-inverse side is used to save the in-memory representation to the database.

You may define a bidirectional one-to-many association by mapping a one-to-many association to the same table column(s) as a many-to-one association and declaring the many-valued end `inverse="true"`.

```
<class name="Parent">
  <id name="id" column="parent_id"/>
  ....
  <set name="children" inverse="true">
    <key column="parent_id"/>
    <one-to-many class="Child"/>
  </set>
</class>

<class name="eg.Child">
  <id name="id" column="id"/>
  ....
  <many-to-one name="parent"
    class="Parent"
    column="parent_id"
    not-null="true"/>
</class>
```

Mapping one end of an association with `inverse="true"` doesn't affect the operation of cascades, these are orthogonal concepts!

### 6.3.3. Ternary associations

There are three possible approaches to mapping a ternary association. One is to use a `Map` with an association as its index:

```
<map name="contracts">
  <key column="employer_id" not-null="true"/>
  <map-key-many-to-many column="employee_id" class="Employee"/>
  <one-to-many class="Contract"/>
</map>
```

```
<map name="connections">
  <key column="incoming_node_id"/>
  <map-key-many-to-many column="outgoing_node_id" class="Node"/>
  <many-to-many column="connection_id" class="Connection"/>
</map>
```

A second approach is to simply remodel the association as an entity class. This is the approach we use most commonly.

A final alternative is to use composite elements, which we will discuss later.

### 6.3.4. Using an `<idbag>`

If you've fully embraced our view that composite keys are a bad thing and that entities should have synthetic identifiers (surrogate keys), then you might find it a bit odd that the many to many associations and collections of values that we've shown so far all map to tables with composite keys! Now, this point is quite arguable; a pure association table doesn't seem to benefit much from a surrogate key (though a collection of composite values *might*). Nevertheless, Hibernate provides a feature that allows you to map many to many associations and collections of values to a table with a surrogate key.

The `<idbag>` element lets you map a `List` (or `Collection`) with bag semantics.

```
<idbag name="lovers" table="LOVERS">
  <collection-id column="ID" type="long">
    <generator class="sequence"/>
  </collection-id>
  <key column="PERSON1"/>
  <many-to-many column="PERSON2" class="eg.Person" outer-join="true"/>
</idbag>
```

As you can see, an `<idbag>` has a synthetic id generator, just like an entity class! A different surrogate key is assigned to each collection row. Hibernate does not provide any mechanism to discover the surrogate key value of a particular row, however.

Note that the update performance of an `<idbag>` is *much* better than a regular `<bag>`! Hibernate can locate individual rows efficiently and update or delete them individually, just like a list, map or set.

In the current implementation, the native identifier generation strategy is not supported for `<idbag>` collection identifiers.

## 6.4. Collection examples

The previous sections are pretty confusing. So let's look at an example. This class:

```
package eg;
import java.util.Set;

public class Parent {
    private long id;
    private Set children;

    public long getId() { return id; }
    private void setId(long id) { this.id=id; }

    private Set getChildren() { return children; }
    private void setChildren(Set children) { this.children=children; }

    ....
    ....
}
```

has a collection of `Child` instances. If each child has at most one parent, the most natural mapping is a one-to-many association:

```
<hibernate-mapping>

  <class name="Parent">
    <id name="id">
      <generator class="sequence"/>
    </id>
    <set name="children">
      <key column="parent_id"/>
      <one-to-many class="Child"/>
    </set>
  </class>

  <class name="Child">
    <id name="id">
      <generator class="sequence"/>
    </id>
    <property name="name"/>
  </class>
```

```
</hibernate-mapping>
```

This maps to the following table definitions:

```
create table parent ( id bigint not null primary key )
create table child ( id bigint not null primary key, name varchar(255), parent_id bigint )
alter table child add constraint childfk0 (parent_id) references parent
```

If the parent is *required*, use a bidirectional one-to-many association:

```
<hibernate-mapping>

  <class name="Parent">
    <id name="id">
      <generator class="sequence"/>
    </id>
    <set name="children" inverse="true">
      <key column="parent_id"/>
      <one-to-many class="Child"/>
    </set>
  </class>

  <class name="Child">
    <id name="id">
      <generator class="sequence"/>
    </id>
    <property name="name"/>
    <many-to-one name="parent" class="Parent" column="parent_id" not-null="true"/>
  </class>

</hibernate-mapping>
```

Notice the NOT NULL constraint:

```
create table parent ( id bigint not null primary key )
create table child ( id bigint not null
                    primary key,
                    name varchar(255),
                    parent_id bigint not null )
alter table child add constraint childfk0 (parent_id) references parent
```

Alternatively, if you absolutely insist that this association should be unidirectional, you can declare the NOT NULL constraint on the <key> mapping:

```
<hibernate-mapping>

  <class name="Parent">
    <id name="id">
      <generator class="sequence"/>
    </id>
    <set name="children">
      <key column="parent_id" not-null="true"/>
      <one-to-many class="Child"/>
    </set>
  </class>

  <class name="Child">
    <id name="id">
      <generator class="sequence"/>
    </id>
    <property name="name"/>
  </class>

</hibernate-mapping>
```

On the other hand, if a child might have multiple parents, a many-to-many association is appropriate:

```
<hibernate-mapping>

  <class name="Parent">
    <id name="id">
      <generator class="sequence"/>
    </id>
    <set name="children" table="childset">
      <key column="parent_id"/>
      <many-to-many class="Child" column="child_id"/>
    </set>
  </class>

  <class name="Child">
    <id name="id">
      <generator class="sequence"/>
    </id>
    <property name="name"/>
  </class>

</hibernate-mapping>
```

Table definitions:

```
create table parent ( id bigint not null primary key )
create table child ( id bigint not null primary key, name varchar(255) )
create table childset ( parent_id bigint not null,
                        child_id bigint not null,
                        primary key ( parent_id, child_id ) )
alter table childset add constraint childsetfk0 (parent_id) references parent
alter table childset add constraint childsetfk1 (child_id) references child
```

For more examples and a complete walk-through a parent/child relationship mapping, see Chapter 21, *Example: Parent/Child*.

Even more exotic association mappings are possible, we will catalog all possibilities in the next chapter.

---

# Chapter 7. Association Mappings

## 7.1. Introduction

Association mappings are the often most difficult thing to get right. In this section we'll go through the canonical cases one by one, starting with unidirectional mappings, and then considering the bidirectional cases. We'll use `Person` and `Address` in all the examples.

We'll classify associations by whether or not they map to an intervening join table, and by multiplicity.

Nullable foreign keys are not considered good practice in traditional data modelling, so all our examples use not null foreign keys. This is not a requirement of Hibernate, and the mappings will all work if you drop the nullability constraints.

## 7.2. Unidirectional associations

### 7.2.1. many to one

A *unidirectional many-to-one association* is the most common kind of unidirectional association.

```
<class name="Person">
  <id name="id" column="personId">
    <generator class="native"/>
  </id>
  <many-to-one name="address"
    column="addressId"
    not-null="true"/>
</class>

<class name="Address">
  <id name="id" column="addressId">
    <generator class="native"/>
  </id>
</class>
```

```
create table Person ( personId bigint not null primary key, addressId bigint not null )
create table Address ( addressId bigint not null primary key )
```

### 7.2.2. one to one

A *unidirectional one-to-one association on a foreign key* is almost identical. The only difference is the column unique constraint.

```
<class name="Person">
  <id name="id" column="personId">
    <generator class="native"/>
  </id>
  <many-to-one name="address"
    column="addressId"
    unique="true"
    not-null="true"/>
</class>
```

```
<class name="Address">
  <id name="id" column="addressId">
    <generator class="native"/>
  </id>
</class>
```

```
create table Person ( personId bigint not null primary key, addressId bigint not null unique )
create table Address ( addressId bigint not null primary key )
```

A *unidirectional one-to-one association on a primary key* usually uses a special id generator. (Notice that we've reversed the direction of the association in this example.)

```
<class name="Person">
  <id name="id" column="personId">
    <generator class="native"/>
  </id>
</class>

<class name="Address">
  <id name="id" column="personId">
    <generator class="foreign">
      <param name="property">person</param>
    </generator>
  </id>
  <one-to-one name="person" constrained="true"/>
</class>
```

```
create table Person ( personId bigint not null primary key )
create table Address ( personId bigint not null primary key )
```

### 7.2.3. one to many

A *unidirectional one-to-many association on a foreign key* is a very unusual case, and is not really recommended.

```
<class name="Person">
  <id name="id" column="personId">
    <generator class="native"/>
  </id>
  <set name="addresses">
    <key column="personId"
      not-null="true"/>
    <one-to-many class="Address"/>
  </set>
</class>

<class name="Address">
  <id name="id" column="addressId">
    <generator class="native"/>
  </id>
</class>
```

```
create table Person ( personId bigint not null primary key )
create table Address ( addressId bigint not null primary key, personId bigint not null )
```

We think it's better to use a join table for this kind of association.

## 7.3. Unidirectional associations with join tables

### 7.3.1. one to many

A *unidirectional one-to-many association on a join table* is much preferred. Notice that by specifying `unique="true"`, we have changed the multiplicity from many-to-many to one-to-many.

```
<class name="Person">
  <id name="id" column="personId">
    <generator class="native"/>
  </id>
  <set name="addresses" table="PersonAddress">
    <key column="personId"/>
    <many-to-many column="addressId"
      unique="true"
      class="Address"/>
  </set>
</class>

<class name="Address">
  <id name="id" column="addressId">
    <generator class="native"/>
  </id>
</class>
```

```
create table Person ( personId bigint not null primary key )
create table PersonAddress ( personId not null, addressId bigint not null primary key )
create table Address ( addressId bigint not null primary key )
```

### 7.3.2. many to one

A *unidirectional many-to-one association on a join table* is quite common when the association is optional.

```
<class name="Person">
  <id name="id" column="personId">
    <generator class="native"/>
  </id>
  <join table="PersonAddress"
    optional="true">
    <key column="personId" unique="true"/>
    <many-to-one name="address"
      column="addressId"
      not-null="true"/>
  </join>
</class>

<class name="Address">
  <id name="id" column="addressId">
    <generator class="native"/>
  </id>
</class>
```

```
create table Person ( personId bigint not null primary key )
create table PersonAddress ( personId bigint not null primary key, addressId bigint not null )
create table Address ( addressId bigint not null primary key )
```

### 7.3.3. one to one

A *unidirectional one-to-one association on a join table* is extremely unusual, but possible.

```
<class name="Person">
  <id name="id" column="personId">
    <generator class="native"/>
  </id>
  <join table="PersonAddress"
    optional="true">
    <key column="personId"
      unique="true"/>
    <many-to-one name="address"
      column="addressId"
      not-null="true"
      unique="true"/>
  </join>
</class>

<class name="Address">
  <id name="id" column="addressId">
    <generator class="native"/>
  </id>
</class>
```

```
create table Person ( personId bigint not null primary key )
create table PersonAddress ( personId bigint not null primary key, addressId bigint not null unique )
create table Address ( addressId bigint not null primary key )
```

### 7.3.4. many to many

Finally, we have a *unidirectional many-to-many association*.

```
<class name="Person">
  <id name="id" column="personId">
    <generator class="native"/>
  </id>
  <set name="addresses" table="PersonAddress">
    <key column="personId"/>
    <many-to-many column="addressId"
      class="Address"/>
  </set>
</class>

<class name="Address">
  <id name="id" column="addressId">
    <generator class="native"/>
  </id>
</class>
```

```
create table Person ( personId bigint not null primary key )
create table PersonAddress ( personId bigint not null, addressId bigint not null, primary key (personId, addressId) )
create table Address ( addressId bigint not null primary key )
```

## 7.4. Bidirectional associations



### 7.4.1. one to many / many to one

A *bidirectional many-to-one association* is the most common kind of association. (This is the standard parent/child relationship.)

```
<class name="Person">
  <id name="id" column="personId">
    <generator class="native"/>
  </id>
  <many-to-one name="address"
    column="addressId"
    not-null="true"/>
</class>

<class name="Address">
  <id name="id" column="addressId">
    <generator class="native"/>
  </id>
  <set name="people" inverse="true">
    <key column="addressId"/>
    <one-to-many class="Person"/>
  </set>
</class>
```

```
create table Person ( personId bigint not null primary key, addressId bigint not null )
create table Address ( addressId bigint not null primary key )
```

### 7.4.2. one to one

A *bidirectional one-to-one association on a foreign key* is quite common.

```
<class name="Person">
  <id name="id" column="personId">
    <generator class="native"/>
  </id>
  <many-to-one name="address"
    column="addressId"
    unique="true"
    not-null="true"/>
</class>

<class name="Address">
  <id name="id" column="addressId">
    <generator class="native"/>
  </id>
  <one-to-one name="person"
    property-ref="address"/>
</class>
```

```
create table Person ( personId bigint not null primary key, addressId bigint not null unique )
create table Address ( addressId bigint not null primary key )
```

A *bidirectional one-to-one association on a primary key* uses the special id generator.

```
<class name="Person">
  <id name="id" column="personId">
    <generator class="native"/>
  </id>
  <one-to-one name="address"/>
</class>
```

```
<class name="Address">
  <id name="id" column="personId">
    <generator class="foreign">
      <param name="property">person</param>
    </generator>
  </id>
  <one-to-one name="person"
    constrained="true"/>
</class>
```

```
create table Person ( personId bigint not null primary key )
create table Address ( personId bigint not null primary key )
```

## 7.5. Bidirectional associations with join tables

### 7.5.1. one to many / many to one

A *bidirectional one-to-many association on a join table*. Note that the `inverse="true"` can go on either end of the association, on the collection, or on the join.

```
<class name="Person">
  <id name="id" column="personId">
    <generator class="native"/>
  </id>
  <set name="addresses"
    table="PersonAddress">
    <key column="personId"/>
    <many-to-many column="addressId"
      unique="true"
      class="Address"/>
  </set>
</class>

<class name="Address">
  <id name="id" column="addressId">
    <generator class="native"/>
  </id>
  <join table="PersonAddress"
    inverse="true"
    optional="true">
    <key column="addressId"/>
    <many-to-one name="person"
      column="personId"
      not-null="true"/>
  </join>
</class>
```

```
create table Person ( personId bigint not null primary key )
create table PersonAddress ( personId bigint not null, addressId bigint not null primary key )
create table Address ( addressId bigint not null primary key )
```

### 7.5.2. one to one

A *bidirectional one-to-one association on a join table* is extremely unusual, but possible.

```
<class name="Person">
```

```

<id name="id" column="personId">
  <generator class="native" />
</id>
<join table="PersonAddress"
  optional="true">
  <key column="personId"
    unique="true" />
  <many-to-one name="address"
    column="addressId"
    not-null="true"
    unique="true" />
  </join>
</class>

<class name="Address">
  <id name="id" column="addressId">
    <generator class="native" />
  </id>
  <join table="PersonAddress"
    optional="true"
    inverse="true">
    <key column="addressId"
      unique="true" />
    <many-to-one name="address"
      column="personId"
      not-null="true"
      unique="true" />
    </join>
</class>

```

```

create table Person ( personId bigint not null primary key )
create table PersonAddress ( personId bigint not null primary key, addressId bigint not null unique )
create table Address ( addressId bigint not null primary key )

```

### 7.5.3. many to many

Finally, we have a *bidirectional many-to-many association*.

```

<class name="Person">
  <id name="id" column="personId">
    <generator class="native" />
  </id>
  <set name="addresses">
    <key column="personId" />
    <many-to-many column="addressId"
      class="Address" />
  </set>
</class>

<class name="Address">
  <id name="id" column="addressId">
    <generator class="native" />
  </id>
  <set name="people" inverse="true">
    <key column="addressId" />
    <many-to-many column="personId"
      class="Person" />
  </set>
</class>

```

```

create table Person ( personId bigint not null primary key )
create table PersonAddress ( personId bigint not null, addressId bigint not null, primary key (personId, addressId) )
create table Address ( addressId bigint not null primary key )

```



---

## Chapter 8. Component Mapping

The notion of a *component* is re-used in several different contexts, for different purposes, throughout Hibernate.

### 8.1. Dependent objects

A component is a contained object that is persisted as a value type, not an entity reference. The term "component" refers to the object-oriented notion of composition (not to architecture-level components). For example, you might model a person like this:

```
public class Person {
    private java.util.Date birthday;
    private Name name;
    private String key;
    public String getKey() {
        return key;
    }
    private void setKey(String key) {
        this.key=key;
    }
    public java.util.Date getBirthday() {
        return birthday;
    }
    public void setBirthday(java.util.Date birthday) {
        this.birthday = birthday;
    }
    public Name getName() {
        return name;
    }
    public void setName(Name name) {
        this.name = name;
    }
    .....
    .....
}
```

```
public class Name {
    char initial;
    String first;
    String last;
    public String getFirst() {
        return first;
    }
    void setFirst(String first) {
        this.first = first;
    }
    public String getLast() {
        return last;
    }
    void setLast(String last) {
        this.last = last;
    }
    public char getInitial() {
        return initial;
    }
    void setInitial(char initial) {
        this.initial = initial;
    }
}
```

Now `Name` may be persisted as a component of `Person`. Notice that `Name` defines getter and setter methods for its persistent properties, but doesn't need to declare any interfaces or identifier properties.

Our Hibernate mapping would look like:

```
<class name="eg.Person" table="person">
  <id name="Key" column="pid" type="string">
    <generator class="uuid.hex"/>
  </id>
  <property name="birthday" type="date"/>
  <component name="Name" class="eg.Name"> <!-- class attribute optional -->
    <property name="initial"/>
    <property name="first"/>
    <property name="last"/>
  </component>
</class>
```

The person table would have the columns `pid`, `birthday`, `initial`, `first` and `last`.

Like all value types, components do not support shared references. In other words, two persons could have the same name, but the two person objects would contain two independent name objects, only "the same" by value. The null value semantics of a component are *ad hoc*. When reloading the containing object, Hibernate will assume that if all component columns are null, then the entire component is null. This should be okay for most purposes.

The properties of a component may be of any Hibernate type (collections, many-to-one associations, other components, etc). Nested components should *not* be considered an exotic usage. Hibernate is intended to support a very fine-grained object model.

The `<component>` element allows a `<parent>` subelement that maps a property of the component class as a reference back to the containing entity.

```
<class name="eg.Person" table="person">
  <id name="Key" column="pid" type="string">
    <generator class="uuid.hex"/>
  </id>
  <property name="birthday" type="date"/>
  <component name="Name" class="eg.Name" unique="true">
    <parent name="namedPerson"/> <!-- reference back to the Person -->
    <property name="initial"/>
    <property name="first"/>
    <property name="last"/>
  </component>
</class>
```

## 8.2. Collections of dependent objects

Collections of components are supported (eg. an array of type `Name`). Declare your component collection by replacing the `<element>` tag with a `<composite-element>` tag.

```
<set name="someNames" table="some_names" lazy="true">
  <key column="id"/>
  <composite-element class="eg.Name"> <!-- class attribute required -->
    <property name="initial"/>
    <property name="first"/>
    <property name="last"/>
  </composite-element>
</set>
```

Note: if you define a Set of composite elements, it is very important to implement `equals()` and `hashCode()` correctly.

Composite elements may contain components but not collections. If your composite element itself contains components, use the `<nested-composite-element>` tag. This is a pretty exotic case - a collection of components which themselves have components. By this stage you should be asking yourself if a one-to-many association is more appropriate. Try remodelling the composite element as an entity - but note that even though the Java model is the same, the relational model and persistence semantics are still slightly different.

Please note that a composite element mapping doesn't support null-able properties if you're using a `<set>`. Hibernate has to use each column's value to identify a record when deleting objects (there is no separate primary key column in the composite element table), which is not possible with null values. You have to either use only not-null properties in a composite-element or choose a `<list>`, `<map>`, `<bag>` or `<idbag>`.

A special case of a composite element is a composite element with a nested `<many-to-one>` element. A mapping like this allows you to map extra columns of a many-to-many association table to the composite element class. The following is a many-to-many association from `Order` to `Item` where `purchaseDate`, `price` and `quantity` are properties of the association:

```
<class name="eg.Order" .... >
  ....
  <set name="purchasedItems" table="purchase_items" lazy="true">
    <key column="order_id">
      <composite-element class="eg.Purchase">
        <property name="purchaseDate"/>
        <property name="price"/>
        <property name="quantity"/>
        <many-to-one name="item" class="eg.Item"/> <!-- class attribute is optional -->
      </composite-element>
    </set>
  </class>
```

Of course, there can't be a reference to the purchase on the other side, for bidirectional association navigation. Remember that components are value types and don't allow shared references. A single `Purchase` can be in the set of an `Order`, but it can't be referenced by the `Item` at the same time.

Even ternary (or quaternary, etc) associations are possible:

```
<class name="eg.Order" .... >
  ....
  <set name="purchasedItems" table="purchase_items" lazy="true">
    <key column="order_id">
      <composite-element class="eg.OrderLine">
        <many-to-one name="purchaseDetails" class="eg.Purchase"/>
        <many-to-one name="item" class="eg.Item"/>
      </composite-element>
    </set>
  </class>
```

Composite elements may appear in queries using the same syntax as associations to other entities.

## 8.3. Components as Map indices

The `<composite-map-key>` element lets you map a component class as the key of a `Map`. Make sure you override `hashCode()` and `equals()` correctly on the component class.

## 8.4. Components as composite identifiers

You may use a component as an identifier of an entity class. Your component class must satisfy certain require-

ments:

- It must implement `java.io.Serializable`.
- It must re-implement `equals()` and `hashCode()`, consistently with the database's notion of composite key equality.

*Note: in Hibernate3, the second requirement is not an absolutely hard requirement of Hibernate. But do it anyway.*

You can't use an `IdentifierGenerator` to generate composite keys. Instead the application must assign its own identifiers.

Use the `<composite-id>` tag (with nested `<key-property>` elements) in place of the usual `<id>` declaration. For example, the `OrderLine` class has a primary key that depends upon the (composite) primary key of `Order`.

```
<class name="OrderLine">

  <composite-id name="id" class="OrderLineId">
    <key-property name="lineId"/>
    <key-property name="orderId"/>
    <key-property name="customerId"/>
  </composite-id>

  <property name="name"/>

  <many-to-one name="order" class="Order"
    insert="false" update="false">
    <column name="orderId"/>
    <column name="customerId"/>
  </many-to-one>
  ....
</class>
```

Now, any foreign keys referencing the `OrderLine` table are also composite. You must declare this in your mappings for other classes. An association to `OrderLine` would be mapped like this:

```
<many-to-one name="orderLine" class="OrderLine">
<!-- the "class" attribute is optional, as usual -->
  <column name="lineId"/>
  <column name="orderId"/>
  <column name="customerId"/>
</many-to-one>
```

(Note that the `<column>` tag is an alternative to the `column` attribute everywhere.)

A many-to-many association to `OrderLine` also uses the composite foreign key:

```
<set name="undeliveredOrderLines">
  <key column name="warehouseId"/>
  <many-to-many class="OrderLine">
    <column name="lineId"/>
    <column name="orderId"/>
    <column name="customerId"/>
  </many-to-many>
</set>
```

The collection of `OrderLines` in `Order` would use:

```
<set name="orderLines" inverse="true">
  <key>
    <column name="orderId"/>
    <column name="customerId"/>
```



```

    </key>
    <one-to-many class="OrderLine"/>
  </set>

```

(The `<one-to-many>` element, as usual, declares no columns.)

If `OrderLine` itself owns a collection, it also has a composite foreign key.

```

<class name="OrderLine">
  ....
  ....
  <list name="deliveryAttempts">
    <key>    <!-- a collection inherits the composite key type -->
      <column name="lineId"/>
      <column name="orderId"/>
      <column name="customerId"/>
    </key>
    <list-index column="attemptId" base="1"/>
    <composite-element class="DeliveryAttempt">
      ...
    </composite-element>
  </set>
</class>

```

## 8.5. Dynamic components

You may even map a property of type `Map`:

```

<dynamic-component name="userAttributes">
  <property name="foo" column="FOO"/>
  <property name="bar" column="BAR"/>
  <many-to-one name="baz" class="Baz" column="BAZ_ID"/>
</dynamic-component>

```

The semantics of a `<dynamic-component>` mapping are identical to `<component>`. The advantage of this kind of mapping is the ability to determine the actual properties of the bean at deployment time, just by editing the mapping document. Runtime manipulation of the mapping document is also possible, using a DOM parser. Even better, you can access (and change) Hibernate's configuration-time metamodel via the `Configuration` object.

---

# Chapter 9. Inheritance Mapping

## 9.1. The Three Strategies

Hibernate supports the three basic inheritance mapping strategies:

- table per class hierarchy
- table per subclass
- table per concrete class

In addition, Hibernate supports a fourth, slightly different kind of polymorphism:

- implicit polymorphism

It is possible to use different mapping strategies for different branches of the same inheritance hierarchy, and then make use of implicit polymorphism to achieve polymorphism across the whole hierarchy. However, Hibernate does not support mixing `<subclass>`, and `<joined-subclass>` and `<union-subclass>` mappings under the same root `<class>` element. It is possible to mix together the table per hierarchy and table per subclass strategies, under the the same `<class>` element, by combining the `<subclass>` and `<join>` elements (see below).

### 9.1.1. Table per class hierarchy

Suppose we have an interface `Payment`, with implementors `CreditCardPayment`, `CashPayment`, `ChequePayment`. The table per hierarchy mapping would look like:

```
<class name="Payment" table="PAYMENT">
  <id name="id" type="long" column="PAYMENT_ID">
    <generator class="native"/>
  </id>
  <discriminator column="PAYMENT_TYPE" type="string"/>
  <property name="amount" column="AMOUNT"/>
  ...
  <subclass name="CreditCardPayment" discriminator-value="CREDIT">
    <property name="creditCardType" column="CCTYPE"/>
    ...
  </subclass>
  <subclass name="CashPayment" discriminator-value="CASH">
    ...
  </subclass>
  <subclass name="ChequePayment" discriminator-value="CHEQUE">
    ...
  </subclass>
</class>
```

Exactly one table is required. There is one big limitation of this mapping strategy: columns declared by the subclasses, such as `CCTYPE`, may not have `NOT NULL` constraints.

### 9.1.2. Table per subclass

A table per subclass mapping would look like:

```

<class name="Payment" table="PAYMENT">
  <id name="id" type="long" column="PAYMENT_ID">
    <generator class="native"/>
  </id>
  <property name="amount" column="AMOUNT"/>
  ...
  <joined-subclass name="CreditCardPayment" table="CREDIT_PAYMENT">
    <key column="PAYMENT_ID"/>
    <property name="creditCardType" column="CCTYPE"/>
    ...
  </joined-subclass>
  <joined-subclass name="CashPayment" table="CASH_PAYMENT">
    <key column="PAYMENT_ID"/>
    ...
  </joined-subclass>
  <joined-subclass name="ChequePayment" table="CHEQUE_PAYMENT">
    <key column="PAYMENT_ID"/>
    ...
  </joined-subclass>
</class>

```

Four tables are required. The three subclass tables have primary key associations to the superclass table (so the relational model is actually a one-to-one association).

### 9.1.3. Table per subclass, using a discriminator

Note that Hibernate's implementation of table per subclass requires no discriminator column. Other object/relational mappers use a different implementation of table per subclass which requires a type discriminator column in the superclass table. The approach taken by Hibernate is much more difficult to implement but arguably more correct from a relational point of view. If you would like to use a discriminator column with the table per subclass strategy, you may combine the use of `<subclass>` and `<join>`, as follow:

```

<class name="Payment" table="PAYMENT">
  <id name="id" type="long" column="PAYMENT_ID">
    <generator class="native"/>
  </id>
  <discriminator column="PAYMENT_TYPE" type="string"/>
  <property name="amount" column="AMOUNT"/>
  ...
  <subclass name="CreditCardPayment" discriminator-value="CREDIT">
    <join table="CREDIT_PAYMENT">
      <property name="creditCardType" column="CCTYPE"/>
      ...
    </join>
  </subclass>
  <subclass name="CashPayment" discriminator-value="CASH">
    <join table="CASH_PAYMENT">
      ...
    </join>
  </subclass>
  <subclass name="ChequePayment" discriminator-value="CHEQUE">
    <join table="CHEQUE_PAYMENT" fetch="select">
      ...
    </join>
  </subclass>
</class>

```

The optional `fetch="select"` declaration tells Hibernate not to fetch the `ChequePayment` subclass data using an outer join when querying the superclass.

### 9.1.4. Mixing table per class hierarchy with table per subclass

You may even mix the table per hierarchy and table per subclass strategies using this approach:

```
<class name="Payment" table="PAYMENT">
  <id name="id" type="long" column="PAYMENT_ID">
    <generator class="native"/>
  </id>
  <discriminator column="PAYMENT_TYPE" type="string"/>
  <property name="amount" column="AMOUNT"/>
  ...
  <subclass name="CreditCardPayment" discriminator-value="CREDIT">
    <join table="CREDIT_PAYMENT">
      <property name="creditCardType" column="CCTYPE"/>
      ...
    </join>
  </subclass>
  <subclass name="CashPayment" discriminator-value="CASH">
    ...
  </subclass>
  <subclass name="ChequePayment" discriminator-value="CHEQUE">
    ...
  </subclass>
</class>
```

For any of these mapping strategies, a polymorphic association to the root `Payment` class is mapped using `<many-to-one>`.

```
<many-to-one name="payment" column="PAYMENT_ID" class="Payment"/>
```

### 9.1.5. Table per concrete class

There are two ways we could go about mapping the table per concrete class strategy. The first is to use `<union-subclass>`.

```
<class name="Payment">
  <id name="id" type="long" column="PAYMENT_ID">
    <generator class="native"/>
  </id>
  <property name="amount" column="AMOUNT"/>
  ...
  <union-subclass name="CreditCardPayment" table="CREDIT_PAYMENT">
    <property name="creditCardType" column="CCTYPE"/>
    ...
  </union-subclass>
  <union-subclass name="CashPayment" table="CASH_PAYMENT">
    ...
  </union-subclass>
  <union-subclass name="ChequePayment" table="CHEQUE_PAYMENT">
    ...
  </union-subclass>
</class>
```

Three tables are involved. Each table defines columns for all properties of the class, including inherited properties.

The limitation of this approach is that if a property is mapped on the superclass, the column name must be the same on all subclass tables. (We might relax this in a future release of Hibernate.)

### 9.1.6. Table per concrete class, using implicit polymorphism

An alternative approach is to make use of implicit polymorphism:

```

<class name="CreditCardPayment" table="CREDIT_PAYMENT">
  <id name="id" type="long" column="CREDIT_PAYMENT_ID">
    <generator class="native"/>
  </id>
  <property name="amount" column="CREDIT_AMOUNT"/>
  ...
</class>

<class name="CashPayment" table="CASH_PAYMENT">
  <id name="id" type="long" column="CASH_PAYMENT_ID">
    <generator class="native"/>
  </id>
  <property name="amount" column="CASH_AMOUNT"/>
  ...
</class>

<class name="ChequePayment" table="CHEQUE_PAYMENT">
  <id name="id" type="long" column="CHEQUE_PAYMENT_ID">
    <generator class="native"/>
  </id>
  <property name="amount" column="CHEQUE_AMOUNT"/>
  ...
</class>

```

Notice that nowhere do we mention the `Payment` interface explicitly. Also notice that properties of `Payment` are mapped in each of the subclasses. If you want to avoid duplication, consider using XML entities (e.g. `<!ENTITY allproperties SYSTEM "allproperties.xml"> [` in the DOCTYPE declaration and `&allproperties;` in the mapping).

The disadvantage of this approach is that Hibernate does not generate SQL `UNIONS` when performing polymorphic queries.

For this mapping strategy, a polymorphic association to `Payment` is usually mapped using `<any>`.

```

<any name="payment" meta-type="string" id-type="long">
  <meta-value value="CREDIT" class="CreditCardPayment"/>
  <meta-value value="CASH" class="CashPayment"/>
  <meta-value value="CHEQUE" class="ChequePayment"/>
  <column name="PAYMENT_CLASS"/>
  <column name="PAYMENT_ID"/>
</any>

```

### 9.1.7. Mixing implicit polymorphism with other inheritance mappings

There is one further thing to notice about this mapping. Since the subclasses are each mapped in their own `<class>` element (and since `Payment` is just an interface), each of the subclasses could easily be part of another inheritance hierarchy! (And you can still use polymorphic queries against the `Payment` interface.)

```

<class name="CreditCardPayment" table="CREDIT_PAYMENT">
  <id name="id" type="long" column="CREDIT_PAYMENT_ID">
    <generator class="native"/>
  </id>
  <discriminator column="CREDIT_CARD" type="string"/>
  <property name="amount" column="CREDIT_AMOUNT"/>
  ...
  <subclass name="MasterCardPayment" discriminator-value="MDC"/>
  <subclass name="VisaPayment" discriminator-value="VISA"/>
</class>

<class name="NonelectronicTransaction" table="NONELECTRONIC_TXN">
  <id name="id" type="long" column="TXN_ID">
    <generator class="native"/>
  </id>
  ...

```

```

<joined-subclass name="CashPayment" table="CASH_PAYMENT">
  <key column="PAYMENT_ID" />
  <property name="amount" column="CASH_AMOUNT" />
  ...
</joined-subclass>
<joined-subclass name="ChequePayment" table="CHEQUE_PAYMENT">
  <key column="PAYMENT_ID" />
  <property name="amount" column="CHEQUE_AMOUNT" />
  ...
</joined-subclass>
</class>

```

Once again, we don't mention `Payment` explicitly. If we execute a query against the `Payment` interface - for example, from `Payment` - Hibernate automatically returns instances of `CreditCardPayment` (and its subclasses, since they also implement `Payment`), `CashPayment` and `ChequePayment` but not instances of `NonelectronicTransaction`.

## 9.2. Limitations

There are certain limitations to the "implicit polymorphism" approach to the table per concrete-class mapping strategy. There are somewhat less restrictive limitations to `<union-subclass>` mappings.

The following table shows the limitations of table per concrete-class mappings, and of implicit polymorphism, in Hibernate.

**Table 9.1. Features of inheritance mappings**

Inheritance strategy	Poly-morphic many-to-one	Poly-morphic one-to-one	Poly-morphic one-to-many	Poly-morphic many-to-many	Poly-morphic <code>load()/get()</code>	Poly-morphic queries	Poly-morphic joins
table per class-hierarchy	<code>&lt;many-to-one&gt;</code>	<code>&lt;one-to-one&gt;</code>	<code>&lt;one-to-many&gt;</code>	<code>&lt;many-to-many&gt;</code>	<code>s.get(Payment.class, id)</code>	from Payment p	from Order o join o.payment p
table per subclass	<code>&lt;many-to-one&gt;</code>	<code>&lt;one-to-one&gt;</code>	<code>&lt;one-to-many&gt;</code>	<code>&lt;many-to-many&gt;</code>	<code>s.get(Payment.class, id)</code>	from Payment p	from Order o join o.payment p
table per concrete-class (union-subclass)	<code>&lt;many-to-one&gt;</code>	<code>&lt;one-to-one&gt;</code>	<code>&lt;one-to-many&gt;</code> (for <code>inverse="true"</code> only)	<code>&lt;many-to-many&gt;</code>	<code>s.get(Payment.class, id)</code>	from Payment p	from Order o join o.payment p
table per concrete class (implicit polymorphism)	<code>&lt;any&gt;</code>	<i>not supported</i>	<i>not supported</i>	<code>&lt;many-to-any&gt;</code>	<code>s.createCriteria(Payment.class).add( Restrictions.idEq</code>	from Payment p	<i>not supported</i>

<b>Inherit- ance strategy</b>	<b>Poly- morphic many- to-one</b>	<b>Poly- morphic one-to-one</b>	<b>Poly- morphic one- to-many</b>	<b>Poly- morphic many- to-many</b>	<b>Poly- morphic load() /get ( )</b>	<b>Poly- morphic queries</b>	<b>Poly- morphic joins</b>
					(id) ) .uniqueRe sult()		

---

## Chapter 10. Working with objects

Hibernate is a full object/relational mapping solution that not only shields the developer from the details of the underlying database management system, but also offers *state management* of objects. This is, contrary to the management of SQL statements in common JDBC/SQL persistence layers, a very natural object-oriented view of persistence in Java applications.

In other words, Hibernate application developers should always think about the *state* of their objects, and not necessarily about the execution of SQL statements. This part is taken care of by Hibernate and is only relevant for the application developer when tuning the performance of the system.

### 10.1. Hibernate object states

Hibernate defines and supports the following object states:

- *Transient* - an object is transient if it has just been instantiated using the `new` operator, and it is not associated with a Hibernate `Session`. It has no persistent representation in the database and no identifier value has been assigned. Transient instances will be destroyed by the garbage collector if the application doesn't hold a reference anymore. Use the Hibernate `Session` to make an object persistent (and let Hibernate take care of the SQL statements that need to be executed for this transition).
- *Persistent* - a persistent instance has a representation in the database and an identifier value. It might just have been saved or loaded, however, it is by definition in the scope of a `Session`. Hibernate will detect any changes made to an object in persistent state and synchronize the state with the database when the unit of work completes. Developers don't execute manual `UPDATE` statements, or `DELETE` statements when an object should be made transient.
- *Detached* - a detached instance is an object that has been persistent, but its `Session` has been closed. The reference to the object is still valid, of course, and the detached instance might even be modified in this state. A detached instance can be reattached to a new `Session` at a later point in time, making it (and all the modifications) persistent again. This feature enables a programming model for long running units of work that require user think-time. We call them *application transactions*, i.e. a unit of work from the point of view of the user.

We'll now discuss the states and state transitions (and the Hibernate methods that trigger a transition) in more detail.

### 10.2. Making objects persistent

Newly instantiated instances of a persistent class are considered *transient* by Hibernate. We can make a transient instance *persistent* by associating it with a session:

```
DomesticCat fritz = new DomesticCat();
fritz.setColor(Color.GINGER);
fritz.setSex('M');
fritz.setName("Fritz");
Long generatedId = (Long) sess.save(fritz);
```

If `Cat` has a generated identifier, the identifier is generated and assigned to the `cat` when `save()` is called. If `Cat` has an assigned identifier, or a composite key, the identifier should be assigned to the `cat` instance before calling `save()`. You may also use `persist()` instead of `save()`, with the semantics defined in the EJB3 early



draft.

Alternatively, you may assign the identifier using an overloaded version of `save()`.

```
DomesticCat pk = new DomesticCat();
pk.setColor(Color.TABBY);
pk.setSex('F');
pk.setName("PK");
pk.setKittens( new HashSet() );
pk.addKitten(fritz);
sess.save( pk, new Long(1234) );
```

If the object you make persistent has associated objects (e.g. the `kittens` collection in the previous example), these objects may be made persistent in any order you like unless you have a `NOT NULL` constraint upon a foreign key column. There is never a risk of violating foreign key constraints. However, you might violate a `NOT NULL` constraint if you `save()` the objects in the wrong order.

Usually you don't bother with this detail, as you'll very likely use Hibernate's *transitive persistence* feature to save the associated objects automatically. Then, even `NOT NULL` constraint violations don't occur - Hibernate will take care of everything. Transitive persistence is discussed later in this chapter.

## 10.3. Loading an object

The `load()` methods of `Session` gives you a way to retrieve a persistent instance if you already know its identifier. `load()` takes a class object and will load the state into a newly instantiated instance of that class, in persistent state.

```
Cat fritz = (Cat) sess.load(Cat.class, generatedId);
```

```
// you need to wrap primitive identifiers
long pkId = 1234;
DomesticCat pk = (DomesticCat) sess.load( Cat.class, new Long(pkId) );
```

Alternatively, you can load state into a given instance:

```
Cat cat = new DomesticCat();
// load pk's state into cat
sess.load( cat, new Long(pkId) );
Set kittens = cat.getKittens();
```

Note that `load()` will throw an unrecoverable exception if there is no matching database row. If the class is mapped with a proxy, `load()` just returns an uninitialized proxy and does not actually hit the database until you invoke a method of the proxy. This behaviour is very useful if you wish to create an association to an object without actually loading it from the database. It also allows multiple instances to be loaded as a batch if `batch-size` is defined for the class mapping.

If you are not certain that a matching row exists, you should use the `get()` method, which hits the database immediately and returns null if there is no matching row.

```
Cat cat = (Cat) sess.get(Cat.class, id);
if (cat==null) {
    cat = new Cat();
    sess.save(cat, id);
}
return cat;
```

You may even load an object using an SQL `SELECT ... FOR UPDATE`, using a `LockMode`. See the API docu-

mentation for more information.

```
Cat cat = (Cat) sess.get(Cat.class, id, LockMode.UPGRADE);
```

Note that any associated instances or contained collections are *not* selected FOR UPDATE, unless you decide to specify lock or all as a cascade style for the association.

It is possible to re-load an object and all its collections at any time, using the `refresh()` method. This is useful when database triggers are used to initialize some of the properties of the object.

```
sess.save(cat);
sess.flush(); //force the SQL INSERT
sess.refresh(cat); //re-read the state (after the trigger executes)
```

An important question usually appears at this point: How much does Hibernate load from the database and how many SQL SELECTS will it use? This depends on the *fetching strategy* and is explained in Section 19.1, “Fetching strategies”.

## 10.4. Querying

If you don't know the identifiers of the objects you are looking for, you need a query. Hibernate supports an easy-to-use but powerful object oriented query language (HQL). For programmatic query creation, Hibernate supports a sophisticated Criteria and Example query feature (QBC and QBE). You may also express your query in the native SQL of your database, with optional support from Hibernate for result set conversion into objects.

### 10.4.1. Executing queries

HQL and native SQL queries are represented with an instance of `org.hibernate.Query`. This interface offers methods for parameter binding, result set handling, and for the execution of the actual query. You always obtain a `Query` using the current `Session`:

```
List cats = session.createQuery(
    "from Cat as cat where cat.birthdate < ?")
    .setDate(0, date)
    .list();

List mothers = session.createQuery(
    "select mother from Cat as cat join cat.mother as mother where cat.name = ?")
    .setString(0, name)
    .list();

List kittens = session.createQuery(
    "from Cat as cat where cat.mother = ?")
    .setEntity(0, pk)
    .list();

Cat mother = (Cat) session.createQuery(
    "select cat.mother from Cat as cat where cat = ?")
    .setEntity(0, izi)
    .uniqueResult();
```

A query is usually executed by invoking `list()`, the result of the query will be loaded completely into a collection in memory. Entity instances retrieved by a query are in persistent state. The `uniqueResult()` method offers a shortcut if you know your query will only return a single object.

### Iterating results

Occasionally, you might be able to achieve better performance by executing the query using the `iterate()` method. This will only usually be the case if you expect that the actual entity instances returned by the query will already be in the session or second-level cache. If they are not already cached, `iterate()` will be slower than `list()` and might require many database hits for a simple query, usually *1* for the initial select which only returns identifiers, and *n* additional selects to initialize the actual instances.

```
// fetch ids
Iterator iter = sess.createQuery("from eg.Qux q order by q.likeliness").iterate();
while ( iter.hasNext() ) {
    Qux qux = (Qux) iter.next(); // fetch the object
    // something we couldnt express in the query
    if ( qux.calculateComplicatedAlgorithm() ) {
        // delete the current instance
        iter.remove();
        // dont need to process the rest
        break;
    }
}
```

## Queries that return tuples

Hibernate queries sometimes return tuples of objects, in which case each tuple is returned as an array:

```
Iterator kittensAndMothers = sess.createQuery(
    "select kitten, mother from Cat kitten join kitten.mother mother")
    .list()
    .iterator();

while ( kittensAndMothers.hasNext() ) {
    Object[] tuple = (Object[]) kittensAndMothers.next();
    Cat kitten = tuple[0];
    Cat mother = tuple[1];
    ....
}
```

## Scalar results

Queries may specify a property of a class in the `select` clause. They may even call SQL aggregate functions. Properties or aggregates are considered "scalar" results (and not entities in persistent state).

```
Iterator results = sess.createQuery(
    "select cat.color, min(cat.birthdate), count(cat) from Cat cat " +
    "group by cat.color")
    .list()
    .iterator();

while ( results.hasNext() ) {
    Object[] row = results.next();
    Color type = (Color) row[0];
    Date oldest = (Date) row[1];
    Integer count = (Integer) row[2];
    .....
}
```

## Bind parameters

Methods on `Query` are provided for binding values to named parameters or JDBC-style `?` parameters. *Contrary to JDBC, Hibernate numbers parameters from zero.* Named parameters are identifiers of the form `:name` in the query string. The advantages of named parameters are:

- named parameters are insensitive to the order they occur in the query string

- they may occur multiple times in the same query
- they are self-documenting

```
//named parameter (preferred)
Query q = sess.createQuery("from DomesticCat cat where cat.name = :name");
q.setString("name", "Fritz");
Iterator cats = q.iterate();
```

```
//positional parameter
Query q = sess.createQuery("from DomesticCat cat where cat.name = ?");
q.setString(0, "Izi");
Iterator cats = q.iterate();
```

```
//named parameter list
List names = new ArrayList();
names.add("Izi");
names.add("Fritz");
Query q = sess.createQuery("from DomesticCat cat where cat.name in (:namesList)");
q.setParameterList("namesList", names);
List cats = q.list();
```

## Pagination

If you need to specify bounds upon your result set (the maximum number of rows you want to retrieve and / or the first row you want to retrieve) you should use methods of the `Query` interface:

```
Query q = sess.createQuery("from DomesticCat cat");
q.setFirstResult(20);
q.setMaxResults(10);
List cats = q.list();
```

Hibernate knows how to translate this limit query into the native SQL of your DBMS.

## Scrollable iteration

If your JDBC driver supports scrollable `ResultSets`, the `Query` interface may be used to obtain a `ScrollableResults` object, which allows flexible navigation of the query results.

```
Query q = sess.createQuery("select cat.name, cat from DomesticCat cat " +
    "order by cat.name");
ScrollableResults cats = q.scroll();
if ( cats.first() ) {

    // find the first name on each page of an alphabetical list of cats by name
    firstNamesOfPages = new ArrayList();
    do {
        String name = cats.getString(0);
        firstNamesOfPages.add(name);
    }
    while ( cats.scroll(PAGE_SIZE) );

    // Now get the first page of cats
    pageOfCats = new ArrayList();
    cats.beforeFirst();
    int i=0;
    while( ( PAGE_SIZE > i++ ) && cats.next() ) pageOfCats.add( cats.get(1) );

}
cats.close();
```

Note that an open database connection (and cursor) is required for this functionality, use `setMaxResult()/setFirstResult()` if you need offline pagination functionality.

## Externalizing named queries

You may also define named queries in the mapping document. (Remember to use a `CDATA` section if your query contains characters that could be interpreted as markup.)

```
<query name="eg.DomesticCat.by.name.and.minimum.weight"><![CDATA[
    from eg.DomesticCat as cat
      where cat.name = ?
      and cat.weight > ?
  ]]></query>
```

Parameter binding and executing is done programatically:

```
Query q = sess.getNamedQuery("eg.DomesticCat.by.name.and.minimum.weight");
q.setString(0, name);
q.setInt(1, minWeight);
List cats = q.list();
```

Note that the actual program code is independent of the query language that is used, you may also define native SQL queries in metadata, or migrate existing queries to Hibernate by placing them in mapping files.

### 10.4.2. Filtering collections

A collection *filter* is a special type of query that may be applied to a persistent collection or array. The query string may refer to `this`, meaning the current collection element.

```
Collection blackKittens = session.createFilter(
    pk.getKittens(),
    "where this.color = ?")
    .setParameter( Color.BLACK, Hibernate.custom(ColorUserType.class) )
    .list()
);
```

The returned collection is considered a bag, and it's a copy of the given collection. The original collection is not modified (this is contrary to the implication of the name "filter", but consistent with expected behavior).

Observe that filters do not require a `from` clause (though they may have one if required). Filters are not limited to returning the collection elements themselves.

```
Collection blackKittenMates = session.createFilter(
    pk.getKittens(),
    "select this.mate where this.color = eg.Color.BLACK.intValue")
    .list();
```

Even an empty filter query is useful, e.g. to load a subset of elements in a huge collection:

```
Collection tenKittens = session.createFilter(
    mother.getKittens(), "")
    .setFirstResult(0).setMaxResults(10)
    .list();
```

### 10.4.3. Criteria queries

HQL is extremely powerful but some developers prefer to build queries dynamically, using an object-oriented API, rather than building query strings. Hibernate provides an intuitive `Criteria` query API for these cases:

```
Criteria crit = session.createCriteria(Cat.class);
```

```
crit.add( Expression.eq( "color", eg.Color.BLACK ) );
crit.setMaxResults(10);
List cats = crit.list();
```

The Criteria and the associated `Example` API are discussed in more detail in Chapter 15, *Criteria Queries*.

#### 10.4.4. Queries in native SQL

You may express a query in SQL, using `createSQLQuery()` and let Hibernate take care of the mapping from result sets to objects. Note that you may at any time call `session.connection()` and use the `JDBC Connection` directly. If you chose to use the Hibernate API, you must enclose SQL aliases in braces:

```
List cats = session.createSQLQuery(
    "SELECT {cat.*} FROM CAT {cat} WHERE ROWNUM<10",
    "cat",
    Cat.class
).list();
```

```
List cats = session.createSQLQuery(
    "SELECT {cat}.ID AS {cat.id}, {cat}.SEX AS {cat.sex}, " +
        "{cat}.MATE AS {cat.mate}, {cat}.SUBCLASS AS {cat.class}, ... " +
    "FROM CAT {cat} WHERE ROWNUM<10",
    "cat",
    Cat.class
).list();
```

SQL queries may contain named and positional parameters, just like Hibernate queries. More information about native SQL queries in Hibernate can be found in Chapter 16, *Native SQL*.

### 10.5. Modifying persistent objects

*Transactional persistent instances* (ie. objects loaded, saved, created or queried by the `Session`) may be manipulated by the application and any changes to persistent state will be persisted when the `Session` is *flushed* (discussed later in this chapter). There is no need to call a particular method (like `update()`, which has a different purpose) to make your modifications persistent. So the most straightforward way to update the state of an object is to `load()` it, and then manipulate it directly, while the `Session` is open:

```
DomesticCat cat = (DomesticCat) sess.load( Cat.class, new Long(69) );
cat.setName("PK");
sess.flush(); // changes to cat are automatically detected and persisted
```

Sometimes this programming model is inefficient since it would require both an SQL `SELECT` (to load an object) and an SQL `UPDATE` (to persist its updated state) in the same session. Therefore Hibernate offers an alternate approach, using detached instances.

*Note that Hibernate does not offer its own API for direct execution of `UPDATE` or `DELETE` statements. Hibernate is a state management service, you don't have to think in statements to use it. `JDBC` is a perfect API for executing SQL statements, you can get a `JDBC Connection` at any time by calling `session.connection()`. Furthermore, the notion of mass operations conflicts with object/relational mapping for online transaction processing-oriented applications. Future versions of Hibernate may however provide special mass operation functions. See Chapter 13, *Batch processing* for some possible batch operation tricks.*

### 10.6. Modifying detached objects

Many applications need to retrieve an object in one transaction, send it to the UI layer for manipulation, then save the changes in a new transaction. Applications that use this kind of approach in a high-concurrency environment usually use versioned data to ensure isolation for the "long" unit of work.

Hibernate supports this model by providing for reattachment of detached instances using the `Session.update()` or `Session.merge()` methods:

```
// in the first session
Cat cat = (Cat) firstSession.load(Cat.class, catId);
Cat potentialMate = new Cat();
firstSession.save(potentialMate);

// in a higher layer of the application
cat.setMate(potentialMate);

// later, in a new session
secondSession.update(cat); // update cat
secondSession.update(mate); // update mate
```

If the `Cat` with identifier `catId` had already been loaded by `secondSession` when the application tried to reattach it, an exception would have been thrown.

Use `update()` if you are sure that the session does not contain an already persistent instance with the same identifier, and `merge()` if you want to merge your modifications at any time without consideration of the state of the session. In other words, `update()` is usually the first method you would call in a fresh session, ensuring that reattachment of your detached instances is the first operation that is executed.

The application should individually `update()` detached instances reachable from the given detached instance if and *only* if it wants their state also updated. This can be automated of course, using *transitive persistence*, see Section 10.11, “Transitive persistence”.

The `lock()` method also allows an application to reassociate an object with a new session. However, the detached instance has to be unmodified!

```
//just reassociate:
sess.lock(fritz, LockMode.NONE);
//do a version check, then reassociate:
sess.lock(izi, LockMode.READ);
//do a version check, using SELECT ... FOR UPDATE, then reassociate:
sess.lock(pk, LockMode.UPGRADE);
```

Note that `lock()` can be used with various `LockModes`, see the API documentation and the chapter on transaction handling for more information. Reattachment is not the only usecase for `lock()`.

Other models for long units of work are discussed in Section 11.3, “Optimistic concurrency control”.

## 10.7. Automatic state detection

Hibernate users have requested a general purpose method that either saves a transient instance by generating a new identifier or updates/reattaches the detached instances associated with its current identifier. The `saveOrUpdate()` method implements this functionality.

```
// in the first session
Cat cat = (Cat) firstSession.load(Cat.class, catID);

// in a higher tier of the application
Cat mate = new Cat();
cat.setMate(mate);
```

```
// later, in a new session
secondSession.saveOrUpdate(cat); // update existing state (cat has a non-null id)
secondSession.saveOrUpdate(mate); // save the new instance (mate has a null id)
```

The usage and semantics of `saveOrUpdate()` seems to be confusing for new users. Firstly, so long as you are not trying to use instances from one session in another new session, you should not need to use `update()`, `saveOrUpdate()`, or `merge()`. Some whole applications will never use either of these methods.

Usually `update()` or `saveOrUpdate()` are used in the following scenario:

- the application loads an object in the first session
- the object is passed up to the UI tier
- some modifications are made to the object
- the object is passed back down to the business logic tier
- the application persists these modifications by calling `update()` in a second session

`saveOrUpdate()` does the following:

- if the object is already persistent in this session, do nothing
- if another object associated with the session has the same identifier, throw an exception
- if the object has no identifier property, `save()` it
- if the object's identifier has the value assigned to a newly instantiated object, `save()` it
- if the object is versioned (by a `<version>` or `<timestamp>`), and the version property value is the same value assigned to a newly instantiated object, `save()` it
- otherwise `update()` the object

and `merge()` is very different:

- if there is a persistent instance with the same identifier currently associated with the session, copy the state of the given object onto the persistent instance
- if there is no persistent instance currently associated with the session, try to load it from the database, or create a new persistent instance
- the persistent instance is returned
- the given instance does not become associated with the session, it remains detached

## 10.8. Deleting persistent objects

`Session.delete()` will remove an object's state from the database. Of course, your application might still hold a reference to a deleted object. It's best to think of `delete()` as making a persistent instance transient.

```
sess.delete(cat);
```

You may delete objects in any order you like, without risk of foreign key constraint violations. It is still possible to violate a `NOT NULL` constraint on a foreign key column by deleting objects in the wrong order, e.g. if you delete the parent, but forget to delete the children.

## 10.9. Replicating object between two different datastores

It is occasionally useful to be able to take a graph of persistent instances and make them persistent in a different datastore, without regenerating identifier values.

```
//retrieve a cat from one database
```



```

Session session1 = factory1.openSession();
Transaction tx1 = session1.beginTransaction();
Cat cat = session1.get(Cat.class, catId);
tx1.commit();
session1.close();

//reconcile with a second database
Session session2 = factory2.openSession();
Transaction tx2 = session2.beginTransaction();
session2.replicate(cat, ReplicationMode.LATEST_VERSION);
tx2.commit();
session2.close();

```

The `ReplicationMode` determines how `replicate()` will deal with conflicts with existing rows in the database.

- `ReplicationMode.IGNORE` - ignore the object when there is an existing database row with the same identifier
- `ReplicationMode.OVERWRITE` - overwrite any existing database row with the same identifier
- `ReplicationMode.EXCEPTION` - throw an exception if there is an existing database row with the same identifier
- `ReplicationMode.LATEST_VERSION` - overwrite the row if its version number is earlier than the version number of the object, or ignore the object otherwise

Use cases for this feature include reconciling data entered into different database instances, upgrading system configuration information during product upgrades, rolling back changes made during non-ACID transactions and more.

## 10.10. Flushing the Session

From time to time the `Session` will execute the SQL statements needed to synchronize the JDBC connection's state with the state of objects held in memory. This process, *flush*, occurs by default at the following points

- before some query executions
- from `org.hibernate.Transaction.commit()`
- from `Session.flush()`

The SQL statements are issued in the following order

1. all entity insertions, in the same order the corresponding objects were saved using `Session.save()`
2. all entity updates
3. all collection deletions
4. all collection element deletions, updates and insertions
5. all collection insertions
6. all entity deletions, in the same order the corresponding objects were deleted using `Session.delete()`

(An exception is that objects using `native ID` generation are inserted when they are saved.)

Except when you explicitly `flush()`, there are absolutely no guarantees about *when* the `Session` executes the JDBC calls, only the *order* in which they are executed. However, Hibernate does guarantee that the `Query.list(...)` will never return stale data; nor will they return the wrong data.

It is possible to change the default behavior so that flush occurs less frequently. The `FlushMode` class defines three different modes: only flush at commit time (and only when the Hibernate `Transaction` API is used), flush automatically using the explained routine, or never flush unless `flush()` is called explicitly. The last mode is useful for long running units of work, where a `Session` is kept open and disconnected for a long time (see Sec-

tion 11.3.2, “Long session and automatic versioning”).

```
sess = sf.openSession();
Transaction tx = sess.beginTransaction();
sess.setFlushMode(FlushMode.COMMIT); // allow queries to return stale state

Cat izi = (Cat) sess.load(Cat.class, id);
izi.setName(iznizi);

// might return stale data
sess.find("from Cat as cat left outer join cat.kittens kitten");

// change to izi is not flushed!
...
tx.commit(); // flush occurs
```

During flush, an exception might occur (e.g. if a DML operation violates a constraint). Since handling exceptions involves some understanding of Hibernate's transactional behavior, we discuss it in Chapter 11, *Transactions And Concurrency*.

## 10.11. Transitive persistence

It is quite cumbersome to save, delete, or reattach individual objects, especially if you deal with a graph of associated objects. A common case is a parent/child relationship. Consider the following example:

If the children in a parent/child relationship would be value typed (e.g. a collection of addresses or strings), their lifecycle would depend on the parent and no further action would be required for convenient "cascading" of state changes. When the parent is saved, the value-typed child objects are saved as well, when the parent is deleted, the children will be deleted, etc. This even works for operations such as the removal of a child from the collection; Hibernate will detect this and, since value-typed objects can't have shared references, delete the child from the database.

Now consider the same scenario with parent and child objects being entities, not value-types (e.g. categories and items, or parent and child cats). Entities have their own lifecycle, support shared references (so removing an entity from the collection does not mean it can be deleted), and there is by default no cascading of state from one entity to any other associated entities. Hibernate does not implement *persistence by reachability* by default.

For each basic operation of the Hibernate session - including `persist()`, `merge()`, `saveOrUpdate()`, `delete()`, `lock()`, `refresh()`, `evict()`, `replicate()` - there is a corresponding cascade style. Respectively, the cascade styles are named `create`, `merge`, `save-update`, `delete`, `lock`, `refresh`, `evict`, `replicate`. If you want an operation to be cascaded along an association, you must indicate that in the mapping document. For example:

```
<one-to-one name="person" cascade="create"/>
```

Cascade styles may be combined:

```
<one-to-one name="person" cascade="create,delete,lock"/>
```

You may even use `cascade="all"` to specify that *all* operations should be cascaded along the association. The default `cascade="none"` specifies that no operations are to be cascaded.

A special cascade style, `delete-orphan`, applies only to one-to-many associations, and indicates that the `delete()` operation should be applied to any child object that is removed from the association.

Recommendations:

- It doesn't usually make sense to enable cascade on a `<many-to-one>` or `<many-to-many>` association. Cascade is often useful for `<one-to-one>` and `<one-to-many>` associations.
- If the child object's lifespan is bounded by the lifespan of the of the parent object make it a *lifecycle object* by specifying `cascade="all,delete-orphan"`.
- Otherwise, you might not need cascade at all. But if you think that you will often be working with the parent and children together in the same transaction, and you want to save yourself some typing, consider using `cascade="create,merge,save-update"`.

Mapping an association (either a single valued association, or a collection) with `cascade="all"` marks the association as a *parent/child* style relationship where save/update/delete of the parent results in save/update/delete of the child or children.

Futhermore, a mere reference to a child from a persistent parent will result in save/update of the child. This metaphor is incomplete, however. A child which becomes unreferenced by its parent is *not* automatically deleted, except in the case of a `<one-to-many>` association mapped with `cascade="delete-orphan"`. The precise semantics of cascading operations for a parent/child relationship are as follows:

- If a parent is passed to `persist()`, all children are passed to `persist()`
- If a parent is passed to `merge()`, all children are passed to `merge()`
- If a parent is passed to `save()`, `update()` or `saveOrUpdate()`, all children are passed to `saveOrUpdate()`
- If a transient or detached child becomes referenced by a persistent parent, it is passed to `saveOrUpdate()`
- If a parent is deleted, all children are passed to `delete()`
- If a child is dereferenced by a persistent parent, *nothing special happens* - the application should explicitly delete the child if necessary - unless `cascade="delete-orphan"`, in which case the "orphaned" child is deleted.

## 10.12. Using metadata

Hibernate requires a very rich meta-level model of all entity and value types. From time to time, this model is very useful to the application itself. For example, the application might use Hibernate's metadata to implement a "smart" deep-copy algorithm that understands which objects should be copied (eg. mutable value types) and which should not (eg. immutable value types and, possibly, associated entities).

Hibernate exposes metadata via the `ClassMetadata` and `CollectionMetadata` interfaces and the `Type` hierarchy. Instances of the metadata interfaces may be obtained from the `SessionFactory`.

```
Cat fritz = .....;
ClassMetadata catMeta = sessionFactory.getClassMetadata(Cat.class);

Object[] propertyValues = catMeta.getPropertyValues(fritz);
String[] propertyNames = catMeta.getPropertyNames();
Type[] propertyTypes = catMeta.getPropertyTypes();

// get a Map of all properties which are not collections or associations
Map namedValues = new HashMap();
for ( int i=0; i<propertyNames.length; i++ ) {
    if ( !propertyTypes[i].isEntityType() && !propertyTypes[i].isCollectionType() ) {
        namedValues.put( propertyNames[i], propertyValues[i] );
    }
}
```

---

# Chapter 11. Transactions And Concurrency

The most important point about Hibernate and concurrency control is that it is very easy to understand. Hibernate directly uses JDBC connections and JTA resources without adding any additional locking behavior. We highly recommend you spend some time with the JDBC, ANSI, and transaction isolation specification of your database management system. Hibernate only adds automatic versioning but does not lock objects in memory or change the isolation level of your database transactions. Basically, use Hibernate like you would use direct JDBC (or JTA/CMT) with your database resources.

However, in addition to automatic versioning, Hibernate also offers a (minor) API for pessimistic locking of rows, using the `SELECT FOR UPDATE` syntax. This API is discussed later in this chapter.

We start the discussion of concurrency control in Hibernate with the granularity of `Configuration`, `SessionFactory`, and `Session`, as well as database and long application transactions.

## 11.1. Session and transaction scopes

A `SessionFactory` is an expensive-to-create, threadsafe object intended to be shared by all application threads. It is created once, usually on application startup, from a `Configuration` instance.

A `Session` is an inexpensive, non-threadsafe object that should be used once, for a single business process, a single unit of work, and then discarded. A `Session` will not obtain a `JDBC Connection` (or a `Datasource`) unless it is needed, so you may safely open and close a `Session` even if you are not sure that data access will be needed to serve a particular request. (This becomes important as soon as you are implementing some of the following patterns using request interception.)

To complete this picture you also have to think about database transactions. A database transaction has to be as short as possible, to reduce lock contention in the database. Long database transactions will prevent your application from scaling to highly concurrent load.

What is the scope of a unit of work? Can a single Hibernate `Session` span several database transactions or is this a one-to-one relationship of scopes? When should you open and close a `Session` and how do you demarcate the database transaction boundaries?

### 11.1.1. Unit of work

First, don't use the *session-per-operation* antipattern, that is, don't open and close a `Session` for every simple database call in a single thread! Of course, the same is true for database transactions. Database calls in an application are made using a planned sequence, they are grouped into atomic units of work. (Note that this also means that auto-commit after every single SQL statement is useless in an application, this mode is intended for ad-hoc SQL console work. Hibernate disables, or expects the application server to do so, auto-commit mode immediately.)

The most common pattern in a multi-user client/server application is *session-per-request*. In this model, a request from the client is sent to the server (where the Hibernate persistence layer runs), a new Hibernate `Session` is opened, and all database operations are executed in this unit of work. Once the work has been completed (and the response for the client has been prepared), the session is flushed and closed. You would also use a single database transaction to serve the clients request, starting and committing it when you open and close the `Session`. The relationship between the two is one-to-one and this model is a perfect fit for many applications.

The challenge lies in the implementation: not only has the `Session` and transaction to be started and ended correctly, but they also have to be accessible for data access operations. The demarcation of a unit of work is ideally implemented using an interceptor that runs when a request hits the server and before the response will be send (i.e. a `ServletFilter`). We recommend to bind the `Session` to the thread that serves the request, using a `ThreadLocal` variable. This allows easy access (like accessing a static variable) in all code that runs in this thread. Depending on the database transaction demarcation mechanism you chose, you might also keep the transaction context in a `ThreadLocal` variable. The implementation patterns for this are known as *ThreadLocal Session* and *Open Session in View*. You can easily extend the `HibernateUtil` helper class shown earlier in this documentation to implement this. Of course, you'd have to find a way to implement an interceptor and set it up in your environment. See the Hibernate website for tips and examples.

### 11.1.2. Application transactions

The session-per-request pattern is not the only useful concept you can use to design units of work. Many business processes require a whole series of interactions with the user interleaved with database accesses. In web and enterprise applications it is not acceptable for a database transaction to span a user interaction. Consider the following example:

- The first screen of a dialog opens, the data seen by the user has been loaded in a particular `Session` and database transaction. The user is free to modify the objects.
- The user clicks "Save" after 5 minutes and expects his modifications to be made persistent; he also expects that he was the only person editing this information and that no conflicting modification can occur.

We call this unit of work, from the point of view of the user, a long running *application transaction*. There are many ways how you can implement this in your application.

A first naive implementation might keep the `Session` and database transaction open during user think time, with locks held in the database to prevent concurrent modification, and to guarantee isolation and atomicity. This is of course an anti-pattern, since lock contention would not allow the application to scale with the number of concurrent users.

Clearly, we have to use several database transactions to implement the application transaction. In this case, maintaining isolation of business processes becomes the partial responsibility of the application tier. A single application transaction usually spans several database transactions. It will be atomic if only one of these database transactions (the last one) stores the updated data, all others simply read data (e.g. in a wizard-style dialog spanning several request/response cycles). This is easier to implement than it might sound, especially if you use Hibernate's features:

- *Automatic Versioning* - Hibernate can do automatic optimistic concurrency control for you, it can automatically detect if a concurrent modification occurred during user think time.
- *Detached Objects* - If you decide to use the already discussed *session-per-request* pattern, all loaded instances will be in detached state during user think time. Hibernate allows you to reattach the objects and persist the modifications, the pattern is called *session-per-request-with-detached-objects*. Automatic versioning is used to isolate concurrent modifications.
- *Long Session* - The Hibernate `Session` may be disconnected from the underlying JDBC connection after the database transaction has been committed, and reconnected when a new client request occurs. This pattern is known as *session-per-application-transaction* and makes even reattachment unnecessary. Automatic versioning is used to isolate concurrent modifications.

Both *session-per-request-with-detached-objects* and *session-per-application-transaction* have advantages and disadvantages, we discuss them later in this chapter in the context of optimistic concurrency control.

### 11.1.3. Considering object identity

An application may concurrently access the same persistent state in two different `Session`s. However, an instance of a persistent class is never shared between two `Session` instances. Hence there are two different notions of identity:

#### Database Identity

```
foo.getId().equals( bar.getId() )
```

#### JVM Identity

```
foo==bar
```

Then for objects attached to a *particular* `Session` (i.e. in the scope of a `Session`) the two notions are equivalent, and JVM identity for database identity is guaranteed by Hibernate. However, while the application might concurrently access the "same" (persistent identity) business object in two different sessions, the two instances will actually be "different" (JVM identity). Conflicts are resolved using (automatic versioning) at flush/commit time, using an optimistic approach.

This approach leaves Hibernate and the database to worry about concurrency; it also provides the best scalability, since guaranteeing identity in single-threaded units of work only doesn't need expensive locking or other means of synchronization. The application never needs to synchronize on any business object, as long as it sticks to a single thread per `Session`. Within a `Session` the application may safely use `==` to compare objects.

However, an application that uses `==` outside of a `Session`, might see unexpected results. This might occur even in some unexpected places, for example, if you put two detached instances into the same `Set`. Both might have the same database identity (i.e. they represent the same row), but JVM identity is by definition not guaranteed for instances in detached state. The developer has to override the `equals()` and `hashCode()` methods in persistent classes and implement his own notion of object equality. There is one caveat: Never use the database identifier to implement equality, use a business key, a combination of unique, usually immutable, attributes. The database identifier will change if a transient object is made persistent. If the transient instance (usually together with detached instances) is held in a `Set`, changing the hashCode breaks the contract of the `Set`. Attributes for business keys don't have to be as stable as database primary keys, you only have to guarantee stability as long as the objects are in the same `Set`. See the Hibernate website for a more thorough discussion of this issue. Also note that this is not a Hibernate issue, but simply how Java object identity and equality has to be implemented.

### 11.1.4. Common issues

Never use the anti-patterns *session-per-user-session* or *session-per-application* (of course, there are rare exceptions to this rule). Note that some of the following issues might also appear with the recommended patterns, make sure you understand the implications before making a design decision:

- A `Session` is not thread-safe. Things which are supposed to work concurrently, like HTTP requests, session beans, or Swing workers, will cause race conditions if a `Session` instance would be shared. If you keep your Hibernate `Session` in your `HttpSession` (discussed later), you should consider synchronizing access to your `Http` session. Otherwise, a user that clicks reload fast enough may use the same `Session` in two concurrently running threads.
- An exception thrown by Hibernate means you have to rollback your database transaction and close the `Session`.

sion immediately (discussed later in more detail). If your `Session` is bound to the application, you have to stop the application. Rolling back the database transaction doesn't put your business objects back into the state they were at the start of the transaction. This means the database state and the business objects do get out of sync. Usually this is not a problem, because exceptions are not recoverable and you have to start over after rollback anyway.

- The `Session` caches every object that is in persistent state (watched and checked for dirty state by Hibernate). This means it grows endlessly until you get an `OutOfMemoryException`, if you keep it open for a long time or simply load too much data. One solution for this is to call `clear()` and `evict()` to manage the `Session` cache, but you most likely should consider a Stored Procedure if you need mass data operations. Some solutions are shown in Chapter 13, *Batch processing*. Keeping a `Session` open for the duration of a user session also means a high probability of stale data.

## 11.2. Database transaction demarcation

Datatabase (or system) transaction boundaries are always necessary. No communication with the database can occur outside of a database transaction (this seems to confuse many developers who are used to the auto-commit mode). Always use clear transaction boundaries, even for read-only operations. Depending on your isolation level and database capabilities this might not be required but there is no downside if you always demarcate transactions explicitly.

A Hibernate application can run in non-managed (i.e. standalone, simple Web- or Swing applications) and managed J2EE environments. In a non-managed environment, Hibernate is usually responsible for its own database connection pool. The application developer has to manually set transaction boundaries, in other words, begin, commit, or rollback database transactions himself. A managed environment usually provides container-managed transactions, with the transaction assembly defined declaratively in deployment descriptors of EJB session beans, for example. Programmatic transaction demarcation is then no longer necessary, even flushing the `Session` is done automatically.

However, it is often desirable to keep your persistence layer portable. Hibernate offers a wrapper API called `Transaction` that translates into the native transaction system of your deployment environment. This API is optional (using database transactions is not!) and you don't have to use it if database portability provided by Hibernate is all you need.

Usually, ending a `Session` involves four distinct phases:

- flush the session
- commit the transaction
- close the session
- handle exceptions

Flushing the session has been discussed earlier, we'll now have a closer look at transaction demarcation and exception handling in both managed- and non-managed environments.

### 11.2.1. Non-managed environment

If a Hibernate persistence layer runs in a non-managed environment, database connections are either handled by Hibernate's pooling mechanism or provided by the developer (this case has other implications, esp. with regard to caching):

```
// Session sess = factory.openSession(myConnection);  
Session sess = factory.openSession();
```

```

try {
    // do some work
    ...
    sess.flush();
    sess.connection().commit();
}
catch (RuntimeException e) {
    sess.connection().rollback();
    throw e; // or display error message
}
finally {
    sess.close();
}

```

Note that you will very likely never see this piece of code in a normal application; fatal (system) exceptions should always be caught at the "top". In other words, the code that executes Hibernate calls (in the persistence layer) and the code that handles `RuntimeException` (and usually can only clean up and exit) are in different layers. This can be a challenge to design yourself and you should use J2EE/EJB container services whenever they are available. Exception handling is discussed later in this chapter.

We recommend, even if persistence layer portability is not your primary concern, the `Transaction API`:

```

Session sess = factory.openSession();
Transaction tx = null;
try {
    tx = sess.beginTransaction();

    // do some work
    ...

    tx.commit();
}
catch (RuntimeException e) {
    if (tx != null) tx.rollback();
    throw e; // or display error message
}
finally {
    sess.close();
}

```

Note that you don't have to `flush()` the `Session` explicitly, the call to `commit()` automatically triggers the synchronization. This piece of code is now portable and runs in non-managed and JTA environments. See Section 3.4.5, "Transaction strategy configuration" for the configuration options of the `Transaction API` and how it can be mapped to the underlying resource transaction system.

A call to `close()` marks the end of a session. The main implication of `close()` is that the JDBC connection will be relinquished by the session. If you provided your own connection, `close()` returns a reference to it, so you can manually close it or return it to the pool. Otherwise `close()` returns it to the pool.

### 11.2.2. Using JTA

If your persistence layer runs in an application server (e.g. behind EJB session beans), transaction boundaries are defined in deployment descriptors. Every datasource connection obtained by Hibernate will automatically be part of a global JTA transaction. Hibernate simply joins this transaction, or if a particular session bean method has no mandatory transaction, Hibernate will tell the application server to start and end a transaction directly. (The latter should be considered a very rare case and is offered for consistency reasons. Note that your container might not allow mixed CMT and BMT behavior.)

If you set the properties `hibernate.transaction.flush_before_completion` and `hibernate.transaction.auto_close_session` to `true`, Hibernate will also automatically flush and close the `Session`



for you. The only thing left is exception handling and rollback of the database transaction. Fortunately, even this happens automatically, since an unhandled `RuntimeException` thrown by a session bean method tells the container to set the global transaction to rollback.

In other words, all you have to do in a managed environment is to get a `Session` from the `SessionFactory` (usually bound to JNDI), do your data access work, and leave the rest to the container. Transaction boundaries are set declaratively in the deployment descriptors of your session bean.

### 11.2.3. Exception handling

If the `Session` throws an exception (including any `SQLException`), you should immediately rollback the database transaction, call `Session.close()` and discard the `Session` instance. Certain methods of `Session` will *not* leave the session in a consistent state. No exception thrown by Hibernate can be treated as recoverable. Ensure that the `Session` will be closed by calling `close()` in a `finally` block.

The `HibernateException`, which wraps most of the errors that can occur in a Hibernate persistence layer, is an unchecked exception (it wasn't in older versions of Hibernate). In our opinion, we shouldn't force the application developer to catch an unrecoverable exception at a low layer. In most systems, unchecked and fatal exceptions are only caught in one of the first frames of the method call stack (i.e. in higher layers) and an error message is presented to the application user (or some other appropriate action is taken). Note that Hibernate might also throw other unchecked exceptions (e.g. when detecting stale data in version checks) which are not a `HibernateException`. These are, again, not recoverable and appropriate action should be taken.

A subtype of `HibernateException` is `SQLException` which indicate a specific set of error conditions occurring while Hibernate interacts with the database.

For `SQLExceptions` thrown while interacting with the database, Hibernate will attempt to convert the error into a subclass of `JDBCException`. The underlying `SQLException` is accessible by calling `JDBCException.getCause()`. Hibernate converts the `SQLException` into an appropriate `JDBCException` subclass based on the `SQLExceptionConverter` attached to the `SessionFactory`. By default, the `SQLExceptionConverter` is defined by the configured dialect; however, it is also possible to plug in a custom implementation (see the javadocs for the `SQLExceptionConverterFactory` class for details). The standard `JDBCException` subtypes are:

- `JDBCConnectionException` - indicates an error with the underlying JDBC communication.
- `SQLGrammarException` - indicates a grammar or syntax problem with the issued SQL.
- `ConstraintViolationException` - indicates some form of integrity constraint violation.
- `LockAcquisitionException` - indicates an error acquiring a lock level necessary to perform the requested operation.
- `GenericJDBCException` - a generic exception which did not fall into any of the other categories.

## 11.3. Optimistic concurrency control

The only approach that is consistent with high concurrency and high scalability is optimistic concurrency control with versioning. Version checking uses version numbers, or timestamps, to detect conflicting updates (and to prevent lost updates). Hibernate provides for three possible approaches to writing application code that uses optimistic concurrency. The use cases we show are in the context of long application transactions but version checking also has the benefit of preventing lost updates in single database transactions.

### 11.3.1. Application version checking

In an implementation without much help from Hibernate, each interaction with the database occurs in a new `Session` and the developer is responsible for reloading all persistent instances from the database before manip-

ulating them. This approach forces the application to carry out its own version checking to ensure application transaction isolation. This approach is the least efficient in terms of database access. It is the approach most similar to entity EJBs.

```
// foo is an instance loaded by a previous Session
session = factory.openSession();
Transaction t = session.beginTransaction();
int oldVersion = foo.getVersion();
session.load( foo, foo.getKey() ); // load the current state
if ( oldVersion!=foo.getVersion() ) throw new StaleObjectStateException();
foo.setProperty("bar");
t.commit();
session.close();
```

The `version` property is mapped using `<version>`, and Hibernate will automatically increment it during flush if the entity is dirty.

Of course, if you are operating in a low-data-concurrency environment and don't require version checking, you may use this approach and just skip the version check. In that case, *last commit wins* will be the default strategy for your long application transactions. Keep in mind that this might confuse the users of the application, as they might experience lost updates without error messages or a chance to merge conflicting changes.

Clearly, manual version checking is only feasible in very trivial circumstances and not practical for most applications. Often not only single instances, but complete graphs of modified objects have to be checked. Hibernate offers automatic version checking with either long `Session` or detached instances as the design paradigm.

### 11.3.2. Long session and automatic versioning

A single `Session` instance and its persistent instances are used for the whole application transaction. Hibernate checks instance versions at flush time, throwing an exception if concurrent modification is detected. It's up to the developer to catch and handle this exception (common options are the opportunity for the user to merge changes or to restart the business process with non-stale data).

The `Session` is disconnected from any underlying JDBC connection when waiting for user interaction. This approach is the most efficient in terms of database access. The application need not concern itself with version checking or with reattaching detached instances, nor does it have to reload instances in every database transaction.

```
// foo is an instance loaded earlier by the Session
session.reconnect(); // Obtain a new JDBC connection
Transaction t = session.beginTransaction();
foo.setProperty("bar");
t.commit(); // End database transaction, flushing the change and checking the version
session.disconnect(); // Return JDBC connection
```

The `foo` object still knows which `Session` it was loaded in. `Session.reconnect()` obtains a new connection (or you may supply one) and resumes the session. The method `Session.disconnect()` will disconnect the session from the JDBC connection and return the connection to the pool (unless you provided the connection). After reconnection, to force a version check on data you aren't updating, you may call `Session.lock()` with `LockMode.READ` on any objects that might have been updated by another transaction. You don't need to lock any data that you *are* updating.

This pattern is problematic if the `Session` is too big to be stored during user think time, e.g. an `HttpSession` should be kept as small as possible. As the `Session` is also the (mandatory) first-level cache and contains all loaded objects, we can probably use this strategy only for a few request/response cycles. This is indeed recommended, as the `Session` will soon also have stale data.

Also note that you should keep the disconnected `Session` close to the persistence layer. In other words, use an EJB stateful session bean to hold the `Session` and don't transfer it to the web layer (or even serialize it to a separate tier) to store it in the `HttpSession`.

### 11.3.3. Detached objects and automatic versioning

Each interaction with the persistent store occurs in a new `Session`. However, the same persistent instances are reused for each interaction with the database. The application manipulates the state of detached instances originally loaded in another `Session` and then reattaches them using `Session.update()`, `Session.saveOrUpdate()`, or `Session.merge()`.

```
// foo is an instance loaded by a previous Session
foo.setProperty("bar");
session = factory.openSession();
Transaction t = session.beginTransaction();
session.saveOrUpdate(foo); // Use merge() if "foo" might have been loaded already
t.commit();
session.close();
```

Again, Hibernate will check instance versions during flush, throwing an exception if conflicting updates occurred.

You may also call `lock()` instead of `update()` and use `LockMode.READ` (performing a version check, bypassing all caches) if you are sure that the object has not been modified.

### 11.3.4. Customizing automatic versioning

You may disable Hibernate's automatic version increment for particular properties and collections by setting the `optimistic-lock` mapping attribute to `false`. Hibernate will then no longer increment versions if the property is dirty.

Legacy database schemas are often static and can't be modified. Or, other applications might also access the same database and don't know how to handle version numbers or even timestamps. In both cases, versioning can't rely on a particular column in a table. To force a version check without a version or timestamp property mapping, with a comparison of the state of all fields in a row, turn on `optimistic-lock="all"` in the `<class>` mapping. Note that this conceptually only works if Hibernate can compare the old and new state, i.e. if you use a single long `Session` and not `session-per-request-with-detached-objects`.

Sometimes concurrent modification can be permitted as long as the changes that have been made don't overlap. If you set `optimistic-lock="dirty"` when mapping the `<class>`, Hibernate will only compare dirty fields during flush.

In both cases, with dedicated version/timestamp columns or with full/dirty field comparison, Hibernate uses a single `UPDATE` statement (with an appropriate `WHERE` clause) per entity to execute the version check and update the information. If you use transitive persistence to cascade reattachment to associated entities, Hibernate might execute unnecessary updates. This is usually not a problem, but *on update* triggers in the database might be executed even when no changes have been made to detached instances. You can customize this behavior by setting `select-before-update="true"` in the `<class>` mapping, forcing Hibernate to `SELECT` the instance to ensure that changes did actually occur, before updating the row.

## 11.4. Pessimistic Locking

It is not intended that users spend much time worrying about locking strategies. It's usually enough to specify an isolation level for the JDBC connections and then simply let the database do all the work. However, advanced users may sometimes wish to obtain exclusive pessimistic locks, or re-obtain locks at the start of a new transaction.

Hibernate will always use the locking mechanism of the database, never lock objects in memory!

The `LockMode` class defines the different lock levels that may be acquired by Hibernate. A lock is obtained by the following mechanisms:

- `LockMode.WRITE` is acquired automatically when Hibernate updates or inserts a row.
- `LockMode.UPGRADE` may be acquired upon explicit user request using `SELECT ... FOR UPDATE` on databases which support that syntax.
- `LockMode.UPGRADE_NOWAIT` may be acquired upon explicit user request using a `SELECT ... FOR UPDATE NOWAIT` under Oracle.
- `LockMode.READ` is acquired automatically when Hibernate reads data under Repeatable Read or Serializable isolation level. May be re-acquired by explicit user request.
- `LockMode.NONE` represents the absence of a lock. All objects switch to this lock mode at the end of a Transaction. Objects associated with the session via a call to `update()` or `saveOrUpdate()` also start out in this lock mode.

The "explicit user request" is expressed in one of the following ways:

- A call to `Session.load()`, specifying a `LockMode`.
- A call to `Session.lock()`.
- A call to `Query.setLockMode()`.

If `Session.load()` is called with `UPGRADE` or `UPGRADE_NOWAIT`, and the requested object was not yet loaded by the session, the object is loaded using `SELECT ... FOR UPDATE`. If `load()` is called for an object that is already loaded with a less restrictive lock than the one requested, Hibernate calls `lock()` for that object.

`Session.lock()` performs a version number check if the specified lock mode is `READ`, `UPGRADE` or `UPGRADE_NOWAIT`. (In the case of `UPGRADE` or `UPGRADE_NOWAIT`, `SELECT ... FOR UPDATE` is used.)

If the database does not support the requested lock mode, Hibernate will use an appropriate alternate mode (instead of throwing an exception). This ensures that applications will be portable.

---

## Chapter 12. Interceptors and events

It is often useful for the application to react to certain events that occur inside Hibernate. This allows implementation of certain kinds of generic functionality, and extension of Hibernate functionality.

### 12.1. Interceptors

The `Interceptor` interface provides callbacks from the session to the application allowing the application to inspect and/or manipulate properties of a persistent object before it is saved, updated, deleted or loaded. One possible use for this is to track auditing information. For example, the following `Interceptor` automatically sets the `createTimestamp` when an `Auditable` is created and updates the `lastUpdateTimestamp` property when an `Auditable` is updated.

```
package org.hibernate.test;

import java.io.Serializable;
import java.util.Date;
import java.util.Iterator;

import org.hibernate.Interceptor;
import org.hibernate.type.Type;

public class AuditInterceptor implements Interceptor, Serializable {

    private int updates;
    private int creates;

    public void onDelete(Object entity,
                        Serializable id,
                        Object[] state,
                        String[] propertyNames,
                        Type[] types) {

        // do nothing
    }

    public boolean onFlushDirty(Object entity,
                              Serializable id,
                              Object[] currentState,
                              Object[] previousState,
                              String[] propertyNames,
                              Type[] types) {

        if ( entity instanceof Auditable ) {
            updates++;
            for ( int i=0; i < propertyNames.length; i++ ) {
                if ( "lastUpdateTimestamp".equals( propertyNames[i] ) ) {
                    currentState[i] = new Date();
                    return true;
                }
            }
        }
        return false;
    }

    public boolean onLoad(Object entity,
                        Serializable id,
                        Object[] state,
                        String[] propertyNames,
                        Type[] types) {

        return false;
    }

    public boolean onSave(Object entity,
                        Serializable id,
```

```

        Object[] state,
        String[] propertyNames,
        Type[] types) {

    if ( entity instanceof Auditable ) {
        creates++;
        for ( int i=0; i<propertyNames.length; i++ ) {
            if ( "createTimestamp".equals( propertyNames[i] ) ) {
                state[i] = new Date();
                return true;
            }
        }
    }
    return false;
}

public void postFlush(Iterator entities) {
    System.out.println("Creations: " + creates + ", Updates: " + updates);
}

public void preFlush(Iterator entities) {
    updates=0;
    creates=0;
}

...
}

```

The interceptor would be specified when a session is created.

```
Session session = sf.openSession( new AuditInterceptor() );
```

You may also set an interceptor on a global level, using the Configuration:

```
new Configuration().setInterceptor( new AuditInterceptor() );
```

## 12.2. Event system

If you have to react to particular events in your persistence layer, you may also use the Hibernate3 *event* architecture. The event system can be used in addition or as a replacement for interceptors.

Essentially all of the methods of the `Session` interface correlate to an event. You have a `LoadEvent`, a `FlushEvent`, etc (consult the XML configuration-file DTD or the `org.hibernate.event` package for the full list of defined event types). When a request is made of one of these methods, the Hibernate `Session` generates an appropriate event and passes it to the configured event listener for that type. Out-of-the-box, these listeners implement the same processing in which those methods always resulted. However, you are free to implement a customization of one of the listener interfaces (i.e., the `LoadEvent` is processed by the registered implementation of the `LoadEventListener` interface), in which case their implementation would be responsible for processing any `load()` requests made of the `Session`.

The listeners should be considered effectively singletons; meaning, they are shared between requests, and thus should not save any state as instance variables.

A custom listener should implement the appropriate interface for the event it wants to process and/or extend one of the convenience base classes (or even the default event listeners used by Hibernate out-of-the-box as these are declared non-final for this purpose). Custom listeners can either be registered programmatically through the `Configuration` object, or specified in the Hibernate configuration XML (declarative configuration through the properties file is not supported). Here's an example of a custom load event listener:

```

public class MyLoadListener extends DefaultLoadEventListener {
    // this is the single method defined by the LoadEventListener interface
    public Object onLoad(LoadEvent event, LoadEventListener.LoadType loadType)
        throws HibernateException {
        if ( !MySecurity.isAuthorized( event.getEntityClassName(), event.getEntityId() ) ) {
            throw MySecurityException("Unauthorized access");
        }
        return super.onLoad(event, loadType);
    }
}

```

You also need a configuration entry telling Hibernate to use the listener instead of the default listener:

```

<hibernate-configuration>
  <session-factory>
    ...
    <listener type="load" class="MyLoadListener"/>
  </session-factory>
</hibernate-configuration>

```

Instead, you may register it programmatically:

```

Configuration cfg = new Configuration();
cfg.getSessionEventListenerConfig().setLoadEventListener( new MyLoadListener() );

```

Listeners registered declaratively cannot share instances. If the same class name is used in multiple `<listener/>` elements, each reference will result in a separate instance of that class. If you need the capability to share listener instances between listener types you must use the programmatic registration approach.

Why implement an interface and define the specific type during configuration? Well, a listener implementation could implement multiple event listener interfaces. Having the type additionally defined during registration makes it easier to turn custom listeners on or off during configuration.

## 12.3. Hibernate declarative security

Usually, declarative security in Hibernate applications is managed in a session facade layer. Now, Hibernate3 allows certain actions to be permissioned via JACC, and authorized via JAAS. This is optional functionality built on top of the event architecture.

First, you must configure the appropriate event listeners, to enable the use of JAAS authorization.

```

<listener type="pre-delete" class="org.hibernate.secure.JACCPreDeleteEventListener"/>
<listener type="pre-update" class="org.hibernate.secure.JACCPreUpdateEventListener"/>
<listener type="pre-insert" class="org.hibernate.secure.JACCPreInsertEventListener"/>
<listener type="pre-load" class="org.hibernate.secure.JACCPreLoadEventListener"/>

```

Next, still in `hibernate.cfg.xml`, bind the permissions to roles:

```

<grant role="admin" entity-name="User" actions="insert,update,read"/>
<grant role="su" entity-name="User" actions="*/>

```

The role names are the roles understood by your JACC provider.

---

## Chapter 13. Batch processing

A naive approach to inserting 100 000 rows in the database using Hibernate might look like this:

```
Session session = sessionFactory.openSession();
Transaction tx = session.beginTransaction();
for ( int i=0; i<100000; i++ ) {
    Customer customer = new Customer(.....);
    session.save(customer);
}
tx.commit();
session.close();
```

This would fall over with an `OutOfMemoryException` somewhere around the 50 000th row. That's because Hibernate caches all the newly inserted `Customer` instances in the session-level cache.

In this chapter we'll show you how to avoid this problem. First, however, if you are doing batch processing, it is absolutely critical that you enable the use of JDBC batching, if you intend to achieve reasonable performance. Set the JDBC batch size to a reasonable number (say, 10-50):

```
hibernate.jdbc.batch_size 20
```

You also might like to do this kind of work in a process where interaction with the second-level cache is completely disabled:

```
hibernate.cache.use_second_level_cache false
```

### 13.1. Batch inserts

When making new objects persistent, you must `flush()` and then `clear()` the session regularly, to control the size of the first-level cache.

```
Session session = sessionFactory.openSession();
Transaction tx = session.beginTransaction();

for ( int i=0; i<100000; i++ ) {
    Customer customer = new Customer(.....);
    session.save(customer);
    if ( i % 20 == 0 ) { //20, same as the JDBC batch size
        //flush a batch of inserts and release memory:
        session.flush();
        session.clear();
    }
}

tx.commit();
session.close();
```

### 13.2. Batch updates

For retrieving and updating data the same ideas apply. In addition, you need to use `scroll()` to take advantage of server-side cursors for queries that return many rows of data.

```
Session session = sessionFactory.openSession();
Transaction tx = session.beginTransaction();
```



```
ScrollableResults customers = session.getNamedQuery("GetCustomers")
    .setCacheMode(CacheMode.IGNORE)
    .scroll(ScrollMode.FORWARD_ONLY);
int count=0;
while ( customers.next() ) {
    Customer customer = (Customer) customers.get(0);
    customer.updateStuff(...);
    if ( ++count % 20 == 0 ) {
        //flush a batch of updates and release memory:
        session.flush();
        session.clear();
    }
}

tx.commit();
session.close();
```

---

# Chapter 14. HQL: The Hibernate Query Language

Hibernate is equipped with an extremely powerful query language that (quite intentionally) looks very much like SQL. But don't be fooled by the syntax; HQL is fully object-oriented, understanding notions like inheritance, polymorphism and association.

## 14.1. Case Sensitivity

Queries are case-insensitive, except for names of Java classes and properties. So `seLeCT` is the same as `select` is the same as `SELECT` but `org.hibernate.eg.FOO` is not `org.hibernate.eg.Foo` and `foo.barSet` is not `foo.BARSET`.

This manual uses lowercase HQL keywords. Some users find queries with uppercase keywords more readable, but we find this convention ugly when embedded in Java code.

## 14.2. The from clause

The simplest possible Hibernate query is of the form:

```
from eg.Cat
```

which simply returns all instances of the class `eg.Cat`. We don't usually need to qualify the class name, since `auto-import` is the default. So we almost always just write:

```
from Cat
```

Most of the time, you will need to assign an *alias*, since you will want to refer to the `Cat` in other parts of the query.

```
from Cat as cat
```

This query assigns the alias `cat` to `Cat` instances, so we could use that alias later in the query. The `as` keyword is optional; we could also write:

```
from Cat cat
```

Multiple classes may appear, resulting in a cartesian product or "cross" join.

```
from Formula, Parameter
```

```
from Formula as form, Parameter as param
```

It is considered good practice to name query aliases using an initial lowercase, consistent with Java naming standards for local variables (eg. `domesticCat`).

## 14.3. Associations and joins

We may also assign aliases to associated entities, or even to elements of a collection of values, using a `join`.

```
from Cat as cat
    inner join cat.mate as mate
    left outer join cat.kittens as kitten
```

```
from Cat as cat left join cat.mate.kittens as kittens
```

```
from Formula form full join form.parameter param
```

The supported join types are borrowed from ANSI SQL

- inner join
- left outer join
- right outer join
- full join (not usually useful)

The inner join, left outer join and right outer join constructs may be abbreviated.

```
from Cat as cat
    join cat.mate as mate
    left join cat.kittens as kitten
```

In addition, a "fetch" join allows associations or collections of values to be initialized along with their parent objects, using a single select. This is particularly useful in the case of a collection. It effectively overrides the outer join and lazy declarations of the mapping file for associations and collections. See Section 19.1, "Fetching strategies" for more information.

```
from Cat as cat
    inner join fetch cat.mate
    left join fetch cat.kittens
```

A fetch join does not usually need to assign an alias, because the associated objects should not be used in the where clause (or any other clause). Also, the associated objects are not returned directly in the query results. Instead, they may be accessed via the parent object.

Note that, in the current implementation, only one collection role may be fetched in a query (everything else would be non-performant). Note also that the `fetch` construct may not be used in queries called using `scroll()` or `iterate()`. Finally, note that `full join fetch` and `right join fetch` are not meaningful.

## 14.4. The select clause

The `select` clause picks which objects and properties to return in the query result set. Consider:

```
select mate
from Cat as cat
    inner join cat.mate as mate
```

The query will select `mates` of other `Cats`. Actually, you may express this query more compactly as:

```
select cat.mate from Cat cat
```

Queries may return properties of any value type including properties of component type:

```
select cat.name from DomesticCat cat
where cat.name like 'fri%'
```

```
select cust.name.firstName from Customer as cust
```

Queries may return multiple objects and/or properties as an array of type `Object[]`

```
select mother, offspr, mate.name
from DomesticCat as mother
     inner join mother.mate as mate
     left outer join mother.kittens as offspr
```

or as an actual typesafe Java object

```
select new Family(mother, mate, offspr)
from DomesticCat as mother
     join mother.mate as mate
     left join mother.kittens as offspr
```

assuming that the class `Family` has an appropriate constructor.

## 14.5. Aggregate functions

HQL queries may even return the results of aggregate functions on properties:

```
select avg(cat.weight), sum(cat.weight), max(cat.weight), count(cat)
from Cat cat
```

The supported aggregate functions are

- `avg(...)`, `sum(...)`, `min(...)`, `max(...)`
- `count(*)`
- `count(...)`, `count(distinct ...)`, `count(all...)`

The `distinct` and `all` keywords may be used and have the same semantics as in SQL.

```
select distinct cat.name from Cat cat

select count(distinct cat.name), count(cat) from Cat cat
```

## 14.6. Polymorphic queries

A query like:

```
from Cat as cat
```

returns instances not only of `Cat`, but also of subclasses like `DomesticCat`. Hibernate queries may name *any* Java class or interface in the `from` clause. The query will return instances of all persistent classes that extend that class or implement the interface. The following query would return all persistent objects:

```
from java.lang.Object o
```

The interface `Named` might be implemented by various persistent classes:

```
from Named n, Named m where n.name = m.name
```

Note that these last two queries will require more than one SQL `SELECT`. This means that the `order by` clause does not correctly order the whole result set. (It also means you can't call these queries using `Query.scroll()`.)

## 14.7. The where clause

The `where` clause allows you to narrow the list of instances returned. If no alias exists, you may refer to properties by name:

```
from Cat where name='Fritz'
```

If there is an alias, use a qualified property name:

```
from Cat as cat where cat.name='Fritz'
```

returns instances of `Cat` named 'Fritz'.

```
select foo
from Foo foo, Bar bar
where foo.startDate = bar.date
```

will return all instances of `Foo` for which there exists an instance of `bar` with a `date` property equal to the `startDate` property of the `Foo`. Compound path expressions make the `where` clause extremely powerful. Consider:

```
from Cat cat where cat.mate.name is not null
```

This query translates to an SQL query with a table (inner) join. If you were to write something like

```
from Foo foo
where foo.bar.baz.customer.address.city is not null
```

you would end up with a query that would require four table joins in SQL.

The `=` operator may be used to compare not only properties, but also instances:

```
from Cat cat, Cat rival where cat.mate = rival.mate
```

```
select cat, mate
from Cat cat, Cat mate
where cat.mate = mate
```

The special property (lowercase) `id` may be used to reference the unique identifier of an object. (You may also use its property name.)

```
from Cat as cat where cat.id = 123

from Cat as cat where cat.mate.id = 69
```

The second query is efficient. No table join is required!

Properties of composite identifiers may also be used. Suppose `Person` has a composite identifier consisting of `country` and `medicareNumber`.

```
from bank.Person person
where person.id.country = 'AU'
```

```
and person.id.medicareNumber = 123456
```

```
from bank.Account account
where account.owner.id.country = 'AU'
      and account.owner.id.medicareNumber = 123456
```

Once again, the second query requires no table join.

Likewise, the special property `class` accesses the discriminator value of an instance in the case of polymorphic persistence. A Java class name embedded in the where clause will be translated to its discriminator value.

```
from Cat cat where cat.class = DomesticCat
```

You may also specify properties of components or composite user types (and of components of components, etc). Never try to use a path-expression that ends in a property of component type (as opposed to a property of a component). For example, if `store.owner` is an entity with a component `address`

```
store.owner.address.city    // okay
store.owner.address         // error!
```

An "any" type has the special properties `id` and `class`, allowing us to express a join in the following way (where `AuditLog.item` is a property mapped with `<any>`).

```
from AuditLog log, Payment payment
where log.item.class = 'Payment' and log.item.id = payment.id
```

Notice that `log.item.class` and `payment.class` would refer to the values of completely different database columns in the above query.

## 14.8. Expressions

Expressions allowed in the `where` clause include most of the kind of things you could write in SQL:

- mathematical operators `+`, `-`, `*`, `/`
- binary comparison operators `=`, `>=`, `<=`, `<>`, `!=`, `like`
- logical operations `and`, `or`, `not`
- string concatenation `||`
- SQL scalar functions like `upper()` and `lower()`
- `current_date()`, `current_time()`, `current_timestamp()`
- Any function or operator defined by EJB-QL 3.0
- `case when ... then ... else ... end`
- Parentheses `( )` indicate grouping
- `in`, `not in`, `between`, `is null` `is not null` `is empty` `is not empty` `member of` `not member of`
- JDBC IN parameters `?`
- named parameters `:name`, `:start_date`, `:x1`
- SQL literals `'foo'`, `69`, `'1970-01-01 10:00:01.0'`
- Java public static final constants eg. `Color.TABBY`

`in` and `between` may be used as follows:

```
from DomesticCat cat where cat.name between 'A' and 'B'
```

```
from DomesticCat cat where cat.name in ( 'Foo', 'Bar', 'Baz' )
```

and the negated forms may be written

```
from DomesticCat cat where cat.name not between 'A' and 'B'
```

```
from DomesticCat cat where cat.name not in ( 'Foo', 'Bar', 'Baz' )
```

Likewise, `is null` and `is not null` may be used to test for null values.

Booleans may be easily used in expressions by declaring HQL query substitutions in Hibernate configuration:

```
<property name="hibernate.query.substitutions">true 1, false 0</property>
```

This will replace the keywords `true` and `false` with the literals `1` and `0` in the translated SQL from this HQL:

```
from Cat cat where cat.alive = true
```

You may test the size of a collection with the special property `size`, or the special `size()` function.

```
from Cat cat where cat.kittens.size > 0
```

```
from Cat cat where size(cat.kittens) > 0
```

For indexed collections, you may refer to the minimum and maximum indices using `minindex` and `maxindex` functions. Similarly, you may refer to the minimum and maximum elements of a collection of basic type using the `minelement` and `maxelement` functions.

```
from Calendar cal where maxelement(cal.holidays) > current date
```

```
from Order order where maxindex(order.items) > 100
```

```
from Order order where minelement(order.items) > 10000
```

The SQL functions `any`, `some`, `all`, `exists`, `in` are supported when passed the element or index set of a collection (`elements` and `indices` functions) or the result of a subquery (see below).

```
select mother from Cat as mother, Cat as kit
where kit in elements(foo.kittens)
```

```
select p from NameList list, Person p
where p.name = some elements(list.names)
```

```
from Cat cat where exists elements(cat.kittens)
```

```
from Player p where 3 > all elements(p.scores)
```

```
from Show show where 'fizard' in indices(show.acts)
```

Note that these constructs - `size`, `elements`, `indices`, `minindex`, `maxindex`, `minelement`, `maxelement` - may only be used in the `where` clause in `Hibernate3`.

Elements of indexed collections (arrays, lists, maps) may be referred to by index (in a `where` clause only):

```
from Order order where order.items[0].id = 1234
```

```
select person from Person person, Calendar calendar
where calendar.holidays['national day'] = person.birthDay
and person.nationality.calendar = calendar
```

```
select item from Item item, Order order
where order.items[ order.deliveredItemIndices[0] ] = item and order.id = 11
```

```
select item from Item item, Order order
where order.items[ maxindex(order.items) ] = item and order.id = 11
```

The expression inside `[]` may even be an arithmetic expression.

```
select item from Item item, Order order
where order.items[ size(order.items) - 1 ] = item
```

HQL also provides the built-in `index()` function, for elements of a one-to-many association or collection of values.

```
select item, index(item) from Order order
join order.items item
where index(item) < 5
```

Scalar SQL functions supported by the underlying database may be used

```
from DomesticCat cat where upper(cat.name) like 'FRI%'
```

If you are not yet convinced by all this, think how much longer and less readable the following query would be in SQL:

```
select cust
from Product prod,
Store store
inner join store.customers cust
where prod.name = 'widget'
and store.location.name in ( 'Melbourne', 'Sydney' )
and prod = all elements(cust.currentOrder.lineItems)
```

*Hint:* something like

```
SELECT cust.name, cust.address, cust.phone, cust.id, cust.current_order
FROM customers cust,
stores store,
locations loc,
store_customers sc,
product prod
WHERE prod.name = 'widget'
AND store.loc_id = loc.id
AND loc.name IN ( 'Melbourne', 'Sydney' )
AND sc.store_id = store.id
AND sc.cust_id = cust.id
AND prod.id = ALL(
SELECT item.prod_id
FROM line_items item, orders o
WHERE item.order_id = o.id
AND cust.current_order = o.id
)
```

## 14.9. The order by clause



The list returned by a query may be ordered by any property of a returned class or components:

```
from DomesticCat cat
order by cat.name asc, cat.weight desc, cat.birthdate
```

The optional `asc` or `desc` indicate ascending or descending order respectively.

## 14.10. The group by clause

A query that returns aggregate values may be grouped by any property of a returned class or components:

```
select cat.color, sum(cat.weight), count(cat)
from Cat cat
group by cat.color
```

```
select foo.id, avg(name), max(name)
from Foo foo join foo.names name
group by foo.id
```

A `having` clause is also allowed.

```
select cat.color, sum(cat.weight), count(cat)
from Cat cat
group by cat.color
having cat.color in (eg.Color.TABBY, eg.Color.BLACK)
```

SQL functions and aggregate functions are allowed in the `having` and `order by` clauses, if supported by the underlying database (eg. not in MySQL).

```
select cat
from Cat cat
    join cat.kittens kitten
group by cat
having avg(kitten.weight) > 100
order by count(kitten) asc, sum(kitten.weight) desc
```

Note that neither the `group by` clause nor the `order by` clause may contain arithmetic expressions.

## 14.11. Subqueries

For databases that support subselects, Hibernate supports subqueries within queries. A subquery must be surrounded by parentheses (often by an SQL aggregate function call). Even correlated subqueries (subqueries that refer to an alias in the outer query) are allowed.

```
from Cat as fatcat
where fatcat.weight > (
    select avg(cat.weight) from DomesticCat cat
)
```

```
from DomesticCat as cat
where cat.name = some (
    select name.nickName from Name as name
)
```

```
from Cat as cat
where not exists (
```

```

    from Cat as mate where mate.mate = cat
)

```

```

from DomesticCat as cat
where cat.name not in (
    select name.nickName from Name as name
)

```

## 14.12. HQL examples

Hibernate queries can be quite powerful and complex. In fact, the power of the query language is one of Hibernate's main selling points. Here are some example queries very similar to queries that I used on a recent project. Note that most queries you will write are much simpler than these!

The following query returns the order id, number of items and total value of the order for all unpaid orders for a particular customer and given minimum total value, ordering the results by total value. In determining the prices, it uses the current catalog. The resulting SQL query, against the `ORDER`, `ORDER_LINE`, `PRODUCT`, `CATALOG` and `PRICE` tables has four inner joins and an (uncorrelated) subselect.

```

select order.id, sum(price.amount), count(item)
from Order as order
    join order.lineItems as item
    join item.product as product,
    Catalog as catalog
    join catalog.prices as price
where order.paid = false
    and order.customer = :customer
    and price.product = product
    and catalog.effectiveDate < sysdate
    and catalog.effectiveDate >= all (
        select cat.effectiveDate
        from Catalog as cat
        where cat.effectiveDate < sysdate
    )
group by order
having sum(price.amount) > :minAmount
order by sum(price.amount) desc

```

What a monster! Actually, in real life, I'm not very keen on subqueries, so my query was really more like this:

```

select order.id, sum(price.amount), count(item)
from Order as order
    join order.lineItems as item
    join item.product as product,
    Catalog as catalog
    join catalog.prices as price
where order.paid = false
    and order.customer = :customer
    and price.product = product
    and catalog = :currentCatalog
group by order
having sum(price.amount) > :minAmount
order by sum(price.amount) desc

```

The next query counts the number of payments in each status, excluding all payments in the `AWAITING_APPROVAL` status where the most recent status change was made by the current user. It translates to an SQL query with two inner joins and a correlated subselect against the `PAYMENT`, `PAYMENT_STATUS` and `PAYMENT_STATUS_CHANGE` tables.

```

select count(payment), status.name
from Payment as payment

```

```

    join payment.currentStatus as status
    join payment.statusChanges as statusChange
where payment.status.name <> PaymentStatus.AWAITING_APPROVAL
    or (
        statusChange.timeStamp = (
            select max(change.timeStamp)
            from PaymentStatusChange change
            where change.payment = payment
        )
        and statusChange.user <> :currentUser
    )
group by status.name, status.sortOrder
order by status.sortOrder

```

If I would have mapped the `statusChanges` collection as a list, instead of a set, the query would have been much simpler to write.

```

select count(payment), status.name
from Payment as payment
    join payment.currentStatus as status
where payment.status.name <> PaymentStatus.AWAITING_APPROVAL
    or payment.statusChanges[ maxIndex(payment.statusChanges) ].user <> :currentUser
group by status.name, status.sortOrder
order by status.sortOrder

```

The next query uses the MS SQL Server `isNull()` function to return all the accounts and unpaid payments for the organization to which the current user belongs. It translates to an SQL query with three inner joins, an outer join and a subselect against the `ACCOUNT`, `PAYMENT`, `PAYMENT_STATUS`, `ACCOUNT_TYPE`, `ORGANIZATION` and `ORG_USER` tables.

```

select account, payment
from Account as account
    left outer join account.payments as payment
where :currentUser in elements(account.holder.users)
    and PaymentStatus.UNPAID = isNull(payment.currentStatus.name, PaymentStatus.UNPAID)
order by account.type.sortOrder, account.accountNumber, payment.dueDate

```

For some databases, we would need to do away with the (correlated) subselect.

```

select account, payment
from Account as account
    join account.holder.users as user
    left outer join account.payments as payment
where :currentUser = user
    and PaymentStatus.UNPAID = isNull(payment.currentStatus.name, PaymentStatus.UNPAID)
order by account.type.sortOrder, account.accountNumber, payment.dueDate

```

## 14.13. Tips & Tricks

You can count the number of query results without actually returning them:

```

( (Integer) session.iterate("select count(*) from ...").next() ).intValue()

```

To order a result by the size of a collection, use the following query:

```

select usr.id, usr.name
from User as usr
    left join usr.messages as msg
group by usr.id, usr.name
order by count(msg)

```

If your database supports subselects, you can place a condition upon selection size in the where clause of your query:

```
from User usr where size(usr.messages) >= 1
```

If your database doesn't support subselects, use the following query:

```
select usr.id, usr.name
from User usr
      join usr.messages msg
group by usr.id, usr.name
having count(msg) >= 1
```

As this solution can't return a `User` with zero messages because of the inner join, the following form is also useful:

```
select usr.id, usr.name
from User as usr
      left join usr.messages as msg
group by usr.id, usr.name
having count(msg) = 0
```

Properties of a `JavaBean` can be bound to named query parameters:

```
Query q = s.createQuery("from foo Foo as foo where foo.name=:name and foo.size=:size");
q.setProperties(fooBean); // fooBean has getName() and getSize()
List foos = q.list();
```

Collections are pageable by using the `Query` interface with a filter:

```
Query q = s.createFilter( collection, "" ); // the trivial filter
q.setMaxResults(PAGE_SIZE);
q.setFirstResult(PAGE_SIZE * pageNumber);
List page = q.list();
```

Collection elements may be ordered or grouped using a query filter:

```
Collection orderedCollection = s.filter( collection, "order by this.amount" );
Collection counts = s.filter( collection, "select this.type, count(this) group by this.type" );
```

You can find the size of a collection without initializing it:

```
( (Integer) session.iterate("select count(*) from ....").next() ).intValue();
```

---

# Chapter 15. Criteria Queries

Hibernate features an intuitive, extensible criteria query API.

## 15.1. Creating a criteria instance

The interface `org.hibernate.Criteria` represents a query against a particular persistent class. The `Session` is a factory for `Criteria` instances.

```
Criteria crit = sess.createCriteria(Cat.class);
crit.setMaxResults(50);
List cats = crit.list();
```

## 15.2. Narrowing the result set

An individual query criterion is an instance of the interface `org.hibernate.criterion.Criterion`. The class `org.hibernate.criterion.Restrictions` defines factory methods for obtaining certain built-in `Criterion` types.

```
List cats = sess.createCriteria(Cat.class)
    .add( Restrictions.like("name", "Fritz%") )
    .add( Restrictions.between("weight", minWeight, maxWeight) )
    .list();
```

Restrictions may be grouped logically.

```
List cats = sess.createCriteria(Cat.class)
    .add( Restrictions.like("name", "Fritz%") )
    .add( Restrictions.or(
        Restrictions.eq( "age", new Integer(0) ),
        Restrictions.isNull("age")
    ) )
    .list();
```

```
List cats = sess.createCriteria(Cat.class)
    .add( Restrictions.in( "name", new String[] { "Fritz", "Izi", "Pk" } ) )
    .add( Restrictions.disjunction()
        .add( Restrictions.isNull("age") )
        .add( Restrictions.eq("age", new Integer(0) ) )
        .add( Restrictions.eq("age", new Integer(1) ) )
        .add( Restrictions.eq("age", new Integer(2) ) )
    ) )
    .list();
```

There are quite a range of built-in criterion types (`Restrictions` subclasses), but one that is especially useful lets you specify SQL directly.

```
List cats = sess.createCriteria(Cat.class)
    .add( Restrictions.sql("lower({alias}.name) like lower(?)", "Fritz%", Hibernate.STRING) )
    .list();
```

The `{alias}` placeholder will be replaced by the row alias of the queried entity.

An alternative approach to obtaining a criterion is to get it from a `Property` instance. You can create a `Property` by calling `Property.forName()`.

```
Property age = Property.forName("age");
List cats = sess.createCriteria(Cat.class)
    .add( Restrictions.disjunction()
        .add( age.isNull() )
        .add( age.eq( new Integer(0) ) )
        .add( age.eq( new Integer(1) ) )
        .add( age.eq( new Integer(2) ) )
    ) )
    .add( Property.forName("name").in( new String[] { "Fritz", "Izi", "Pk" } ) )
    .list();
```

## 15.3. Ordering the results

You may order the results using `org.hibernate.criterion.Order`.

```
List cats = sess.createCriteria(Cat.class)
    .add( Restrictions.like("name", "F%")
    .addOrder( Order.asc("name") )
    .addOrder( Order.desc("age") )
    .setMaxResults(50)
    .list();
```

```
List cats = sess.createCriteria(Cat.class)
    .add( Property.forName("name").like("F%") )
    .addOrder( Property.forName("name").asc() )
    .addOrder( Property.forName("age").desc() )
    .setMaxResults(50)
    .list();
```

## 15.4. Associations

You may easily specify constraints upon related entities by navigating associations using `createCriteria()`.

```
List cats = sess.createCriteria(Cat.class)
    .add( Restrictions.like("name", "F%")
    .createCriteria("kittens")
        .add( Restrictions.like("name", "F%")
    .list();
```

note that the second `createCriteria()` returns a new instance of `Criteria`, which refers to the elements of the kittens collection.

The following, alternate form is useful in certain circumstances.

```
List cats = sess.createCriteria(Cat.class)
    .createAlias("kittens", "kt")
    .createAlias("mate", "mt")
    .add( Restrictions.eqProperty("kt.name", "mt.name") )
    .list();
```

(`createAlias()` does not create a new instance of `Criteria`.)

Note that the kittens collections held by the `Cat` instances returned by the previous two queries are *not* pre-filtered by the criteria! If you wish to retrieve just the kittens that match the criteria, you must use `returnMaps()`.

```
List cats = sess.createCriteria(Cat.class)
```

```

        .createCriteria("kittens", "kt")
        .add( Restrictions.eq("name", "F%") )
        .returnMaps()
        .list();
Iterator iter = cats.iterator();
while ( iter.hasNext() ) {
    Map map = (Map) iter.next();
    Cat cat = (Cat) map.get(Criteria.ROOT_ALIAS);
    Cat kitten = (Cat) map.get("kt");
}

```

## 15.5. Dynamic association fetching

You may specify association fetching semantics at runtime using `setFetchMode()`.

```

List cats = sess.createCriteria(Cat.class)
    .add( Restrictions.like("name", "Fritz%") )
    .setFetchMode("mate", FetchMode.EAGER)
    .setFetchMode("kittens", FetchMode.EAGER)
    .list();

```

This query will fetch both `mate` and `kittens` by outer join. See Section 19.1, “Fetching strategies” for more information.

## 15.6. Example queries

The class `org.hibernate.criterion.Example` allows you to construct a query criterion from a given instance.

```

Cat cat = new Cat();
cat.setSex('F');
cat.setColor(Color.BLACK);
List results = session.createCriteria(Cat.class)
    .add( Example.create(cat) )
    .list();

```

Version properties, identifiers and associations are ignored. By default, null valued properties are excluded.

You can adjust how the `Example` is applied.

```

Example example = Example.create(cat)
    .excludeZeroes()           //exclude zero valued properties
    .excludeProperty("color")  //exclude the property named "color"
    .ignoreCase()              //perform case insensitive string comparisons
    .enableLike();             //use like for string comparisons
List results = session.createCriteria(Cat.class)
    .add(example)
    .list();

```

You can even use examples to place criteria upon associated objects.

```

List results = session.createCriteria(Cat.class)
    .add( Example.create(cat) )
    .createCriteria("mate")
        .add( Example.create( cat.getMate() ) )
    .list();

```

## 15.7. Projections, aggregation and grouping

The class `org.hibernate.criterion.Projections` is a factory for `Projection` instances. We apply a projection to a query by calling `setProjection()`.

```
List results = session.createCriteria(Cat.class)
    .setProjection( Projections.rowCount() )
    .add( Restrictions.eq("color", Color.BLACK) )
    .list();
```

```
List results = session.createCriteria(Cat.class)
    .setProjection( Projections.projectionList()
        .add( Projections.rowCount() )
        .add( Projections.avg("weight") )
        .add( Projections.max("weight") )
        .add( Projections.groupProperty("color") )
    )
    .list();
```

There is no explicit "group by" necessary in a criteria query. Certain projection types are defined to be *grouping projections*, which also appear in the SQL `group by` clause.

An alias may optionally be assigned to a projection, so that the projected value may be referred to in restrictions or orderings. Here are two different ways to do this:

```
List results = session.createCriteria(Cat.class)
    .setProjection( Projections.alias( Projections.groupProperty("color"), "colr" ) )
    .addOrder( Order.asc("colr") )
    .list();
```

```
List results = session.createCriteria(Cat.class)
    .setProjection( Projections.groupProperty("color").as("colr") )
    .addOrder( Order.asc("colr") )
    .list();
```

The `alias()` and `as()` methods simply wrap a projection instance in another, aliased, instance of `Projection`. As a shortcut, you can assign an alias when you add the projection to a projection list:

```
List results = session.createCriteria(Cat.class)
    .setProjection( Projections.projectionList()
        .add( Projections.rowCount(), "catCountByColor" )
        .add( Projections.avg("weight"), "avgWeight" )
        .add( Projections.max("weight"), "maxWeight" )
        .add( Projections.groupProperty("color"), "color" )
    )
    .addOrder( Order.desc("catCountByColor") )
    .addOrder( Order.desc("avgWeight") )
    .list();
```

```
List results = session.createCriteria(Domestic.class, "cat")
    .createAlias("kittens", "kit")
    .setProjection( Projections.projectionList()
        .add( Projections.property("cat.name"), "catName" )
        .add( Projections.property("kit.name"), "kitName" )
    )
    .addOrder( Order.asc("catName") )
    .addOrder( Order.asc("kitName") )
    .list();
```

You can also use `Property.forName()` to express projections:



```
List results = session.createCriteria(Cat.class)
    .setProjection( Property.forName("name") )
    .add( Property.forName("color").eq(Color.BLACK) )
    .list();
```

```
List results = session.createCriteria(Cat.class)
    .setProjection( Projections.projectionList()
        .add( Projections.rowCount().as("catCountByColor") )
        .add( Property.forName("weight").avg().as("avgWeight") )
        .add( Property.forName("weight").max().as("maxWeight") )
        .add( Property.forName("color").group().as("color" )
    )
    .addOrder( Order.desc("catCountByColor") )
    .addOrder( Order.desc("avgWeight") )
    .list();
```

## 15.8. Detached queries and subqueries

The `DetachedCriteria` class lets you create a query outside the scope of a session, and then later execute it using some arbitrary `Session`.

```
DetachedCriteria query = DetachedCriteria.forClass(Cat.class)
    .add( Property.forName("sex").eq('F') );

Session session = ....;
Transaction txn = session.beginTransaction();
List results = query.getExecutableCriteria(session).setMaxResults(100).list();
txn.commit();
session.close();
```

A `DetachedCriteria` may also be used to express a subquery. Criterion instances involving subqueries may be obtained via `Subqueries` or `Property`.

```
DetachedCriteria avgWeight = DetachedCriteria.forClass(Cat.class)
    .setProjection( Property.forName("weight").avg() );
session.createCriteria(Cat.class)
    .add( Property.forName("weight").gt(avgWeight) )
    .list();
```

```
DetachedCriteria weights = DetachedCriteria.forClass(Cat.class)
    .setProjection( Property.forName("weight") );
session.createCriteria(Cat.class)
    .add( Subqueries.geAll("weight", weights) )
    .list();
```

Even correlated subqueries are possible:

```
DetachedCriteria avgWeightForSex = DetachedCriteria.forClass(Cat.class, "cat2")
    .setProjection( Property.forName("weight").avg() )
    .add( Property.forName("cat2.sex").eqProperty("cat.sex") );
session.createCriteria(Cat.class, "cat")
    .add( Property.forName("weight").gt(avgWeightForSex) )
    .list();
```

---

## Chapter 16. Native SQL

You may also express queries in the native SQL dialect of your database. This is useful if you want to utilize database specific features such as query hints or the `CONNECT` keyword in Oracle. It also provides a clean migration path from a direct SQL/JDBC based application to Hibernate.

Hibernate3 allows you to specify handwritten SQL (including stored procedures) for all create, update, delete, and load operations.

### 16.1. Creating a native SQL query

SQL queries are controlled via the `SQLQuery` interface, which is obtained by calling `Session.createSQLQuery()`.

```
List cats = sess.createSQLQuery("select {cat.*} from cats cat")
    .addEntity("cat", Cat.class);
    .setMaxResults(50);
    .list();
```

This query specified:

- the SQL query string, with a placeholder for Hibernate to inject the column aliases
- the entity returned by the query, and its SQL table alias

The `addEntity()` method associates SQL table aliases with entity classes, and determines the shape of the query result set.

The `addJoin()` method may be used to load associations to other entities and collections. TODO: examples!

A native SQL query might return a simple scalar value or a combination of scalars and entities.

```
Double max = (Double) sess.createSQLQuery("select max(cat.weight) as maxWeight from cats cat")
    .addScalar("maxWeight", Hibernate.DOUBLE);
    .uniqueResult();
```

### 16.2. Alias and property references

The `{cat.*}` notation used above is a shorthand for "all properties". Alternatively, you may list the columns explicitly, but even then you must let Hibernate inject the SQL column aliases for each property. The placeholder for a column alias is just the property name qualified by the table alias. In the following example, we retrieve Cats from a different table (`cat_log`) to the one declared in the mapping metadata. Notice that we may even use the property aliases in the where clause if we like. The `{}`-syntax is not required for named queries. See more in Section 16.3, "Named SQL queries"

```
String sql = "select cat.originalId as {cat.id}, " +
    "cat.mateid as {cat.mate}, cat.sex as {cat.sex}, " +
    "cat.weight*10 as {cat.weight}, cat.name as {cat.name} " +
    "from cat_log cat where {cat.mate} = :catId"

List loggedCats = sess.createSQLQuery(sql)
    .addEntity("cat", Cat.class)
    .setLong("catId", catId)
    .list();
```

*Note:* if you list each property explicitly, you must include all properties of the class *and its subclasses*!

## 16.3. Named SQL queries

Named SQL queries may be defined in the mapping document and called in exactly the same way as a named HQL query. In this case, we do *not* need to call `addEntity()`.

```
<sql-query name="mySqlQuery">
  <return alias="person" class="eg.Person"/>
  SELECT person.NAME AS {person.name},
         person.AGE AS {person.age},
         person.SEX AS {person.sex}
  FROM PERSON person WHERE person.NAME LIKE 'Hiber%'
</sql-query>
```

```
List people = sess.getNamedQuery("mySqlQuery")
    .setMaxResults(50)
    .list();
```

A named SQL query may return a scalar value. You must specify the column alias and Hibernate type using the `<return-scalar>` element:

```
<sql-query name="mySqlQuery">
  <return-scalar column="name" type="string"/>
  <return-scalar column="age" type="long"/>
  SELECT p.NAME AS name,
         p.AGE AS age,
  FROM PERSON p WHERE p.NAME LIKE 'Hiber%'
</sql-query>
```

The `<return-join>` and `<load-collection>` elements are used to join associations and define queries which initialize collections, respectively. **TODO!**

### 16.3.1. Using return-property to explicitly specify column/alias names

With `<return-property>` you can explicitly tell Hibernate what columns to use as opposed to use `{}`-syntax to let Hibernate inject its own aliases.

```
<sql-query name="mySqlQuery">
  <return alias="person" class="eg.Person">
    <return-property name="name" column="myName"/>
    <return-property name="age" column="myAge"/>
    <return-property name="sex" column="mySex"/>
  </return>
  SELECT person.NAME AS myName,
         person.AGE AS myAge,
         person.SEX AS mySex,
  FROM PERSON person WHERE person.NAME LIKE :name
</sql-query>
```

`<return-property>` also works with multiple columns. This solves a limitation with the `{}`-syntax which can not allow fine grained control of multi-column properties.

```
<sql-query name="organizationCurrentEmployments">
  <return alias="emp" class="Employment">
    <return-property name="salary">
      <return-column name="VALUE"/>
      <return-column name="CURRENCY"/>
    </return-property>
    <return-property name="endDate" column="myEndDate"/>
  </return>
</sql-query>
```

```

    </return>
    SELECT EMPLOYEE AS {emp.employee}, EMPLOYER AS {emp.employer},
    STARTDATE AS {emp.startDate}, ENDDATE AS {emp.endDate},
    REGIONCODE as {emp.regionCode}, EID AS {emp.id}, VALUE, CURRENCY
    FROM EMPLOYMENT
    WHERE EMPLOYER = :id AND ENDDATE IS NULL
    ORDER BY STARTDATE ASC
</sql-query>

```

Notice that in this example we used `<return-property>` in combination with the `{}`-syntax for injection. Allowing users to choose how they want to refer column and properties.

If your mapping has a discriminator you must use `<return-discriminator>` to specify the discriminator column.

### 16.3.2. Using stored procedures for querying

Hibernate 3 introduces support for queries via stored procedures. The stored procedures must return a resultset as the first out-parameter to be able to work with Hibernate. An example of such a stored procedure in Oracle 9 and higher is as follows:

```

CREATE OR REPLACE FUNCTION selectAllEmployments
    RETURN SYS_REFCURSOR
AS
    st_cursor SYS_REFCURSOR;
BEGIN
    OPEN st_cursor FOR
    SELECT EMPLOYEE, EMPLOYER,
    STARTDATE, ENDDATE,
    REGIONCODE, EID, VALUE, CURRENCY
    FROM EMPLOYMENT;
    RETURN st_cursor;
END;

```

To use this query in Hibernate you need to map it via a named query.

```

<sql-query name="selectAllEmployees_SP" callable="true">
    <return alias="emp" class="Employment">
        <return-property name="employee" column="EMPLOYEE"/>
        <return-property name="employer" column="EMPLOYER"/>
        <return-property name="startDate" column="STARTDATE"/>
        <return-property name="endDate" column="ENDDATE"/>
        <return-property name="regionCode" column="REGIONCODE"/>
        <return-property name="id" column="EID"/>
        <return-property name="salary">
            <return-column name="VALUE"/>
            <return-column name="CURRENCY"/>
        </return-property>
    </return>
    { ? = call selectAllEmployments() }
</sql-query>

```

Notice stored procedures currently only return scalars and entities. `<return-join>` and `<load-collection>` are not supported.

### Rules/limitations for using stored procedures

To use stored procedures with Hibernate the procedures have to follow some rules. If they do not follow those rules they are not usable with Hibernate. If you still want to use these procedures you have to execute them via `session.connection()`. The rules are different for each database, since database vendors have different stored procedure semantics/syntax.

For Oracle the following rules apply:

- The procedure must return a result set. This is done by returning a `SYS_REFCURSOR` in Oracle 9 or 10. In Oracle you need to define a `REF CURSOR` type.
- Recommended form is `{ ? = call procName(<parameters>) } or { ? = call procName }` (This is more an Oracle rule than a Hibernate rule.)

For Sybase or MS SQL server the following rules apply:

- The procedure must return a result set. Note that since these servers can/will return multiple result sets and update counts, Hibernate will iterate the results and take the first result that is a result set as its return value. Everything else will be discarded.
- If you can enable `SET NOCOUNT ON` in your procedure it will probably be more efficient, but this is not a requirement.

## 16.4. Custom SQL for create, update and delete

Hibernate3 can use custom SQL statements for create, update, and delete operations. The class and collection persisters in Hibernate already contain a set of configuration time generated strings (`insertsql`, `deletesql`, `updatesql` etc.). The mapping tags `<sql-insert>`, `<sql-delete>`, and `<sql-update>` override these strings:

```
<class name="Person">
  <id name="id">
    <generator class="increment"/>
  </id>
  <property name="name" not-null="true"/>
  <sql-insert>INSERT INTO PERSON (NAME, ID) VALUES ( UPPER(?), ? )</sql-insert>
  <sql-update>UPDATE PERSON SET NAME=UPPER(?) WHERE ID=?</sql-update>
  <sql-delete>DELETE FROM PERSON WHERE ID=?</sql-delete>
</class>
```

The SQL is directly executed in your database, so you are free to use any dialect you like. This will of course reduce the portability of your mapping if you use database specific SQL.

Stored procedures are supported if the `callable` attribute is set:

```
<class name="Person">
  <id name="id">
    <generator class="increment"/>
  </id>
  <property name="name" not-null="true"/>
  <sql-insert callable="true">{call createPerson (?, ?)}</sql-insert>
  <sql-delete callable="true">{? = call deletePerson (?)}</sql-delete>
  <sql-update callable="true">{? = call updatePerson (?, ?)}</sql-update>
</class>
```

The order of the positional parameters are currently vital, as they must be in the same sequence as Hibernate expects them.

You can see the expected order by enabling debug logging for the `org.hiberante.persister.entity` level. With this level enabled Hibernate will print out the static SQL that is used to create, update, delete etc. entities. To see the expected sequence, remember to not include your custom SQL in the mapping files as that will override the Hibernate generated static sql.

The stored procedures are in most cases (read: better do it than not) required to return the number of rows inserted/updated/deleted, as Hibernate has some runtime checks for the success of the statement. Hibernate always registers the first statement parameter as a numeric output parameter for the CUD operations:

```
CREATE OR REPLACE FUNCTION updatePerson (uid IN NUMBER, uname IN VARCHAR2)
RETURN NUMBER IS
BEGIN

    update PERSON
    set
        NAME = uname,
    where
        ID = uid;

    return SQL%ROWCOUNT;

END updatePerson;
```

## 16.5. Custom SQL for loading

You may also declare your own SQL (or HQL) queries for entity loading:

```
<sql-query name="person">
    <return alias="p" class="Person" lock-mode="upgrade"/>
    SELECT NAME AS {p.name}, ID AS {p.id} FROM PERSON WHERE ID=? FOR UPDATE
</sql-query>
```

This is just a named query declaration, as discussed earlier. You may reference this named query in a class mapping:

```
<class name="Person">
    <id name="id">
        <generator class="increment"/>
    </id>
    <property name="name" not-null="true"/>
    <loader query-ref="person"/>
</class>
```

And this also works with stored procedures.

TODO: Document the following example for collection loader.

```
<sql-query name="organizationEmployments">
    <load-collection alias="empcol" role="Organization.employments"/>
    SELECT {empcol.*}
    FROM EMPLOYMENT empcol
    WHERE EMPLOYER = :id
    ORDER BY STARTDATE ASC, EMPLOYEE ASC
</sql-query>

<sql-query name="organizationCurrentEmployments">
    <return alias="emp" class="Employment"/>
    <synchronize table="EMPLOYMENT"/>
    SELECT EMPLOYEE AS {emp.employee}, EMPLOYER AS {emp.employer},
        STARTDATE AS {emp.startDate}, ENDDATE AS {emp.endDate},
        REGIONCODE as {emp.regionCode}, ID AS {emp.id}
    FROM EMPLOYMENT
    WHERE EMPLOYER = :id AND ENDDATE IS NULL
    ORDER BY STARTDATE ASC
</sql-query>
```

---

## Chapter 17. Filtering data

Hibernate3 provides an innovative new approach to handling data with "visibility" rules. A *Hibernate filter* is a global, named, parameterized filter that may be enabled or disabled for a particular Hibernate session.

### 17.1. Hibernate filters

Hibernate3 adds the ability to pre-define filter criteria and attach those filters at both a class and a collection level. A filter criteria is the ability to define a restriction clause very similar to the existing "where" attribute available on the class and various collection elements. Except these filter conditions can be parameterized. The application can then make the decision at runtime whether given filters should be enabled and what their parameter values should be. Filters can be used like database views, but parameterized inside the application.

In order to use filters, they must first be defined and then attached to the appropriate mapping elements. To define a filter, use the `<filter-def/>` element within a `<hibernate-mapping/>` element:

```
<filter-def name="myFilter">
  <filter-param name="myFilterParam" type="string"/>
</filter-def>
```

Then, this filter can be attached to a class:

```
<class name="myClass" ...>
  ...
  <filter name="myFilter" condition=":myFilterParam = MY_FILTERED_COLUMN"/>
</class>
```

or, to a collection:

```
<set ...>
  <filter name="myFilter" condition=":myFilterParam = MY_FILTERED_COLUMN"/>
</set>
```

or, even to both (or multiples of each) at the same time.

The methods on `Session` are: `enableFilter(String filterName)`, `getEnabledFilter(String filterName)`, and `disableFilter(String filterName)`. By default, filters are *not* enabled for a given session; they must be explicitly enabled through use of the `Session.enableFilter()` method, which returns an instance of the `Filter` interface. Using the simple filter defined above, this would look like:

```
session.enableFilter("myFilter").setParameter("myFilterParam", "some-value");
```

Note that methods on the `org.hibernate.Filter` interface do allow the method-chaining common to much of Hibernate.

A full example, using temporal data with an effective record date pattern:

```
<filter-def name="effectiveDate">
  <filter-param name="asOfDate" type="date"/>
</filter-def>

<class name="Employee" ...>
  ...
  <many-to-one name="department" column="dept_id" class="Department"/>
  <property name="effectiveStartDate" type="date" column="eff_start_dt"/>
  <property name="effectiveEndDate" type="date" column="eff_end_dt"/>
```

```

...
    <!--
        Note that this assumes non-terminal records have an eff_end_dt set to
        a max db date for simplicity-sake
    -->
    <filter name="effectiveDate"
        condition=":asOfDate BETWEEN eff_start_dt and eff_end_dt"/>
</class>

<class name="Department" ...>
...
    <set name="employees" lazy="true">
        <key column="dept_id"/>
        <one-to-many class="Employee"/>
        <filter name="effectiveDate"
            condition=":asOfDate BETWEEN eff_start_dt and eff_end_dt"/>
    </set>
</class>

```

Then, in order to ensure that you always get back currently effective records, simply enable the filter on the session prior to retrieving employee data:

```

Session session = ...;
session.enabledFilter("effectiveDate").setParameter("asOfDate", new Date());
List results = session.createQuery("from Employee as e where e.salary > :targetSalary")
    .setLong("targetSalary", new Long(1000000))
    .list();

```

In the HQL above, even though we only explicitly mentioned a salary constraint on the results, because of the enabled filter the query will return only currently active employees who have a salary greater than a million dollars.

Note: if you plan on using filters with outer joining (either through HQL or load fetching) be careful of the direction of the condition expression. Its safest to set this up for left outer joining; in general, place the parameter first followed by the column name(s) after the operator.



---

# Chapter 18. XML Mapping

*Note that this is an experimental feature in Hibernate 3.0 and is under extremely active development.*

## 18.1. Working with XML data

Hibernate lets you work with persistent XML data in much the same way you work with persistent POJOs. A parsed XML tree can be thought of as just another way to represent the relational data at the object level, instead of POJOs.

Hibernate supports dom4j as API for manipulating XML trees. You can write queries that retrieve dom4j trees from the database and have any modification you make to the tree automatically synchronized to the database. You can even take an XML document, parse it using dom4j, and write it to the database with any of Hibernate's basic operations: `persist()`, `saveOrUpdate()`, `merge()`, `delete()`, `replicate()` (merging is not yet supported).

This feature has many applications including data import/export, externalization of entity data via JMS or SOAP and XSLT-based reporting.

A single mapping may be used to simultaneously map properties of a class and nodes of an XML document to the database, or, if there is no class to map, it may be used to map just the XML.

### 18.1.1. Specifying XML and class mapping together

Here is an example of mapping a POJO and XML simultaneously:

```
<class name="Account"
      table="ACCOUNTS"
      node="account">

  <id name="accountId"
      column="ACCOUNT_ID"
      node="@id"/>

  <many-to-one name="customer"
      column="CUSTOMER_ID"
      node="customer/@id"
      embed-xml="false"/>

  <property name="balance"
      column="BALANCE"
      node="balance"/>

  ...

</class>
```

### 18.1.2. Specifying only an XML mapping

Here is an example where there is no POJO class:

```
<class entity-name="Account"
      table="ACCOUNTS"
      node="account">

  <id name="id"
      column="ACCOUNT_ID"
```

```

        node="@id"
        type="string" />

    <many-to-one name="customerId"
        column="CUSTOMER_ID"
        node="customer/@id"
        embed-xml="false"
        entity-name="Customer" />

    <property name="balance"
        column="BALANCE"
        node="balance"
        type="big_decimal" />

    ...

</class>

```

This mapping allows you to access the data as a dom4j tree, or as a graph of property name/value pairs (java Maps). The property names are purely logical constructs that may be referred to in HQL queries.

## 18.2. XML mapping metadata

Many Hibernate mapping elements accept the `node` attribute. This lets you specify the name of an XML attribute or element that holds the property or entity data. The format of the `node` attribute must be one of the following:

- `"element-name"` - map to the named XML element
- `"@attribute-name"` - map to the named XML attribute
- `"."` - map to the parent element
- `"element-name/@attribute-name"` - map to the named attribute of the named element

For collections and single valued associations, there is an additional `embed-xml` attribute. If `embed-xml="true"`, the default, the XML tree for the associated entity (or collection of value type) will be embedded directly in the XML tree for the entity that owns the association. Otherwise, if `embed-xml="false"`, then only the referenced identifier value will appear in the XML for single point associations and collections will simply not appear at all.

You should be careful not to leave `embed-xml="true"` for too many associations, since XML does not deal well with circularity!

```

<class name="Customer"
    table="CUSTOMER"
    node="customer">

    <id name="id"
        column="CUST_ID"
        node="@id" />

    <map name="accounts"
        node="."
        embed-xml="true">
        <key column="CUSTOMER_ID"
            not-null="true" />
        <map-key column="SHORT_DESC"
            node="@short-desc"
            type="string" />
        <one-to-many entity-name="Account"
            embed-xml="false"
            node="account/@id" />
    </map>

```

```

    <component name="name"
      node="name">
      <property name="firstName"
        node="first-name"/>
      <property name="initial"
        node="initial"/>
      <property name="lastName"
        node="last-name"/>
    </component>

    ...

</class>

```

in this case, we have decided to embed the collection of account ids, but not the actual account data. The following HQL query:

```
from Customer c left join fetch c.accounts where c.lastName like :lastName
```

Would return datasets such as this:

```

<customer id="123456789">
  <account id="987632567" short-desc="Savings"/>
  <account id="985612323" short-desc="Credit Card"/>
  <name>
    <first-name>Gavin</first-name>
    <initial>A</initial>
    <last-name>King</last-name>
  </name>
  ...
</customer>

```

If you set `embed-xml="true"` on the `<one-to-many>` mapping, the data might look more like this:

```

<customer id="123456789">
  <account id="987632567" short-desc="Savings">
    <customer id="123456789"/>
    <balance>100.29</balance>
  </account>
  <account id="985612323" short-desc="Credit Card">
    <customer id="123456789"/>
    <balance>-2370.34</balance>
  </account>
  <name>
    <first-name>Gavin</first-name>
    <initial>A</initial>
    <last-name>King</last-name>
  </name>
  ...
</customer>

```

## 18.3. Manipulating XML data

Let's reread and update XML documents in the application. We do this by obtaining a dom4j session:

```

Document doc = ....;

Session session = factory.openSession();
Session dom4jSession = session.getSession(EntityMode.DOM4J);
Transaction tx = session.beginTransaction();

List results = dom4jSession
    .createQuery("from Customer c left join fetch c.accounts where c.lastName like :lastName")

```

```
.list();
for ( int i=0; i<results.size(); i++ ) {
    //add the customer data to the XML document
    Element customer = (Element) results.get(i);
    doc.add(customer);
}

tx.commit();
session.close();
```

```
Session session = factory.openSession();
Session dom4jSession = session.getSession(EntityMode.DOM4J);
Transaction tx = session.beginTransaction();

Element cust = (Element) dom4jSession.get("Customer", customerId);
for ( int i=0; i<results.size(); i++ ) {
    Element customer = (Element) results.get(i);
    //change the customer name in the XML and database
    Element name = customer.element("name");
    name.element("first-name").setText(firstName);
    name.element("initial").setText(initial);
    name.element("last-name").setText(lastName);
}

tx.commit();
session.close();
```

It is extremely useful to combine this feature with Hibernate's `replicate()` operation to implement XML-based data import/export.

---

# Chapter 19. Improving performance

## 19.1. Fetching strategies

A *fetching strategy* is the strategy Hibernate will use for retrieving associated objects if the application needs to navigate the association. Fetch strategies may be declared in the O/R mapping metadata, or over-ridden by a particular HQL or Criteria query.

Hibernate3 defines the following fetching strategies:

- *Join fetching* - Hibernate retrieves the associated instance or collection in the same `SELECT`, using an `OUTER JOIN`.
- *Select fetching* - a second `SELECT` is used to retrieve the associated entity or collection. Unless you explicitly disable lazy fetching by specifying `lazy="false"`, this second select will only be executed when you actually access the association.
- *Subselect fetching* - a second `SELECT` is used to retrieve the associated collections for all entities retrieved in a previous query or fetch. Unless you explicitly disable lazy fetching by specifying `lazy="false"`, this second select will only be executed when you actually access the association.
- *Batch fetching* - an optimization strategy for select fetching - Hibernate retrieves a batch of entity instances or collections in a single `SELECT`, by specifying a list of primary keys or foreign keys.

By default, Hibernate3 uses lazy select fetching, which is the best choice for most entities and collections in most applications. If you set `hibernate.default_batch_fetch_size`, Hibernate will use the batch fetch optimization to lazy fetching (this optimization may also be enabled at a more granular level).

However, there is one problem to be aware of. Access to a lazy association outside of the context of an open Hibernate session will result in an exception. For example:

```
s = sessions.openSession();
Transaction tx = s.beginTransaction();

User u = (User) s.createQuery("from User u where u.name=:userName")
    .setString("userName", userName).uniqueResult();
Map permissions = u.getPermissions();

tx.commit();
s.close();

Integer accessLevel = (Integer) permissions.get("accounts"); // Error!
```

Since the permissions collection was not initialized when the `Session` was closed, the collection will not be able to load its state. *Hibernate does not support lazy initialization for detached objects*. The fix is to move the code that reads from the collection to just before the commit. (There are other more advanced ways to solve this problem, some are discussed later.)

It is also possible to use a non-lazy collection or join fetching, which is non-lazy by nature. However, it is intended that lazy initialization be used for almost all collections, especially for collections of entity references. If you define too many non-lazy associations in your object model, Hibernate will end up needing to fetch the entire database into memory in every transaction!

On the other hand, we often want to choose join fetching (which is non-lazy by nature) instead of select fetch-

ing in a particular transaction. We'll now see how to customize the fetching strategy. In Hibernate3, the mechanisms for choosing a fetch strategy are identical for single-valued associations and collections.

### 19.1.1. Tuning fetch strategies

Select fetching (the default) is extremely vulnerable to N+1 selects problems, so we might want to enable join fetching in the mapping document:

```
<set name="permissions"
    fetch="join">
  <key column="userId"/>
  <one-to-many class="Permission"/>
</set>
```

```
<many-to-one name="mother" class="Cat" fetch="join"/>
```

The fetch strategy defined in the mapping document affects:

- retrieval via `get()` or `load()`
- retrieval that happens implicitly when an association is navigated (lazy fetching)
- Criteria queries

*Specifying join as the fetch strategy in the mapping document does not affect HQL queries.*

Usually, we don't use the mapping document to customize fetching. Instead, we keep the default behavior, and override it for a particular transaction, using the HQL `LEFT JOIN FETCH` clause. This tells Hibernate to fetch the association eagerly in the first select, using an outer join. In the Criteria query API, you would use `setFetchMode(FetchMode.JOIN)`.

If you ever feel like you wish you could change the fetching strategy used by `get()` or `load()`, simply use a Criteria query, for example:

```
User user = (User) session.createCriteria(User.class)
    .setFetchMode("permissions", FetchMode.JOIN)
    .add( Restrictions.idEq(userId) )
    .uniqueResult();
```

(This is Hibernate's equivalent of what some ORM solutions call a "fetch plan".)

Join fetching for collections has one limitation: you may only set one collection role per persistent class or query to be fetched per outer join. Hibernate forbids Cartesian products when possible, `SELECTING` two collections per outer join would create one. This would almost always be slower than two (lazy or non-deferred) `SELECTS`. The restriction to a single outer-joined collection applies to both the mapping fetching strategies and to HQL/Criteria queries.

If you run into this limitation, you should use subselect or batch fetching to achieve acceptable performance. This is common when retrieving a tree of collection-valued associations.

There are no restrictions to join fetching of single-ended associations.

### 19.1.2. Single-ended association proxies

Lazy fetching for collections is implemented using Hibernate's own implementation of persistent collections. However, a different mechanism is needed for lazy behavior in single-ended associations. The target entity of the association must be proxied. Hibernate implements lazy initializing proxies for persistent objects using runtime bytecode enhancement (via the excellent CGLIB library).

By default, Hibernate3 generates proxies (at startup) for all persistent classes and uses them to enable lazy fetching of many-to-one and one-to-one associations.

The mapping file may declare an interface to use as the proxy interface for that class, with the `proxy` attribute. By default, Hibernate uses a subclass of the class. *Note that the proxied class must implement a default constructor with at least package visibility. We recommend this constructor for all persistent classes!*

There are some gotchas to be aware of when extending this approach to polymorphic classes, eg.

```
<class name="Cat" proxy="Cat">
    .....
    <subclass name="DomesticCat">
        .....
    </subclass>
</class>
```

Firstly, instances of `Cat` will never be castable to `DomesticCat`, even if the underlying instance is an instance of `DomesticCat`:

```
Cat cat = (Cat) session.load(Cat.class, id); // instantiate a proxy (does not hit the db)
if ( cat.isDomesticCat() ) {                // hit the db to initialize the proxy
    DomesticCat dc = (DomesticCat) cat;      // Error!
    ....
}
```

Secondly, it is possible to break `proxy ==`.

```
Cat cat = (Cat) session.load(Cat.class, id); // instantiate a Cat proxy
DomesticCat dc =
    (DomesticCat) session.load(DomesticCat.class, id); // acquire new DomesticCat proxy!
System.out.println(cat==dc); // false
```

However, the situation is not quite as bad as it looks. Even though we now have two references to different proxy objects, the underlying instance will still be the same object:

```
cat.setWeight(11.0); // hit the db to initialize the proxy
System.out.println( dc.getWeight() ); // 11.0
```

Third, you may not use a CGLIB proxy for a `final` class or a class with any `final` methods.

Finally, if your persistent object acquires any resources upon instantiation (eg. in initializers or default constructor), then those resources will also be acquired by the proxy. The proxy class is an actual subclass of the persistent class.

These problems are all due to fundamental limitations in Java's single inheritance model. If you wish to avoid these problems your persistent classes must each implement an interface that declares its business methods. You should specify these interfaces in the mapping file. eg.

```
<class name="CatImpl" proxy="Cat">
    .....
    <subclass name="DomesticCatImpl" proxy="DomesticCat">
        .....
    </subclass>
</class>
```

where `CatImpl` implements the interface `Cat` and `DomesticCatImpl` implements the interface `DomesticCat`. Then proxies for instances of `Cat` and `DomesticCat` may be returned by `load()` or `iterate()`. (Note that `list()` does not usually return proxies.)

```
Cat cat = (Cat) session.load(CatImpl.class, catid);
Iterator iter = session.iterate("from CatImpl as cat where cat.name='fritz'");
Cat fritz = (Cat) iter.next();
```

Relationships are also lazily initialized. This means you must declare any properties to be of type `Cat`, not `CatImpl`.

Certain operations do *not* require proxy initialization

- `equals()`, if the persistent class does not override `equals()`
- `hashCode()`, if the persistent class does not override `hashCode()`
- The identifier getter method

Hibernate will detect persistent classes that override `equals()` or `hashCode()`.

### 19.1.3. Initializing collections and proxies

A `LazyInitializationException` will be thrown by Hibernate if an uninitialized collection or proxy is accessed outside of the scope of the `Session`, ie. when the entity owning the collection or having the reference to the proxy is in the detached state.

Sometimes we need to ensure that a proxy or collection is initialized before closing the `Session`. Of course, we can always force initialization by calling `cat.getSex()` or `cat.getKittens().size()`, for example. But that is confusing to readers of the code and is not convenient for generic code.

The static methods `Hibernate.initialize()` and `Hibernate.isInitialized()` provide the application with a convenient way of working with lazily initialized collections or proxies. `Hibernate.initialize(cat)` will force the initialization of a proxy, `cat`, as long as its `Session` is still open. `Hibernate.initialize(cat.getKittens())` has a similar effect for the collection of kittens.

Another option is to keep the `Session` open until all needed collections and proxies have been loaded. In some application architectures, particularly where the code that accesses data using Hibernate, and the code that uses it are in different application layers or different physical processes, it can be a problem to ensure that the `Session` is open when a collection is initialized. There are two basic ways to deal with this issue:

- In a web-based application, a servlet filter can be used to close the `Session` only at the very end of a user request, once the rendering of the view is complete (the *Open Session in View* pattern). Of course, this places heavy demands on the correctness of the exception handling of your application infrastructure. It is vitally important that the `Session` is closed and the transaction ended before returning to the user, even when an exception occurs during rendering of the view. The servlet filter has to be able to access the `Session` for this approach. We recommend that a `ThreadLocal` variable be used to hold the current `Session` (see chapter 1, Section 1.4, “Playing with cats”, for an example implementation).
- In an application with a separate business tier, the business logic must “prepare” all collections that will be needed by the web tier before returning. This means that the business tier should load all the data and return all the data already initialized to the presentation/web tier that is required for a particular use case. Usually, the application calls `Hibernate.initialize()` for each collection that will be needed in the web tier (this call must occur before the session is closed) or retrieves the collection eagerly using a Hibernate query with a `FETCH` clause or a `FetchMode.JOIN` in `Criteria`. This is usually easier if you adopt the *Command* pattern



instead of a *Session Facade*.

- You may also attach a previously loaded object to a new `Session` with `merge()` or `lock()` before accessing uninitialized collections (or other proxies). No, Hibernate does not, and certainly *should* not do this automatically, since it would introduce ad hoc transaction semantics!

Sometimes you don't want to initialize a large collection, but still need some information about it (like its size) or a subset of the data.

You can use a collection filter to get the size of a collection without initializing it:

```
( (Integer) s.createFilter( collection, "select count(*)" ).list().get(0) ).intValue()
```

The `createFilter()` method is also used to efficiently retrieve subsets of a collection without needing to initialize the whole collection:

```
s.createFilter( lazyCollection, "" ).setFirstResult(0).setMaxResults(10).list();
```

### 19.1.4. Using batch fetching

Hibernate can make efficient use of batch fetching, that is, Hibernate can load several uninitialized proxies if one proxy is accessed (or collections. Batch fetching is an optimization of the lazy select fetching strategy. There are two ways you can tune batch fetching: on the class and the collection level.

Batch fetching for classes/entities is easier to understand. Imagine you have the following situation at runtime: You have 25 `Cat` instances loaded in a `Session`, each `Cat` has a reference to its owner, a `Person`. The `Person` class is mapped with a proxy, `lazy="true"`. If you now iterate through all cats and call `getOwner()` on each, Hibernate will by default execute 25 `SELECT` statements, to retrieve the proxied owners. You can tune this behavior by specifying a `batch-size` in the mapping of `Person`:

```
<class name="Person" batch-size="10">...</class>
```

Hibernate will now execute only three queries, the pattern is 10, 10, 5.

You may also enable batch fetching of collections. For example, if each `Person` has a lazy collection of `Cats`, and 10 persons are currently loaded in the `Session`, iterating through all persons will generate 10 `SELECT`s, one for every call to `getCats()`. If you enable batch fetching for the `cats` collection in the mapping of `Person`, Hibernate can pre-fetch collections:

```
<class name="Person">
  <set name="cats" batch-size="3">
    ...
  </set>
</class>
```

With a `batch-size` of 8, Hibernate will load 3, 3, 3, 1 collections in four `SELECT`s. Again, the value of the attribute depends on the expected number of uninitialized collections in a particular `Session`.

Batch fetching of collections is particularly useful if you have a nested tree of items, ie. the typical bill-of-materials pattern. (Although a *nested set* or a *materialized path* might be a better option for read-mostly trees.)

### 19.1.5. Using lazy property fetching

Hibernate3 supports the lazy fetching of individual properties. This optimization technique is also known as *fetch groups*. Please note that this is mostly a marketing feature, as in practice, optimizing row reads is much more important than optimization of column reads. However, only loading some properties of a class might be useful in extreme cases, when legacy tables have hundreds of columns and the data model can not be improved.

To enable lazy property loading, set the `lazy` attribute on your particular property mappings:

```
<class name="Document">
  <id name="id">
    <generator class="native"/>
  </id>
  <property name="name" not-null="true" length="50"/>
  <property name="summary" not-null="true" length="200" lazy="true"/>
  <property name="text" not-null="true" length="2000" lazy="true"/>
</class>
```

Lazy property loading requires buildtime bytecode instrumentation! If your persistent classes are not enhanced, Hibernate will silently ignore lazy property settings and fall back to immediate fetching.

For bytecode instrumentation, use the following Ant task:

```
<target name="instrument" depends="compile">
  <taskdef name="instrument" classname="org.hibernate.tool.instrument.InstrumentTask">
    <classpath path="{jar.path}"/>
    <classpath path="{classes.dir}"/>
    <classpath refid="lib.class.path"/>
  </taskdef>

  <instrument verbose="true">
    <fileset dir="{testclasses.dir}/org/hibernate/auction/model">
      <include name="*.class"/>
    </fileset>
  </instrument>
</target>
```

A different (better?) way to avoid unnecessary column reads, at least for read-only transactions is to use the projection features of HQL or Criteria queries. This avoids the need for buildtime bytecode processing and is certainly a preferred solution.

A completely different way to avoid problems with N+1 selects is to use the second-level cache.

## 19.2. The Second Level Cache

A Hibernate `Session` is a transaction-level cache of persistent data. It is possible to configure a cluster or JVM-level (`SessionFactory`-level) cache on a class-by-class and collection-by-collection basis. You may even plug in a clustered cache. Be careful. Caches are never aware of changes made to the persistent store by another application (though they may be configured to regularly expire cached data).

By default, Hibernate uses `EHCACHE` for JVM-level caching. (JCS support is now deprecated and will be removed in a future version of Hibernate.) You may choose a different implementation by specifying the name of a class that implements `org.hibernate.cache.CacheProvider` using the property `hibernate.cache.provider_class`.

**Table 19.1. Cache Providers**

Cache	Provider class	Type	Cluster Safe	Query Cache Supported
Hashtable (not intended for production use)	org.hibernate.cache.HashtableCacheProvider	memory		yes
EHCache	org.hibernate.cache.EhCacheProvider	memory, disk		yes
OSCache	org.hibernate.cache.OSCacheProvider	memory, disk		yes
SwarmCache	org.hibernate.cache.SwarmCacheProvider	clustered (ip multicast)	yes (clustered invalidation)	
JBoss TreeCache	org.hibernate.cache.TreeCacheProvider	clustered (ip multicast), transactional	yes (replication)	yes (clock sync req.)

### 19.2.1. Cache mappings

The `<cache>` element of a class or collection mapping has the following form:

```
<cache
  usage="transactional|read-write|nonstrict-read-write|read-only" (1)
/>
```

(1) `usage` specifies the caching strategy: `transactional`, `read-write`, `nonstrict-read-write` or `read-only`

Alternatively (preferably?), you may specify `<class-cache>` and `<collection-cache>` elements in `hibernate.cfg.xml`.

The `usage` attribute specifies a *cache concurrency strategy*.

### 19.2.2. Strategy: read only

If your application needs to read but never modify instances of a persistent class, a `read-only` cache may be used. This is the simplest and best performing strategy. It's even perfectly safe for use in a cluster.

```
<class name="eg.Immutable" mutable="false">
  <cache usage="read-only"/>
  ....
</class>
```

### 19.2.3. Strategy: read/write

If the application needs to update data, a `read-write` cache might be appropriate. This cache strategy should never be used if serializable transaction isolation level is required. If the cache is used in a JTA environment, you must specify the property `hibernate.transaction.manager_lookup_class`, naming a strategy for obtaining the JTA `TransactionManager`. In other environments, you should ensure that the transaction is completed when `Session.close()` or `Session.disconnect()` is called. If you wish to use this strategy in a cluster, you should ensure that the underlying cache implementation supports locking. The built-in cache providers do *not*.

```
<class name="eg.Cat" .... >
  <cache usage="read-write"/>
```

```

    ....
    <set name="kittens" ... >
        <cache usage="read-write"/>
        ....
    </set>
</class>

```

### 19.2.4. Strategy: nonstrict read/write

If the application only occasionally needs to update data (ie. if it is extremely unlikely that two transactions would try to update the same item simultaneously) and strict transaction isolation is not required, a `nonstrict-read-write` cache might be appropriate. If the cache is used in a JTA environment, you must specify `hibernate.transaction.manager_lookup_class`. In other environments, you should ensure that the transaction is completed when `Session.close()` or `Session.disconnect()` is called.

### 19.2.5. Strategy: transactional

The `transactional` cache strategy provides support for fully transactional cache providers such as JBoss TreeCache. Such a cache may only be used in a JTA environment and you must specify `hibernate.transaction.manager_lookup_class`.

None of the cache providers support all of the cache concurrency strategies. The following table shows which providers are compatible with which concurrency strategies.

**Table 19.2. Cache Concurrency Strategy Support**

Cache	read-only	nonstrict-read-write	read-write	transactional
Hashtable (not intended for production use)	yes	yes	yes	
EHCache	yes	yes	yes	
OSCache	yes	yes	yes	
SwarmCache	yes	yes		
JBoss TreeCache	yes			yes

## 19.3. Managing the caches

Whenever you pass an object to `save()`, `update()` or `saveOrUpdate()` and whenever you retrieve an object using `load()`, `get()`, `list()`, `iterate()` or `scroll()`, that object is added to the internal cache of the `Session`.

When `flush()` is subsequently called, the state of that object will be synchronized with the database. If you do not want this synchronization to occur or if you are processing a huge number of objects and need to manage memory efficiently, the `evict()` method may be used to remove the object and its collections from the first-level cache.

```

ScrollableResult cats = sess.createQuery("from Cat as cat").scroll(); //a huge result set
while ( cats.next() ) {

```

```
Cat cat = (Cat) cats.get(0);
doSomethingWithACat(cat);
sess.evict(cat);
}
```

The `Session` also provides a `contains()` method to determine if an instance belongs to the session cache.

To completely evict all objects from the session cache, call `Session.clear()`

For the second-level cache, there are methods defined on `SessionFactory` for evicting the cached state of an instance, entire class, collection instance or entire collection role.

```
sessionFactory.evict(Cat.class, catId); //evict a particular Cat
sessionFactory.evict(Cat.class); //evict all Cats
sessionFactory.evictCollection("Cat.kittens", catId); //evict a particular collection of kittens
sessionFactory.evictCollection("Cat.kittens"); //evict all kitten collections
```

The `CacheMode` controls how a particular session interacts with the second-level cache.

- `CacheMode.NORMAL` - read items from and write items to the second-level cache
- `CacheMode.GET` - read items from the second-level cache, but don't write to the second-level cache except when updating data
- `CacheMode.PUT` - write items to the second-level cache, but don't read from the second-level cache
- `CacheMode.REFRESH` - write items to the second-level cache, but don't read from the second-level cache, bypass the effect of `hibernate.cache.use_minimal_puts`, forcing a refresh of the second-level cache for all items read from the database

To browse the contents of a second-level or query cache region, use the `Statistics` API:

```
Map cacheEntries = sessionFactory.getStatistics()
    .getSecondLevelCacheStatistics(regionName)
    .getEntries();
```

You'll need to enable statistics, and, optionally, force Hibernate to keep the cache entries in a more human-understandable format:

```
hibernate.generate_statistics true
hibernate.cache.use_structured_entries true
```

## 19.4. The Query Cache

Query result sets may also be cached. This is only useful for queries that are run frequently with the same parameters. To use the query cache you must first enable it:

```
hibernate.cache.use_query_cache true
```

This setting causes the creation of two new cache regions - one holding cached query result sets (`org.hibernate.cache.StandardQueryCache`), the other holding timestamps of the most recent updates to queryable tables (`org.hibernate.cache.UpdateTimestampsCache`). Note that the query cache does not cache the state of the actual entities in the result set; it caches only identifier values and results of value type. So the query cache should always be used in conjunction with the second-level cache.

Most queries do not benefit from caching, so by default queries are not cached. To enable caching, call `Query.setCacheable(true)`. This call allows the query to look for existing cache results or add its results to the cache when it is executed.

If you require fine-grained control over query cache expiration policies, you may specify a named cache region for a particular query by calling `Query.setCacheRegion()`.

```
List blogs = sess.createQuery("from Blog blog where blog.blogger = :blogger")
    .setEntity("blogger", blogger)
    .setMaxResults(15)
    .setCacheable(true)
    .setCacheRegion("frontpages")
    .list();
```

If the query should force a refresh of its query cache region, you should call `Query.setCacheMode(CacheMode.REFRESH)`. This is particularly useful in cases where underlying data may have been updated via a separate process (i.e., not modified through Hibernate) and allows the application to selectively refresh particular query result sets. This is a more efficient alternative to eviction of a query cache region via `SessionFactory.evictQueries()`.

## 19.5. Understanding Collection performance

We've already spent quite some time talking about collections. In this section we will highlight a couple more issues about how collections behave at runtime.

### 19.5.1. Taxonomy

Hibernate defines three basic kinds of collections:

- collections of values
- one to many associations
- many to many associations

This classification distinguishes the various table and foreign key relationships but does not tell us quite everything we need to know about the relational model. To fully understand the relational structure and performance characteristics, we must also consider the structure of the primary key that is used by Hibernate to update or delete collection rows. This suggests the following classification:

- indexed collections
- sets
- bags

All indexed collections (maps, lists, arrays) have a primary key consisting of the `<key>` and `<index>` columns. In this case collection updates are usually extremely efficient - the primary key may be efficiently indexed and a particular row may be efficiently located when Hibernate tries to update or delete it.

Sets have a primary key consisting of `<key>` and element columns. This may be less efficient for some types of collection element, particularly composite elements or large text or binary fields; the database may not be able

to index a complex primary key as efficiently. On the other hand, for one to many or many to many associations, particularly in the case of synthetic identifiers, it is likely to be just as efficient. (Side-note: if you want `SchemaExport` to actually create the primary key of a `<set>` for you, you must declare all columns as `not-null="true"`.)

`<idbag>` mappings define a surrogate key, so they are always very efficient to update. In fact, they are the best case.

Bags are the worst case. Since a bag permits duplicate element values and has no index column, no primary key may be defined. Hibernate has no way of distinguishing between duplicate rows. Hibernate resolves this problem by completely removing (in a single `DELETE`) and recreating the collection whenever it changes. This might be very inefficient.

Note that for a one-to-many association, the "primary key" may not be the physical primary key of the database table - but even in this case, the above classification is still useful. (It still reflects how Hibernate "locates" individual rows of the collection.)

### 19.5.2. Lists, maps, idbags and sets are the most efficient collections to update

From the discussion above, it should be clear that indexed collections and (usually) sets allow the most efficient operation in terms of adding, removing and updating elements.

There is, arguably, one more advantage that indexed collections have over sets for many to many associations or collections of values. Because of the structure of a `Set`, Hibernate doesn't ever `UPDATE` a row when an element is "changed". Changes to a `Set` always work via `INSERT` and `DELETE` (of individual rows). Once again, this consideration does not apply to one to many associations.

After observing that arrays cannot be lazy, we would conclude that lists, maps and idbags are the most performant (non-inverse) collection types, with sets not far behind. Sets are expected to be the most common kind of collection in Hibernate applications. This is because the "set" semantics are most natural in the relational model.

However, in well-designed Hibernate domain models, we usually see that most collections are in fact one-to-many associations with `inverse="true"`. For these associations, the update is handled by the many-to-one end of the association, and so considerations of collection update performance simply do not apply.

### 19.5.3. Bags and lists are the most efficient inverse collections

Just before you ditch bags forever, there is a particular case in which bags (and also lists) are much more performant than sets. For a collection with `inverse="true"` (the standard bidirectional one-to-many relationship idiom, for example) we can add elements to a bag or list without needing to initialize (fetch) the bag elements! This is because `Collection.add()` or `Collection.addAll()` must always return `true` for a bag or `List` (unlike a `Set`). This can make the following common code much faster.

```
Parent p = (Parent) sess.load(Parent.class, id);
Child c = new Child();
c.setParent(p);
p.getChildren().add(c); //no need to fetch the collection!
sess.flush();
```

### 19.5.4. One shot delete

Occasionally, deleting collection elements one by one can be extremely inefficient. Hibernate isn't completely stupid, so it knows not to do that in the case of an newly-empty collection (if you called `list.clear()`, for example). In this case, Hibernate will issue a single `DELETE` and we are done!

Suppose we add a single element to a collection of size twenty and then remove two elements. Hibernate will issue one `INSERT` statement and two `DELETE` statements (unless the collection is a bag). This is certainly desirable.

However, suppose that we remove eighteen elements, leaving two and then add three new elements. There are two possible ways to proceed

- delete eighteen rows one by one and then insert three rows
- remove the whole collection (in one SQL `DELETE`) and insert all five current elements (one by one)

Hibernate isn't smart enough to know that the second option is probably quicker in this case. (And it would probably be undesirable for Hibernate to be that smart; such behaviour might confuse database triggers, etc.)

Fortunately, you can force this behaviour (ie. the second strategy) at any time by discarding (ie. dereferencing) the original collection and returning a newly instantiated collection with all the current elements. This can be very useful and powerful from time to time.

Of course, one-shot-delete does not apply to collections mapped `inverse="true"`.

## 19.6. Monitoring performance

Optimization is not much use without monitoring and access to performance numbers. Hibernate provides a full range of figures about its internal operations. Statistics in Hibernate are available per `SessionFactory`.

### 19.6.1. Monitoring a SessionFactory

You can access `SessionFactory` metrics in two ways. Your first option is to call `sessionFactory.getStatistics()` and read or display the `Statistics` yourself.

Hibernate can also use JMX to publish metrics if you enable the `StatisticsService` MBean. You may enable a single MBean for all your `SessionFactory` or one per factory. See the following code for minimalistic configuration examples:

```
// MBean service registration for a specific SessionFactory
Hashtable tb = new Hashtable();
tb.put("type", "statistics");
tb.put("sessionFactory", "myFinancialApp");
ObjectName on = new ObjectName("hibernate", tb); // MBean object name

StatisticsService stats = new StatisticsService(); // MBean implementation
stats.setSessionFactory(sessionFactory); // Bind the stats to a SessionFactory
server.registerMBean(stats, on); // Register the Mbean on the server
```

```
// MBean service registration for all SessionFactory's
Hashtable tb = new Hashtable();
tb.put("type", "statistics");
tb.put("sessionFactory", "all");
ObjectName on = new ObjectName("hibernate", tb); // MBean object name

StatisticsService stats = new StatisticsService(); // MBean implementation
server.registerMBean(stats, on); // Register the MBean on the server
```



TODO: This doesn't make sense: In the first case, we retrieve and use the MBean directly. In the second one, we must give the JNDI name in which the session factory is held before using it. Use `hibernateStatsBean.setSessionFactoryJNDIName("my/JNDI/Name")`

You can (de)activate the monitoring for a `SessionFactory`

- at configuration time, set `hibernate.generate_statistics` to false
- at runtime: `sf.getStatistics().setStatisticsEnabled(true)` or `hibernateStatsBean.setStatisticsEnabled(true)`

Statistics can be reset programatically using the `clear()` method. A summary can be sent to a logger (info level) using the `logSummary()` method.

## 19.6.2. Metrics

Hibernate provides a number of metrics, from very basic to the specialized information only relevant in certain scenarios. All available counters are described in the `Statistics` interface API, in three categories:

- Metrics related to the general `Session` usage, such as number of open sessions, retrieved JDBC connections, etc.
- Metrics related to the entities, collections, queries, and caches as a whole (aka global metrics),
- Detailed metrics related to a particular entity, collection, query or cache region.

For example you can check the cache hit, miss, and put ratio of entities, collections and queries, and the average time a query needs. Beware that the number of milliseconds is subject to approximation in Java. Hibernate is tied to the JVM precision, on some platforms this might even only be accurate to 10 seconds.

Simple getters are used to access the global metrics (i.e. not tied to a particular entity, collection, cache region, etc.). You can access the metrics of a particular entity, collection or cache region through its name, and through its HQL or SQL representation for queries. Please refer to the `Statistics`, `EntityStatistics`, `CollectionStatistics`, `SecondLevelCacheStatistics`, and `QueryStatistics` API Javadoc for more information. The following code shows a simple example:

```
Statistics stats = HibernateUtil.sessionFactory.getStatistics();

double queryCacheHitCount = stats.getQueryCacheHitCount();
double queryCacheMissCount = stats.getQueryCacheMissCount();
double queryCacheHitRatio =
    queryCacheHitCount / (queryCacheHitCount + queryCacheMissCount);

log.info("Query Hit ratio:" + queryCacheHitRatio);

EntityStatistics entityStats =
    stats.getEntityStatistics( Cat.class.getName() );
long changes =
    entityStats.getInsertCount()
    + entityStats.getUpdateCount()
    + entityStats.getDeleteCount();
log.info(Cat.class.getName() + " changed " + changes + "times" );
```

To work on all entities, collections, queries and region caches, you can retrieve the list of names of entities, collections, queries and region caches with the following methods: `getQueries()`, `getEntityNames()`, `getCol-`

`lectionRoleNames()`, and `getSecondLevelCacheRegionNames()`.

---

## Chapter 20. Toolset Guide

Roundtrip engineering with Hibernate is possible using a set of Eclipse plugins, commandline tools, as well as Ant tasks.

The *Hibernate Tools* currently include plugins for the Eclipse IDE as well as Ant tasks for reverse engineering of existing databases:

- *Mapping Editor*: An editor for Hibernate XML mapping files, supporting auto-completion and syntax highlighting. It also supports semantic auto-completion for class names and property/field names, making it much more versatile than a normal XML editor.
- *Console*: The console is a new view in Eclipse. In addition to a tree overview of your console configurations, you also get an interactive view of your persistent classes and their relationships. The console allows you to execute HQL queries against your database and browse the result directly in Eclipse.
- *Development Wizards*: Several wizards are provided with the Hibernate Eclipse tools; you can use a wizard to quickly generate Hibernate configuration (cfg.xml) files, or you may even completely reverse engineer an existing database schema into POJO source files and Hibernate mapping files. The reverse engineering wizard supports customizable templates.
- *Ant Tasks*:

Please refer to the *Hibernate Tools* package and its documentation for more information.

However, the Hibernate main package comes bundled with an integrated tool (it can even be used from "inside" Hibernate on-the-fly): *SchemaExport* aka `hbm2ddl`.

### 20.1. Automatic schema generation

DDL may be generated from your mapping files by a Hibernate utility. The generated schema includes referential integrity constraints (primary and foreign keys) for entity and collection tables. Tables and sequences are also created for mapped identifier generators.

You *must* specify a SQL `Dialect` via the `hibernate.dialect` property when using this tool, as DDL is highly vendor specific.

First, customize your mapping files to improve the generated schema.

#### 20.1.1. Customizing the schema

Many Hibernate mapping elements define an optional attribute named `length`. You may set the length of a column with this attribute. (Or, for numeric/decimal data types, the precision.)

Some tags also accept a `not-null` attribute (for generating a `NOT NULL` constraint on table columns) and a `unique` attribute (for generating `UNIQUE` constraint on table columns).

Some tags accept an `index` attribute for specifying the name of an index for that column. A `unique-key` attribute can be used to group columns in a single unit key constraint. Currently, the specified value of the `unique-key` attribute is *not* used to name the constraint, only to group the columns in the mapping file.

## Examples:

```
<property name="foo" type="string" length="64" not-null="true"/>

<many-to-one name="bar" foreign-key="fk_foo_bar" not-null="true"/>

<element column="serial_number" type="long" not-null="true" unique="true"/>
```

Alternatively, these elements also accept a child `<column>` element. This is particularly useful for multi-column types:

```
<property name="foo" type="string">
  <column name="foo" length="64" not-null="true" sql-type="text"/>
</property>

<property name="bar" type="my.customtypes.MultiColumnType"/>
  <column name="fee" not-null="true" index="bar_idx"/>
  <column name="fi" not-null="true" index="bar_idx"/>
  <column name="fo" not-null="true" index="bar_idx"/>
</property>
```

The `sql-type` attribute allows the user to override the default mapping of Hibernate type to SQL datatype.

The `check` attribute allows you to specify a check constraint.

```
<property name="foo" type="integer">
  <column name="foo" check="foo > 10"/>
</property>

<class name="Foo" table="foos" check="bar < 100.0">
  ...
  <property name="bar" type="float"/>
</class>
```

**Table 20.1. Summary**

Attribute	Values	Interpretation
length	number	column length/decimal precision
not-null	true false	specifies that the column should be non-nullable
unique	true false	specifies that the column should have a unique constraint
index	index_name	specifies the name of a (multi-column) index
unique-key	unique_key_name	specifies the name of a multi-column unique constraint
foreign-key	foreign_key_name	specifies the name of the foreign key constraint generated for an association, use it on <code>&lt;one-to-one&gt;</code> , <code>&lt;many-to-one&gt;</code> , <code>&lt;key&gt;</code> , and <code>&lt;many-to-many&gt;</code> mapping elements. Note that <code>inverse="true"</code> sides will not be considered by <code>SchemaExport</code> .
sql-type	column_type	overrides the default column type (attribute of <code>&lt;column&gt;</code> element only)
check	SQL expression	create an SQL check constraint on either column or table

## 20.1.2. Running the tool

The `SchemaExport` tool writes a DDL script to standard out and/or executes the DDL statements.

```
java -cp hibernate_classpaths org.hibernate.tool.hbm2ddl.SchemaExport options mapping_files
```

**Table 20.2. SchemaExport Command Line Options**

Option	Description
<code>--quiet</code>	don't output the script to stdout
<code>--drop</code>	only drop the tables
<code>--text</code>	don't export to the database
<code>--output=my_schema.ddl</code>	output the ddl script to a file
<code>--config=hibernate.cfg.xml</code>	read Hibernate configuration from an XML file
<code>--properties=hibernate.properties</code>	read database properties from a file
<code>--format</code>	format the generated SQL nicely in the script
<code>--delimiter=x</code>	set an end of line delimiter for the script

You may even embed `SchemaExport` in your application:

```
Configuration cfg = ....;
new SchemaExport(cfg).create(false, true);
```

## 20.1.3. Properties

Database properties may be specified

- as system properties with `-D<property>`
- in `hibernate.properties`
- in a named properties file with `--properties`

The needed properties are:

**Table 20.3. SchemaExport Connection Properties**

Property Name	Description
<code>hibernate.connection.driver_class</code>	jdbc driver class
<code>hibernate.connection.url</code>	jdbc url
<code>hibernate.connection.username</code>	database user
<code>hibernate.connection.password</code>	user password
<code>hibernate.dialect</code>	dialect

## 20.1.4. Using Ant

You can call `SchemaExport` from your Ant build script:

```
<target name="schemaexport">
  <taskdef name="schemaexport"
    classname="org.hibernate.tool.hbm2ddl.SchemaExportTask"
    classpathref="class.path" />

  <schemaexport
    properties="hibernate.properties"
    quiet="no"
    text="no"
    drop="no"
    delimiter=";"
    output="schema-export.sql">
    <fileset dir="src">
      <include name="**/*.hbm.xml" />
    </fileset>
  </schemaexport>
</target>
```

## 20.1.5. Incremental schema updates

The `SchemaUpdate` tool will update an existing schema with "incremental" changes. Note that `SchemaUpdate` depends heavily upon the JDBC metadata API, so it will not work with all JDBC drivers.

```
java -cp hibernate_classpaths org.hibernate.tool.hbm2ddl.SchemaUpdate options mapping_files
```

**Table 20.4.** `SchemaUpdate` Command Line Options

Option	Description
<code>--quiet</code>	don't output the script to stdout
<code>--properties=hibernate.properties</code>	read database properties from a file

You may embed `SchemaUpdate` in your application:

```
Configuration cfg = ....;
new SchemaUpdate(cfg).execute(false);
```

## 20.1.6. Using Ant for incremental schema updates

You can call `SchemaUpdate` from the Ant script:

```
<target name="schemaupdate">
  <taskdef name="schemaupdate"
    classname="org.hibernate.tool.hbm2ddl.SchemaUpdateTask"
    classpathref="class.path" />

  <schemaupdate
    properties="hibernate.properties"
    quiet="no">
    <fileset dir="src">
      <include name="**/*.hbm.xml" />
    </fileset>
  </schemaupdate>
</target>
```

---

## Chapter 21. Example: Parent/Child

One of the very first things that new users try to do with Hibernate is to model a parent / child type relationship. There are two different approaches to this. For various reasons the most convenient approach, especially for new users, is to model both `Parent` and `Child` as entity classes with a `<one-to-many>` association from `Parent` to `Child`. (The alternative approach is to declare the `Child` as a `<composite-element>`.) Now, it turns out that default semantics of a one to many association (in Hibernate) are much less close to the usual semantics of a parent / child relationship than those of a composite element mapping. We will explain how to use a *bidirectional one to many association with cascades* to model a parent / child relationship efficiently and elegantly. It's not at all difficult!

### 21.1. A note about collections

Hibernate collections are considered to be a logical part of their owning entity; never of the contained entities. This is a crucial distinction! It has the following consequences:

- When we remove / add an object from / to a collection, the version number of the collection owner is incremented.
- If an object that was removed from a collection is an instance of a value type (eg, a composite element), that object will cease to be persistent and its state will be completely removed from the database. Likewise, adding a value type instance to the collection will cause its state to be immediately persistent.
- On the other hand, if an entity is removed from a collection (a one-to-many or many-to-many association), it will not be deleted, by default. This behaviour is completely consistent - a change to the internal state of another entity should not cause the associated entity to vanish! Likewise, adding an entity to a collection does not cause that entity to become persistent, by default.

Instead, the default behaviour is that adding an entity to a collection merely creates a link between the two entities, while removing it removes the link. This is very appropriate for all sorts of cases. Where it is not appropriate at all is the case of a parent / child relationship, where the life of the child is bound to the lifecycle of the parent.

### 21.2. Bidirectional one-to-many

Suppose we start with a simple `<one-to-many>` association from `Parent` to `Child`.

```
<set name="children">
  <key column="parent_id"/>
  <one-to-many class="Child"/>
</set>
```

If we were to execute the following code

```
Parent p = .....;
Child c = new Child();
p.getChildren().add(c);
session.save(c);
session.flush();
```

Hibernate would issue two SQL statements:

- an `INSERT` to create the record for `c`
- an `UPDATE` to create the link from `p` to `c`

This is not only inefficient, but also violates any `NOT NULL` constraint on the `parent_id` column. We can fix the nullability constraint violation by specifying `not-null="true"` in the collection mapping:

```
<set name="children">
  <key column="parent_id" not-null="true"/>
  <one-to-many class="Child"/>
</set>
```

However, this is not the recommended solution.

The underlying cause of this behaviour is that the link (the foreign key `parent_id`) from `p` to `c` is not considered part of the state of the `Child` object and is therefore not created in the `INSERT`. So the solution is to make the link part of the `Child` mapping.

```
<many-to-one name="parent" column="parent_id" not-null="true"/>
```

(We also need to add the `parent` property to the `Child` class.)

Now that the `Child` entity is managing the state of the link, we tell the collection not to update the link. We use the `inverse` attribute.

```
<set name="children" inverse="true">
  <key column="parent_id"/>
  <one-to-many class="Child"/>
</set>
```

The following code would be used to add a new `Child`

```
Parent p = (Parent) session.load(Parent.class, pid);
Child c = new Child();
c.setParent(p);
p.getChildren().add(c);
session.save(c);
session.flush();
```

And now, only one SQL `INSERT` would be issued!

To tighten things up a bit, we could create an `addChild()` method of `Parent`.

```
public void addChild(Child c) {
    c.setParent(this);
    children.add(c);
}
```

Now, the code to add a `Child` looks like

```
Parent p = (Parent) session.load(Parent.class, pid);
Child c = new Child();
p.addChild(c);
session.save(c);
session.flush();
```

## 21.3. Cascading lifecycle



The explicit call to `save()` is still annoying. We will address this by using cascades.

```
<set name="children" inverse="true" cascade="all">
  <key column="parent_id"/>
  <one-to-many class="Child"/>
</set>
```

This simplifies the code above to

```
Parent p = (Parent) session.load(Parent.class, pid);
Child c = new Child();
p.addChild(c);
session.flush();
```

Similarly, we don't need to iterate over the children when saving or deleting a `Parent`. The following removes `p` and all its children from the database.

```
Parent p = (Parent) session.load(Parent.class, pid);
session.delete(p);
session.flush();
```

However, this code

```
Parent p = (Parent) session.load(Parent.class, pid);
Child c = (Child) p.getChildren().iterator().next();
p.getChildren().remove(c);
c.setParent(null);
session.flush();
```

will not remove `c` from the database; it will only remove the link to `p` (and cause a `NOT NULL` constraint violation, in this case). You need to explicitly `delete()` the `Child`.

```
Parent p = (Parent) session.load(Parent.class, pid);
Child c = (Child) p.getChildren().iterator().next();
p.getChildren().remove(c);
session.delete(c);
session.flush();
```

Now, in our case, a `Child` can't really exist without its parent. So if we remove a `Child` from the collection, we really do want it to be deleted. For this, we must use `cascade="all-delete-orphan"`.

```
<set name="children" inverse="true" cascade="all-delete-orphan">
  <key column="parent_id"/>
  <one-to-many class="Child"/>
</set>
```

Note: even though the collection mapping specifies `inverse="true"`, cascades are still processed by iterating the collection elements. So if you require that an object be saved, deleted or updated by cascade, you must add it to the collection. It is not enough to simply call `setParent()`.

## 21.4. Cascades and `unsaved-value`

Suppose we loaded up a `Parent` in one `Session`, made some changes in a UI action and wish to persist these changes in a new session by calling `update()`. The `Parent` will contain a collection of children and, since cascading update is enabled, Hibernate needs to know which children are newly instantiated and which represent existing rows in the database. Lets assume that both `Parent` and `Child` have autogenerated identifier properties of type `Long`. Hibernate will use the identifier and version/timestamp property value to determine which of the

children are new. (See Section 10.7, “Automatic state detection”.) *In Hibernate3, it is no longer necessary to specify an `unsaved-value` explicitly.*

The following code will update `parent` and `child` and insert `newChild`.

```
//parent and child were both loaded in a previous session
parent.addChild(child);
Child newChild = new Child();
parent.addChild(newChild);
session.update(parent);
session.flush();
```

Well, that's all very well for the case of a generated identifier, but what about assigned identifiers and composite identifiers? This is more difficult, since Hibernate can't use the identifier property to distinguish between a newly instantiated object (with an identifier assigned by the user) and an object loaded in a previous session. In this case, Hibernate will either use the timestamp or version property, or will actually query the second-level cache or, worst case, the database, to see if the row exists.

## 21.5. Conclusion

There is quite a bit to digest here and it might look confusing first time around. However, in practice, it all works out very nicely. Most Hibernate applications use the parent / child pattern in many places.

We mentioned an alternative in the first paragraph. None of the above issues exist in the case of `<composite-element>` mappings, which have exactly the semantics of a parent / child relationship. Unfortunately, there are two big limitations to composite element classes: composite elements may not own collections, and they should not be the child of any entity other than the unique parent.

---

# Chapter 22. Example: Weblog Application

## 22.1. Persistent Classes

The persistent classes represent a weblog, and an item posted in a weblog. They are to be modelled as a standard parent/child relationship, but we will use an ordered bag, instead of a set.

```
package eg;

import java.util.List;

public class Blog {
    private Long _id;
    private String _name;
    private List _items;

    public Long getId() {
        return _id;
    }
    public List.getItems() {
        return _items;
    }
    public String getName() {
        return _name;
    }
    public void setId(Long long1) {
        _id = long1;
    }
    public void setItems(List list) {
        _items = list;
    }
    public void setName(String string) {
        _name = string;
    }
}
```

```
package eg;

import java.text.DateFormat;
import java.util.Calendar;

public class BlogItem {
    private Long _id;
    private Calendar _datetime;
    private String _text;
    private String _title;
    private Blog _blog;

    public Blog getBlog() {
        return _blog;
    }
    public Calendar getDatetime() {
        return _datetime;
    }
    public Long getId() {
        return _id;
    }
    public String getText() {
        return _text;
    }
    public String getTitle() {
        return _title;
    }
    public void setBlog(Blog blog) {
        _blog = blog;
    }
}
```

```

    }
    public void setDatetime(Calendar calendar) {
        _datetime = calendar;
    }
    public void setId(Long long1) {
        _id = long1;
    }
    public void setText(String string) {
        _text = string;
    }
    public void setTitle(String string) {
        _title = string;
    }
}

```

## 22.2. Hibernate Mappings

The XML mappings should now be quite straightforward.

```

<?xml version="1.0"?>
<!DOCTYPE hibernate-mapping PUBLIC
    "-//Hibernate/Hibernate Mapping DTD 3.0//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">

<hibernate-mapping package="eg">

    <class
        name="Blog"
        table="BLOGS">

        <id
            name="id"
            column="BLOG_ID">

            <generator class="native"/>

        </id>

        <property
            name="name"
            column="NAME"
            not-null="true"
            unique="true"/>

        <bag
            name="items"
            inverse="true"
            order-by="DATE_TIME"
            cascade="all">

            <key column="BLOG_ID"/>
            <one-to-many class="BlogItem"/>

        </bag>

    </class>

</hibernate-mapping>

```

```

<?xml version="1.0"?>
<!DOCTYPE hibernate-mapping PUBLIC
    "-//Hibernate/Hibernate Mapping DTD 3.0//EN"
    "http://hibernate.sourceforge.net/hibernate-mapping-3.0.dtd">

<hibernate-mapping package="eg">

```

```

<class
  name="BlogItem"
  table="BLOG_ITEMS"
  dynamic-update="true">

  <id
    name="id"
    column="BLOG_ITEM_ID">

    <generator class="native"/>

  </id>

  <property
    name="title"
    column="TITLE"
    not-null="true"/>

  <property
    name="text"
    column="TEXT"
    not-null="true"/>

  <property
    name="datetime"
    column="DATE_TIME"
    not-null="true"/>

  <many-to-one
    name="blog"
    column="BLOG_ID"
    not-null="true"/>

</class>
</hibernate-mapping>

```

## 22.3. Hibernate Code

The following class demonstrates some of the kinds of things we can do with these classes, using Hibernate.

```

package eg;

import java.util.ArrayList;
import java.util.Calendar;
import java.util.Iterator;
import java.util.List;

import org.hibernate.HibernateException;
import org.hibernate.Query;
import org.hibernate.Session;
import org.hibernate.SessionFactory;
import org.hibernate.Transaction;
import org.hibernate.cfg.Configuration;
import org.hibernate.tool.hbm2ddl.SchemaExport;

public class BlogMain {

    private SessionFactory _sessions;

    public void configure() throws HibernateException {
        _sessions = new Configuration()
            .addClass(Blog.class)
            .addClass(BlogItem.class)
            .buildSessionFactory();
    }
}

```

```

public void exportTables() throws HibernateException {
    Configuration cfg = new Configuration()
        .addClass(Blog.class)
        .addClass(BlogItem.class);
    new SchemaExport(cfg).create(true, true);
}

public Blog createBlog(String name) throws HibernateException {

    Blog blog = new Blog();
    blog.setName(name);
    blog.setItems( new ArrayList() );

    Session session = _sessions.openSession();
    Transaction tx = null;
    try {
        tx = session.beginTransaction();
        session.persist(blog);
        tx.commit();
    }
    catch (HibernateException he) {
        if (tx!=null) tx.rollback();
        throw he;
    }
    finally {
        session.close();
    }
    return blog;
}

public BlogItem createBlogItem(Blog blog, String title, String text)
    throws HibernateException {

    BlogItem item = new BlogItem();
    item.setTitle(title);
    item.setText(text);
    item.setBlog(blog);
    item.setDatetime( Calendar.getInstance() );
    blog.getItems().add(item);

    Session session = _sessions.openSession();
    Transaction tx = null;
    try {
        tx = session.beginTransaction();
        session.update(blog);
        tx.commit();
    }
    catch (HibernateException he) {
        if (tx!=null) tx.rollback();
        throw he;
    }
    finally {
        session.close();
    }
    return item;
}

public BlogItem createBlogItem(Long blogid, String title, String text)
    throws HibernateException {

    BlogItem item = new BlogItem();
    item.setTitle(title);
    item.setText(text);
    item.setDatetime( Calendar.getInstance() );

    Session session = _sessions.openSession();
    Transaction tx = null;
    try {
        tx = session.beginTransaction();
        Blog blog = (Blog) session.load(Blog.class, blogid);
        item.setBlog(blog);

```

```

        blog.getItems().add(item);
        tx.commit();
    }
    catch (HibernateException he) {
        if (tx!=null) tx.rollback();
        throw he;
    }
    finally {
        session.close();
    }
    return item;
}

public void updateBlogItem(BlogItem item, String text)
    throws HibernateException {

    item.setText(text);

    Session session = _sessions.openSession();
    Transaction tx = null;
    try {
        tx = session.beginTransaction();
        session.update(item);
        tx.commit();
    }
    catch (HibernateException he) {
        if (tx!=null) tx.rollback();
        throw he;
    }
    finally {
        session.close();
    }
}

public void updateBlogItem(Long itemid, String text)
    throws HibernateException {

    Session session = _sessions.openSession();
    Transaction tx = null;
    try {
        tx = session.beginTransaction();
        BlogItem item = (BlogItem) session.load(BlogItem.class, itemid);
        item.setText(text);
        tx.commit();
    }
    catch (HibernateException he) {
        if (tx!=null) tx.rollback();
        throw he;
    }
    finally {
        session.close();
    }
}

public List listAllBlogNamesAndItemCounts(int max)
    throws HibernateException {

    Session session = _sessions.openSession();
    Transaction tx = null;
    List result = null;
    try {
        tx = session.beginTransaction();
        Query q = session.createQuery(
            "select blog.id, blog.name, count(blogItem) " +
            "from Blog as blog " +
            "left outer join blog.items as blogItem " +
            "group by blog.name, blog.id " +
            "order by max(blogItem.datetime)"
        );
        q.setMaxResults(max);
        result = q.list();
    }

```

```

        tx.commit();
    }
    catch (HibernateException he) {
        if (tx!=null) tx.rollback();
        throw he;
    }
    finally {
        session.close();
    }
    return result;
}

public Blog getBlogAndAllItems(Long blogid)
    throws HibernateException {

    Session session = _sessions.openSession();
    Transaction tx = null;
    Blog blog = null;
    try {
        tx = session.beginTransaction();
        Query q = session.createQuery(
            "from Blog as blog " +
            "left outer join fetch blog.items " +
            "where blog.id = :blogid"
        );
        q.setParameter("blogid", blogid);
        blog = (Blog) q.uniqueResult();
        tx.commit();
    }
    catch (HibernateException he) {
        if (tx!=null) tx.rollback();
        throw he;
    }
    finally {
        session.close();
    }
    return blog;
}

public List listBlogsAndRecentItems() throws HibernateException {

    Session session = _sessions.openSession();
    Transaction tx = null;
    List result = null;
    try {
        tx = session.beginTransaction();
        Query q = session.createQuery(
            "from Blog as blog " +
            "inner join blog.items as blogItem " +
            "where blogItem.datetime > :minDate"
        );

        Calendar cal = Calendar.getInstance();
        cal.roll(Calendar.MONTH, false);
        q.setCalendar("minDate", cal);

        result = q.list();
        tx.commit();
    }
    catch (HibernateException he) {
        if (tx!=null) tx.rollback();
        throw he;
    }
    finally {
        session.close();
    }
    return result;
}
}

```




# Chapter 23. Example: Various Mappings

This chapters shows off some more complex association mappings.

## 23.1. Employer/Employee

The following model of the relationship between `Employer` and `Employee` uses an actual entity class (`Employment`) to represent the association. This is done because there might be more than one period of employment for the same two parties. Components are used to model monetary values and employee names.



Heres a possible mapping document:

```
<hibernate-mapping>

  <class name="Employer" table="employers">
    <id name="id">
      <generator class="sequence">
        <param name="sequence">employer_id_seq</param>
      </generator>
    </id>
    <property name="name"/>
  </class>

  <class name="Employment" table="employment_periods">

    <id name="id">
      <generator class="sequence">
        <param name="sequence">employment_id_seq</param>
      </generator>
    </id>
    <property name="startDate" column="start_date"/>
    <property name="endDate" column="end_date"/>

    <component name="hourlyRate" class="MonetaryAmount">
      <property name="amount">
        <column name="hourly_rate" sql-type="NUMERIC(12, 2)"/>
      </property>
      <property name="currency" length="12"/>
    </component>

    <many-to-one name="employer" column="employer_id" not-null="true"/>
    <many-to-one name="employee" column="employee_id" not-null="true"/>

  </class>

  <class name="Employee" table="employees">
    <id name="id">
```

```
        <generator class="sequence">
            <param name="sequence">employee_id_seq</param>
        </generator>
    </id>
    <property name="taxfileNumber"/>
    <component name="name" class="Name">
        <property name="firstName"/>
        <property name="initial"/>
        <property name="lastName"/>
    </component>
</class>

</hibernate-mapping>
```

And heres the table schema generated by SchemaExport.

```
create table employers (
    id BIGINT not null,
    name VARCHAR(255),
    primary key (id)
)


create table employment_periods (
    id BIGINT not null,
    hourly_rate NUMERIC(12, 2),
    currency VARCHAR(12),
    employee_id BIGINT not null,
    employer_id BIGINT not null,
    end_date TIMESTAMP,
    start_date TIMESTAMP,
    primary key (id)
)

create table employees (
    id BIGINT not null,
    firstName VARCHAR(255),
    initial CHAR(1),
    lastName VARCHAR(255),
    taxfileNumber VARCHAR(255),
    primary key (id)
)

alter table employment_periods
    add constraint employment_periodsFK0 foreign key (employer_id) references employers
alter table employment_periods
    add constraint employment_periodsFK1 foreign key (employee_id) references employees
create sequence employee_id_seq
create sequence employment_id_seq
create sequence employer_id_seq
```

## 23.2. Author/Work

Consider the following model of the relationships between `Work`, `Author` and `Person`. We represent the relationship between `Work` and `Author` as a many-to-many association. We choose to represent the relationship between `Author` and `Person` as one-to-one association. Another possibility would be to have `Author` extend `Person`.



The following mapping document correctly represents these relationships:

```

<hibernate-mapping>

  <class name="Work" table="works" discriminator-value="W">

    <id name="id" column="id">
      <generator class="native"/>
    </id>
    <discriminator column="type" type="character"/>

    <property name="title"/>
    <set name="authors" table="author_work">
      <key column name="work_id"/>
      <many-to-many class="Author" column name="author_id"/>
    </set>

    <subclass name="Book" discriminator-value="B">
      <property name="text"/>
    </subclass>

    <subclass name="Song" discriminator-value="S">
      <property name="tempo"/>
      <property name="genre"/>
    </subclass>

  </class>

  <class name="Author" table="authors">

    <id name="id" column="id">
      <!-- The Author must have the same identifier as the Person -->
      <generator class="assigned"/>
    </id>

    <property name="alias"/>
    <one-to-one name="person" constrained="true"/>

    <set name="works" table="author_work" inverse="true">
      <key column="author_id"/>
      <many-to-many class="Work" column="work_id"/>
    </set>

  </class>

```

```
<class name="Person" table="persons">
  <id name="id" column="id">
    <generator class="native"/>
  </id>
  <property name="name"/>
</class>

</hibernate-mapping>
```

There are four tables in this mapping. `works`, `authors` and `persons` hold work, author and person data respectively. `author_work` is an association table linking authors to works. Heres the table schema, as generated by SchemaExport.

```
create table works (
  id BIGINT not null generated by default as identity,
  tempo FLOAT,
  genre VARCHAR(255),
  text INTEGER,
  title VARCHAR(255),
  type CHAR(1) not null,
  primary key (id)
)

create table author_work (
  author_id BIGINT not null,
  work_id BIGINT not null,
  primary key (work_id, author_id)
)


create table authors (
  id BIGINT not null generated by default as identity,
  alias VARCHAR(255),
  primary key (id)
)

create table persons (
  id BIGINT not null generated by default as identity,
  name VARCHAR(255),
  primary key (id)
)

alter table authors
  add constraint authorsFK0 foreign key (id) references persons
alter table author_work
  add constraint author_workFK0 foreign key (author_id) references authors
alter table author_work
  add constraint author_workFK1 foreign key (work_id) references works
```

## 23.3. Customer/Order/Product

Now consider a model of the relationships between `Customer`, `Order` and `LineItem` and `Product`. There is a one-to-many association between `Customer` and `Order`, but how should we represent `Order` / `LineItem` / `Product`? I've chosen to map `LineItem` as an association class representing the many-to-many association between `Order` and `Product`. In Hibernate, this is called a composite element.



The mapping document:

```

<hibernate-mapping>

  <class name="Customer" table="customers">
    <id name="id">
      <generator class="native"/>
    </id>
    <property name="name"/>
    <set name="orders" inverse="true">
      <key column="customer_id"/>
      <one-to-many class="Order"/>
    </set>
  </class>

  <class name="Order" table="orders">
    <id name="id">
      <generator class="native"/>
    </id>
    <property name="date"/>
    <many-to-one name="customer" column="customer_id"/>
    <list name="lineItems" table="line_items">
      <key column="order_id"/>
      <list-index column="line_number"/>
      <composite-element class="LineItem">
        <property name="quantity"/>
        <many-to-one name="product" column="product_id"/>
      </composite-element>
    </list>
  </class>

  <class name="Product" table="products">
    <id name="id">
      <generator class="native"/>
    </id>
    <property name="serialNumber"/>
  </class>

</hibernate-mapping>

```

customers, orders, line\_items and products hold customer, order, order line item and product data respectively. line\_items also acts as an association table linking orders with products.

```

create table customers (
  id BIGINT not null generated by default as identity,
  name VARCHAR(255),
  primary key (id)
)

create table orders (
  id BIGINT not null generated by default as identity,
  customer_id BIGINT,
  date TIMESTAMP,
  primary key (id)
)

create table line_items (
  line_number INTEGER not null,

```

```
    order_id BIGINT not null,
    product_id BIGINT,
    quantity INTEGER,
    primary key (order_id, line_number)
)

create table products (
    id BIGINT not null generated by default as identity,
    serialNumber VARCHAR(255),
    primary key (id)
)

alter table orders
    add constraint ordersFK0 foreign key (customer_id) references customers
alter table line_items
    add constraint line_itemsFK0 foreign key (product_id) references products
alter table line_items
    add constraint line_itemsFK1 foreign key (order_id) references orders
```

## 23.4. Miscellaneous example mappings

These examples are all taken from the Hibernate test suite. You will find many other useful example mappings there. Have a look at the `src` folder of the Hibernate distribution.

TODO: put words around this stuff

### 23.4.1. "Typed" one-to-one association

```
<class name="Person">
  <id name="name"/>
  <one-to-one name="address"
    cascade="all">
    <formula>name</formula>
    <formula>'HOME'</formula>
  </one-to-one>
  <one-to-one name="mailingAddress"
    cascade="all">
    <formula>name</formula>
    <formula>'MAILING'</formula>
  </one-to-one>
</class>

<class name="Address" batch-size="2"
  check="addressType in ('MAILING', 'HOME', 'BUSINESS')">
  <composite-id>
    <key-many-to-one name="person"
      column="personName"/>
    <key-property name="type"
      column="addressType"/>
  </composite-id>
  <property name="street" type="text"/>
  <property name="state"/>
  <property name="zip"/>
</class>
```

### 23.4.2. Composite key example

```
<class name="Customer">

  <id name="customerId"
    length="10">
    <generator class="assigned"/>
```

```

</id>

<property name="name" not-null="true" length="100"/>
<property name="address" not-null="true" length="200"/>

<list name="orders"
      inverse="true"
      cascade="save-update">
  <key column="customerId"/>
  <index column="orderNumber"/>
  <one-to-many class="Order"/>
</list>

</class>

<class name="Order" table="CustomerOrder" lazy="true">
  <synchronize table="LineItem"/>
  <synchronize table="Product"/>

  <composite-id name="id"
                class="Order$Id">
    <key-property name="customerId" length="10"/>
    <key-property name="orderNumber"/>
  </composite-id>

  <property name="orderDate"
            type="calendar_date"
            not-null="true"/>

  <property name="total">
    <formula>
      ( select sum(li.quantity*p.price)
        from LineItem li, Product p
        where li.productId = p.productId
              and li.customerId = customerId
              and li.orderNumber = orderNumber )
    </formula>
  </property>

  <many-to-one name="customer"
               column="customerId"
               insert="false"
               update="false"
               not-null="true"/>

  <bag name="lineItems"
       fetch="join"
       inverse="true"
       cascade="save-update">
    <key>
      <column name="customerId"/>
      <column name="orderNumber"/>
    </key>
    <one-to-many class="LineItem"/>
  </bag>

</class>

<class name="LineItem">

  <composite-id name="id"
                class="LineItem$Id">
    <key-property name="customerId" length="10"/>
    <key-property name="orderNumber"/>
    <key-property name="productId" length="10"/>
  </composite-id>

  <property name="quantity"/>

  <many-to-one name="order"
               insert="false"

```

```

        update="false"
        not-null="true">
        <column name="customerId"/>
        <column name="orderId"/>
    </many-to-one>

    <many-to-one name="product"
        insert="false"
        update="false"
        not-null="true"
        column="productId"/>

</class>

<class name="Product">
    <synchronize table="LineItem"/>

    <id name="productId"
        length="10">
        <generator class="assigned"/>
    </id>

    <property name="description"
        not-null="true"
        length="200"/>
    <property name="price" length="3"/>
    <property name="numberAvailable"/>

    <property name="numberOrdered">
        <formula>
            ( select sum(li.quantity)
              from LineItem li
              where li.productId = productId )
        </formula>
    </property>

</class>

```

### 23.4.3. Content based discrimination

```

<class name="Person"
    discriminator-value="P">

    <id name="id"
        column="person_id"
        unsaved-value="0">
        <generator class="native"/>
    </id>

    <discriminator
        type="character">
        <formula>
            case
                when title is not null then 'E'
                when salesperson is not null then 'C'
                else 'P'
            end
        </formula>
    </discriminator>

    <property name="name"
        not-null="true"
        length="80"/>

    <property name="sex"
        not-null="true"
        update="false"/>

```



```

<component name="address">
  <property name="address"/>
  <property name="zip"/>
  <property name="country"/>
</component>

<subclass name="Employee"
  discriminator-value="E">
  <property name="title"
    length="20"/>
  <property name="salary"/>
  <many-to-one name="manager"/>
</subclass>

<subclass name="Customer"
  discriminator-value="C">
  <property name="comments"/>
  <many-to-one name="salesperson"/>
</subclass>

</class>

```

### 23.4.4. Associations on alternate keys

```

<class name="Person">

  <id name="id">
    <generator class="hilo"/>
  </id>

  <property name="name" length="100"/>

  <one-to-one name="address"
    property-ref="person"
    cascade="all"
    fetch="join"/>

  <set name="accounts"
    inverse="true">
    <key column="userId"
      property-ref="userId"/>
    <one-to-many class="Account"/>
  </set>

  <property name="userId" length="8"/>
</class>

<class name="Address">

  <id name="id">
    <generator class="hilo"/>
  </id>

  <property name="address" length="300"/>
  <property name="zip" length="5"/>
  <property name="country" length="25"/>
  <many-to-one name="person" unique="true" not-null="true"/>
</class>

<class name="Account">
  <id name="accountId" length="32">
    <generator class="uuid.hex"/>
  </id>

  <many-to-one name="user"

```

```
        column="userId"  
        property-ref="userId"/>  
  
    <property name="type" not-null="true"/>  
  
</class>
```

---

## Chapter 24. Best Practices

Write fine-grained classes and map them using `<component>`.

Use an `Address` class to encapsulate `street`, `suburb`, `state`, `postcode`. This encourages code reuse and simplifies refactoring.

Declare identifier properties on persistent classes.

Hibernate makes identifier properties optional. There are all sorts of reasons why you should use them. We recommend that identifiers be 'synthetic' (generated, with no business meaning). It doesn't make a difference if you use `long` or `java.lang.Long`; primitives might be syntactically easier to handle though.

Place each class mapping in its own file.

Don't use a single monolithic mapping document. Map `com.eg.Foo` in the file `com/eg/Foo.hbm.xml`. This makes particularly good sense in a team environment.

Load mappings as resources.

Deploy the mappings along with the classes they map.

Consider externalising query strings.

This is a good practice if your queries call non-ANSI-standard SQL functions. Externalising the query strings to mapping files will make the application more portable.

Use bind variables.

As in JDBC, always replace non-constant values by `"?"`. Never use string manipulation to bind a non-constant value in a query! Even better, consider using named parameters in queries.

Don't manage your own JDBC connections.

Hibernate lets the application manage JDBC connections. This approach should be considered a last-resort. If you can't use the built-in connections providers, consider providing your own implementation of `org.hibernate.connection.ConnectionProvider`.

Consider using a custom type.

Suppose you have a Java type, say from some library, that needs to be persisted but doesn't provide the accessors needed to map it as a component. You should consider implementing `org.hibernate.UserType`. This approach frees the application code from implementing transformations to / from a Hibernate type.

Use hand-coded JDBC in bottlenecks.

In performance-critical areas of the system, some kinds of operations (eg. mass update / delete) might benefit from direct JDBC. But please, wait until you *know* something is a bottleneck. And don't assume that direct JDBC is necessarily faster. If need to use direct JDBC, it might be worth opening a Hibernate `Session` and using that SQL connection. That way you can still use the same transaction strategy and underlying connection provider.

Understand `Session` flushing.

From time to time the `Session` synchronizes its persistent state with the database. Performance will be affected if this process occurs too often. You may sometimes minimize unnecessary flushing by disabling automatic flushing or even by changing the order of queries and other operations within a particular transaction.

In a three tiered architecture, consider using `saveOrUpdate()`.

When using a servlet / session bean architecture, you could pass persistent objects loaded in the session bean to and from the servlet / JSP layer. Use a new session to service each request. Use `Session.update()` or `Session.saveOrUpdate()` to update the persistent state of an object.

In a two tiered architecture, consider using session disconnection.

Database Transactions have to be as short as possible for best scalability. However, it is often necessary to implement long running Application Transactions, a single unit-of-work from the point of view of a user. This Application Transaction might span several client requests and response cycles. Either use Detached Objects or, in two tiered architectures, simply disconnect the Hibernate Session from the JDBC connection and reconnect it for each subsequent request. Never use a single Session for more than one Application Transaction usecase, otherwise, you will run into stale data.

Don't treat exceptions as recoverable.

This is more of a necessary practice than a "best" practice. When an exception occurs, roll back the `Transaction` and close the `Session`. If you don't, Hibernate can't guarantee that in-memory state accurately represents persistent state. As a special case of this, do not use `Session.load()` to determine if an instance with the given identifier exists on the database; use `find()` instead.

Prefer lazy fetching for associations.

Use eager (outer-join) fetching sparingly. Use proxies and/or lazy collections for most associations to classes that are not cached at the JVM-level. For associations to cached classes, where there is a high probability of a cache hit, explicitly disable eager fetching using `outer-join="false"`. When an outer-join fetch is appropriate to a particular use case, use a query with a `left join`.

Consider abstracting your business logic from Hibernate.

Hide (Hibernate) data-access code behind an interface. Combine the *DAO* and *Thread Local Session* patterns. You can even have some classes persisted by handcoded JDBC, associated to Hibernate via a `UserType`. (This advice is intended for "sufficiently large" applications; it is not appropriate for an application with five tables!)

Implement `equals()` and `hashCode()` using a unique business key.

If you compare objects outside of the Session scope, you have to implement `equals()` and `hashCode()`. Inside the Session scope, Java object identity is guaranteed. If you implement these methods, never ever use the database identifier! A transient object doesn't have an identifier value and Hibernate would assign a value when the object is saved. If the object is in a Set while being saved, the hash code changes, breaking the contract. To implement `equals()` and `hashCode()`, use a unique business key, that is, compare a unique combination of class properties. Remember that this key has to be stable and unique only while the object is in a Set, not for the whole lifetime (not as stable as a database primary key). Never use collections in the `equals()` comparison (lazy loading) and be careful with other associated classes that might be proxied.

Don't use exotic association mappings.

Good usecases for a real many-to-many associations are rare. Most of the time you need additional information stored in the "link table". In this case, it is much better to use two one-to-many associations to an intermediate link class. In fact, we think that most associations are one-to-many and many-to-one, you should be careful when using any other association style and ask yourself if it is really necessary.